

# CS5344: CLASSIFICATION II: TREES ENSEMBLES, LOGISTICS REGRESSION AND DEEP LEARNING

Anthony Tung  
School of Computing  
National University of Singapore

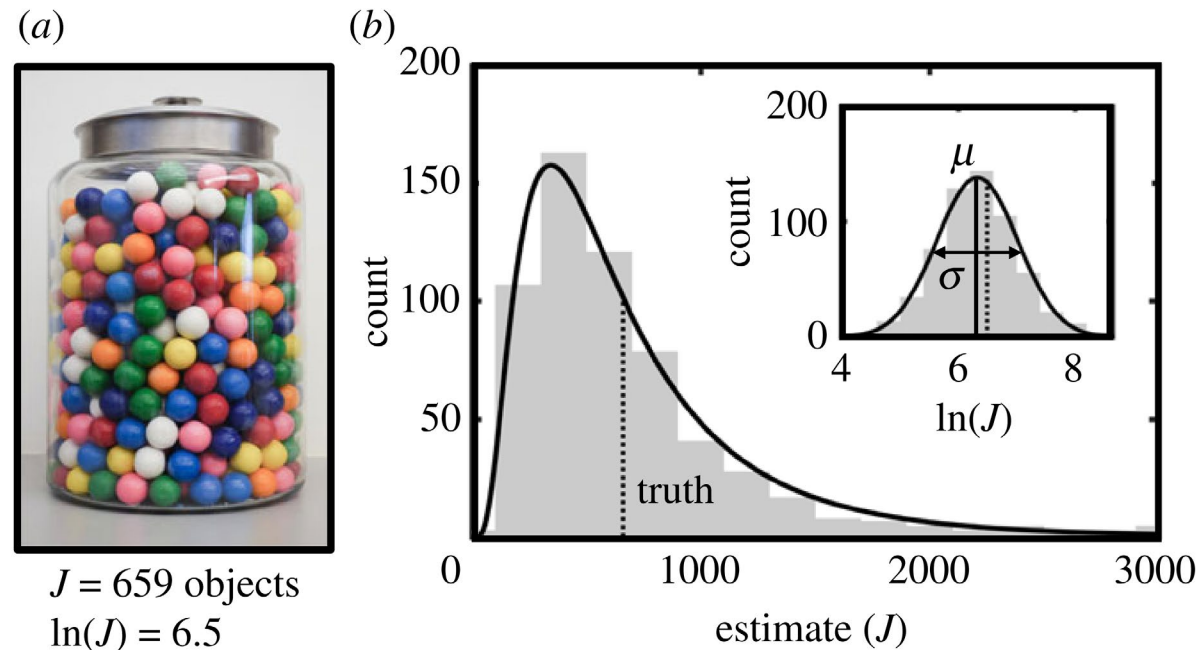
Slide Credit: Bryan Hooi, Wang Wei, Ng See Kiong,  
Wynne Hsu, Chris von der Weth

# CLASSIFICATION OVERVIEW

1. Problem Setup
2. Evaluating Classifiers
3. Nearest Neighbor Methods
4. **Trees and Ensembles**
  - a) Decision Trees
  - b) Bagging and Random Forests
  - c) Boosting and Gradient Boosting Machines
5. Logistic Regression
6. Deep Learning

# BAGGING: MOTIVATION ('WISDOM OF THE CROWDS')

**Wisdom of the crowds** refers to the idea that large groups of people are collectively smarter than individual experts.



# BOOTSTRAP SAMPLING

**Goal:** generate randomly sampled “versions” of the dataset

ID	Home Owner	Marital Status	Annual Income	Defaulted Borrower
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	170K	No
4	Yes	Married	120K	No
5	No	Single	75K	Yes
6	No	Married	160K	No
7	No	Single	50K	Yes

**Original Data**

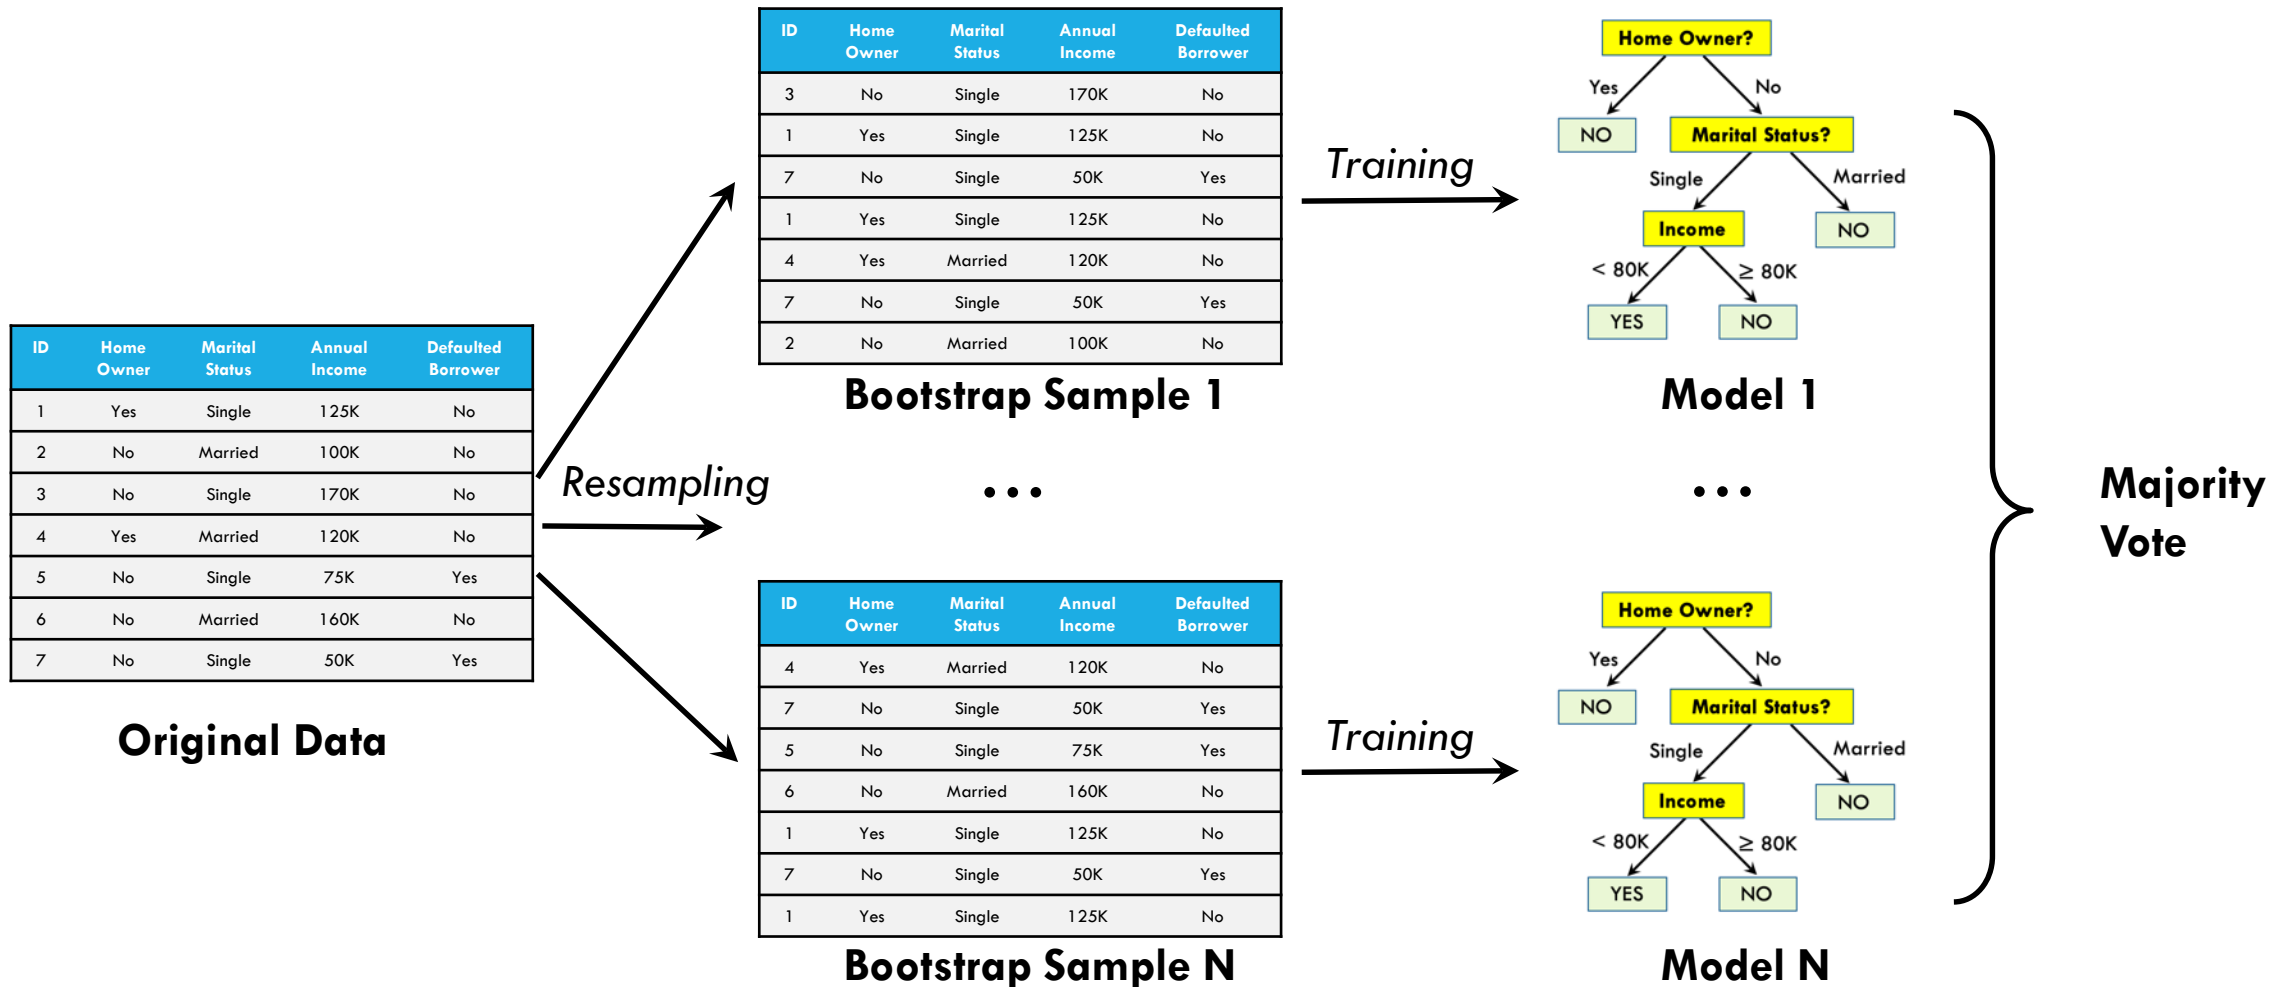


Resample same-sized dataset, **with replacement**

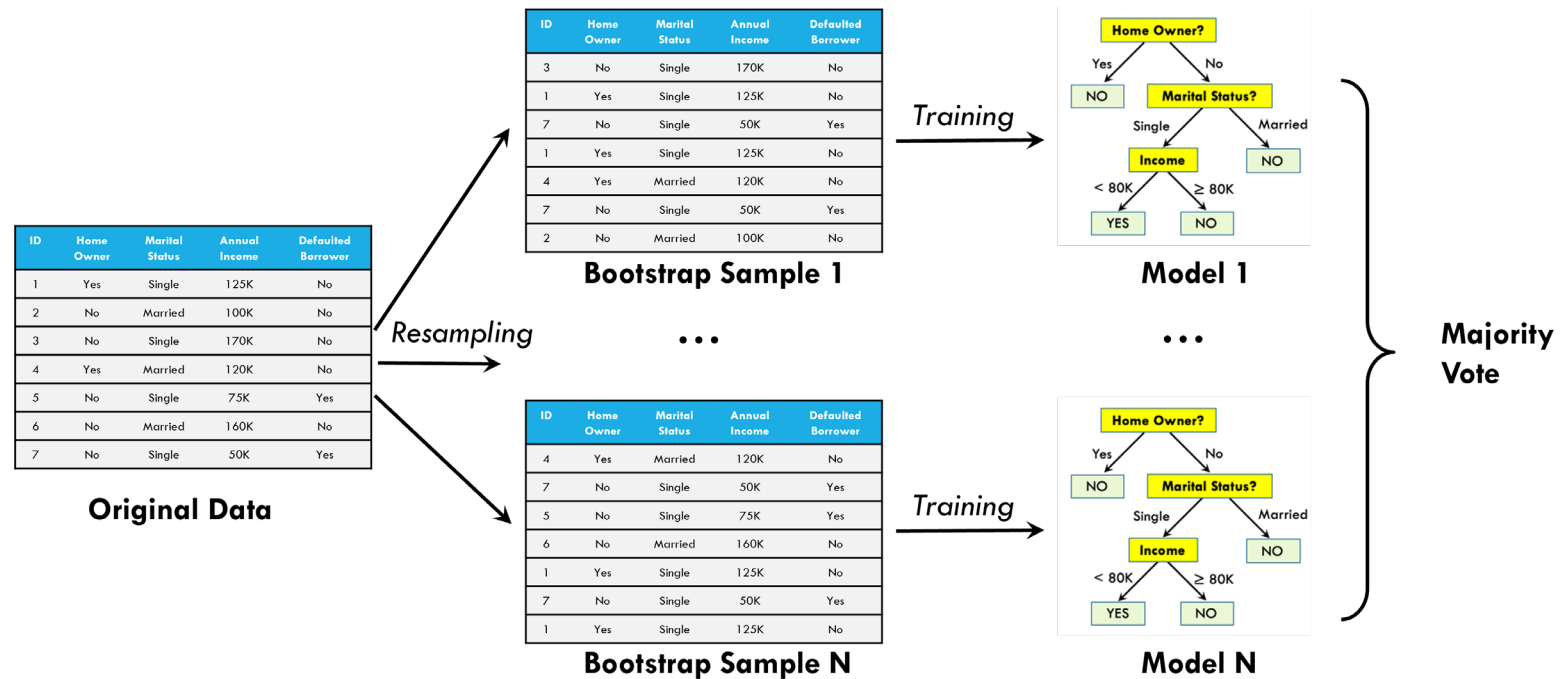
ID	Home Owner	Marital Status	Annual Income	Defaulted Borrower
3	No	Single	170K	No
1	Yes	Single	125K	No
7	No	Single	50K	Yes
1	Yes	Single	125K	No
4	Yes	Married	120K	No
7	No	Single	50K	Yes
2	No	Married	100K	No

**A Bootstrap Sample**

# BOOTSTRAP AGGREGATION (“BAGGING”)



# RANDOM FORESTS: BOOTSTRAP RESAMPLING ("ROW SAMPLING")

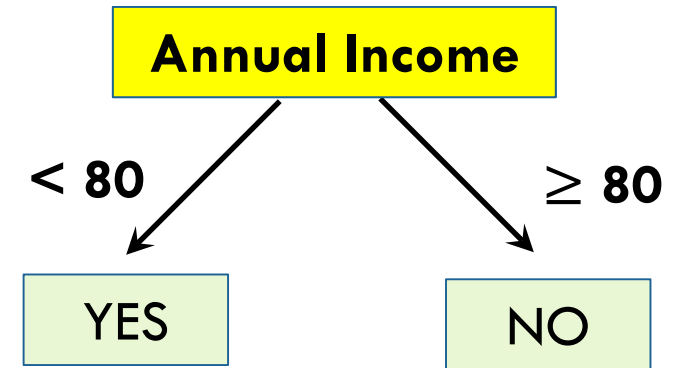


**Bootstrap resampling:** Random forests train an ensemble of decision trees on N bootstrap samples. (Sometimes, we take bootstrap samples of size smaller than the original dataset size, e.g., based on a user-specified parameter).

# RANDOM FORESTS: FEATURE SAMPLING ("COLUMN SAMPLING")

$d$

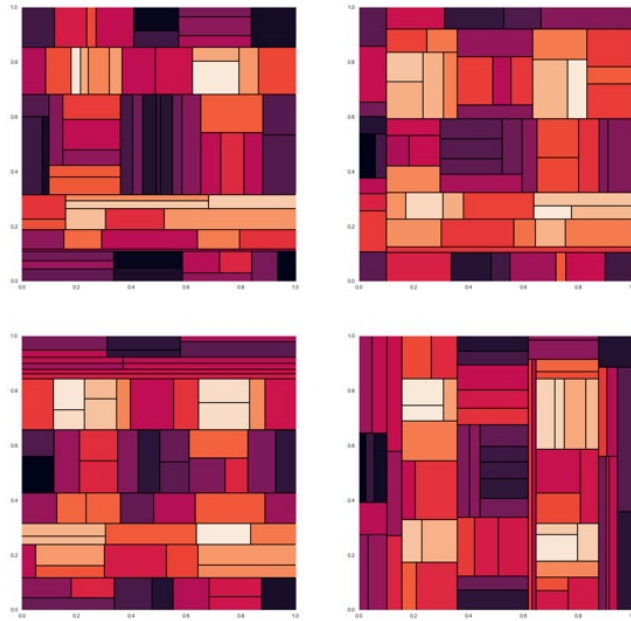
ID	Home Owner	Marital Status	Annual Income	Defaulted Borrower
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	170K	No
4	Yes	Married	120K	No
5	No	Single	75K	Yes
6	No	Married	160K	No
7	No	Single	50K	Yes



**Feature sampling:** During each split, instead of considering **all**  $d$  features, we only consider a **random subset** of features, usually of size  $\sqrt{d}$ .

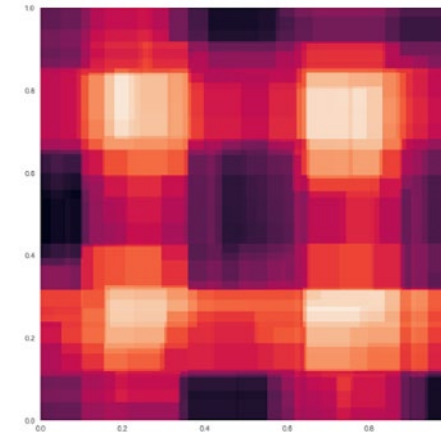
# WHY DO RANDOM FORESTS WORK WELL?

Individual decision tree predictions



Averaging

Random Forest predictions



**Variance Reduction:** individual trees can overfit and give highly variable output. But when averaging them, the predictions are smoother and perform better on test data.

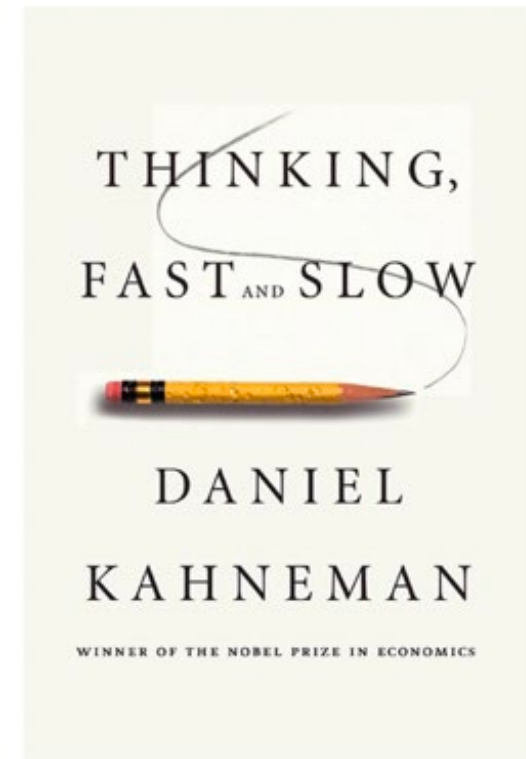
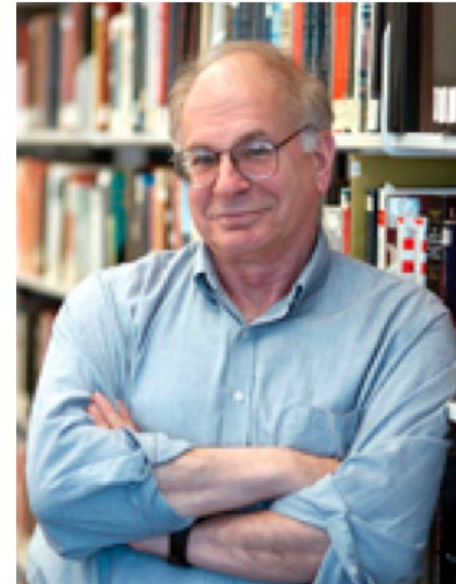
Randomization is important in making the decision trees **decorrelated**.



# IMPORTANCE OF DECORRELATION

“To derive the most useful information from multiple sources of evidence, you should always try to make these sources independent of each other.”

“A simple rule can help: before an issue is discussed, all members of the committee should be asked to write a very brief summary of their position. This procedure makes good use of the value of the **diversity of knowledge and opinion in the group**. The standard practice of open discussion gives too much weight to the opinions of those who speak early and assertively, causing others to line up behind them.”



Daniel Kahnemann, *Thinking Fast and Slow*

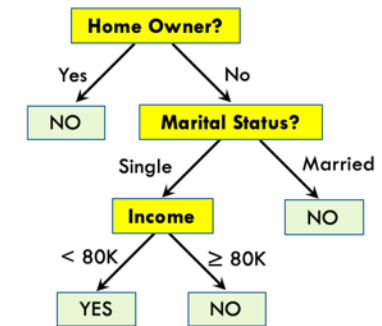
# PROS AND CONS (VS. DECISION TREES)

## Pros

- *Variance Reduction*: Ensembling many decision trees leads to more stable and accurate predictions
- Accuracy is fairly close to state of the art
- Parallelizable
- Not much tuning required

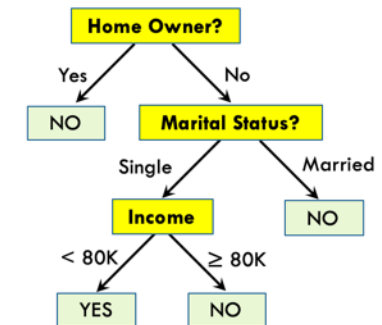
## Cons

- Less interpretable than decision trees
- Slower than decision trees



**Model 1**

...



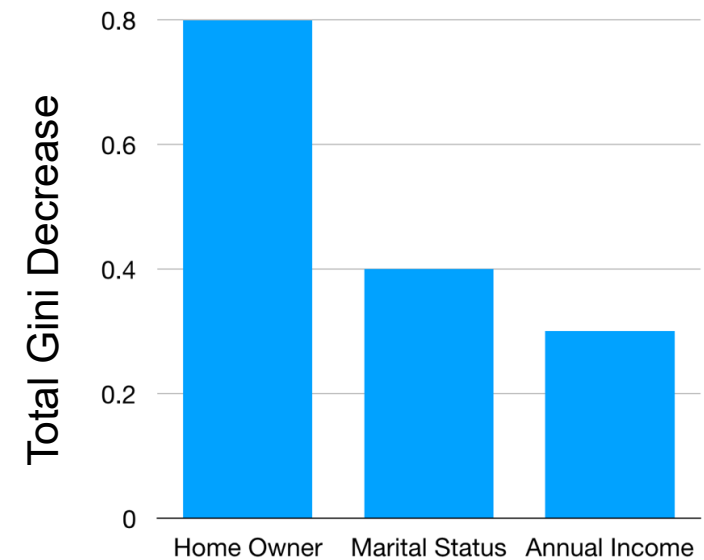
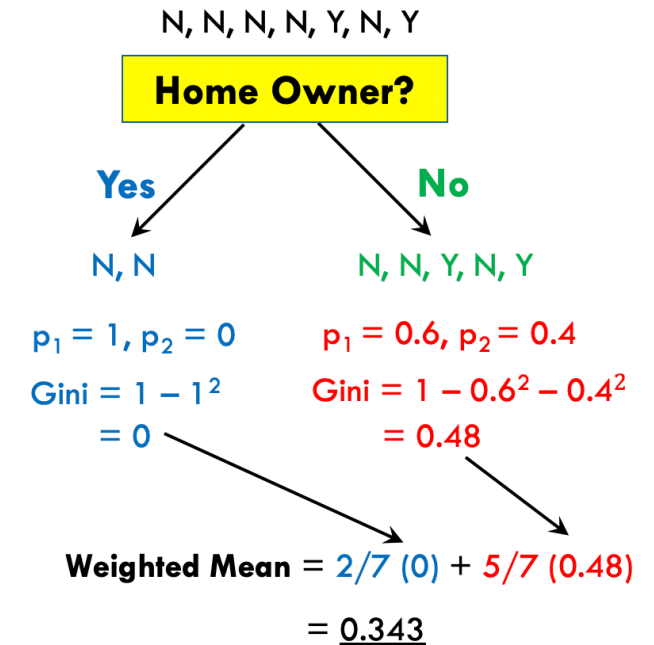
**Model N**

# VARIABLE IMPORTANCE PLOTS

The importance of each variable is measured by the **total reduction in Gini index** brought about by that feature.

More important variables result in greater decrease in Gini index on average.

Computing variable importance is useful for **feature selection**.

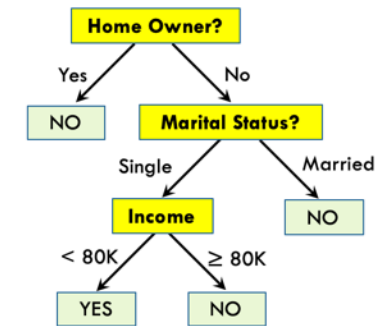




# QUIZ: RANDOM FOREST HYPERPARAMETERS

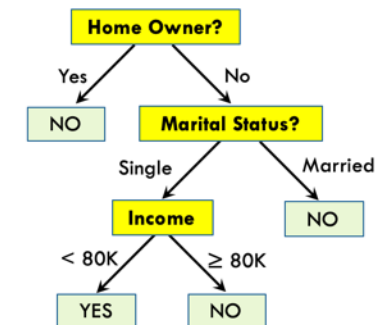
**Q:** Which of the following tends to reduce overfitting?

1. Increasing number of trees
2. Increasing depth of trees
3. Decreasing the number of features considered when selecting each split



**Model 1**

...



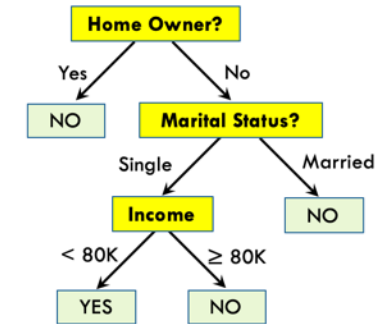
**Model N**



# QUIZ: RANDOM FOREST HYPERPARAMETERS

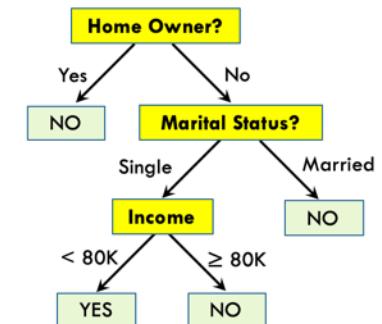
**Q:** Which of the following tends to reduce overfitting?

1. Increasing number of trees
2. Increasing depth of trees
3. Decreasing the number of features considered when selecting each split



**Model 1**

...



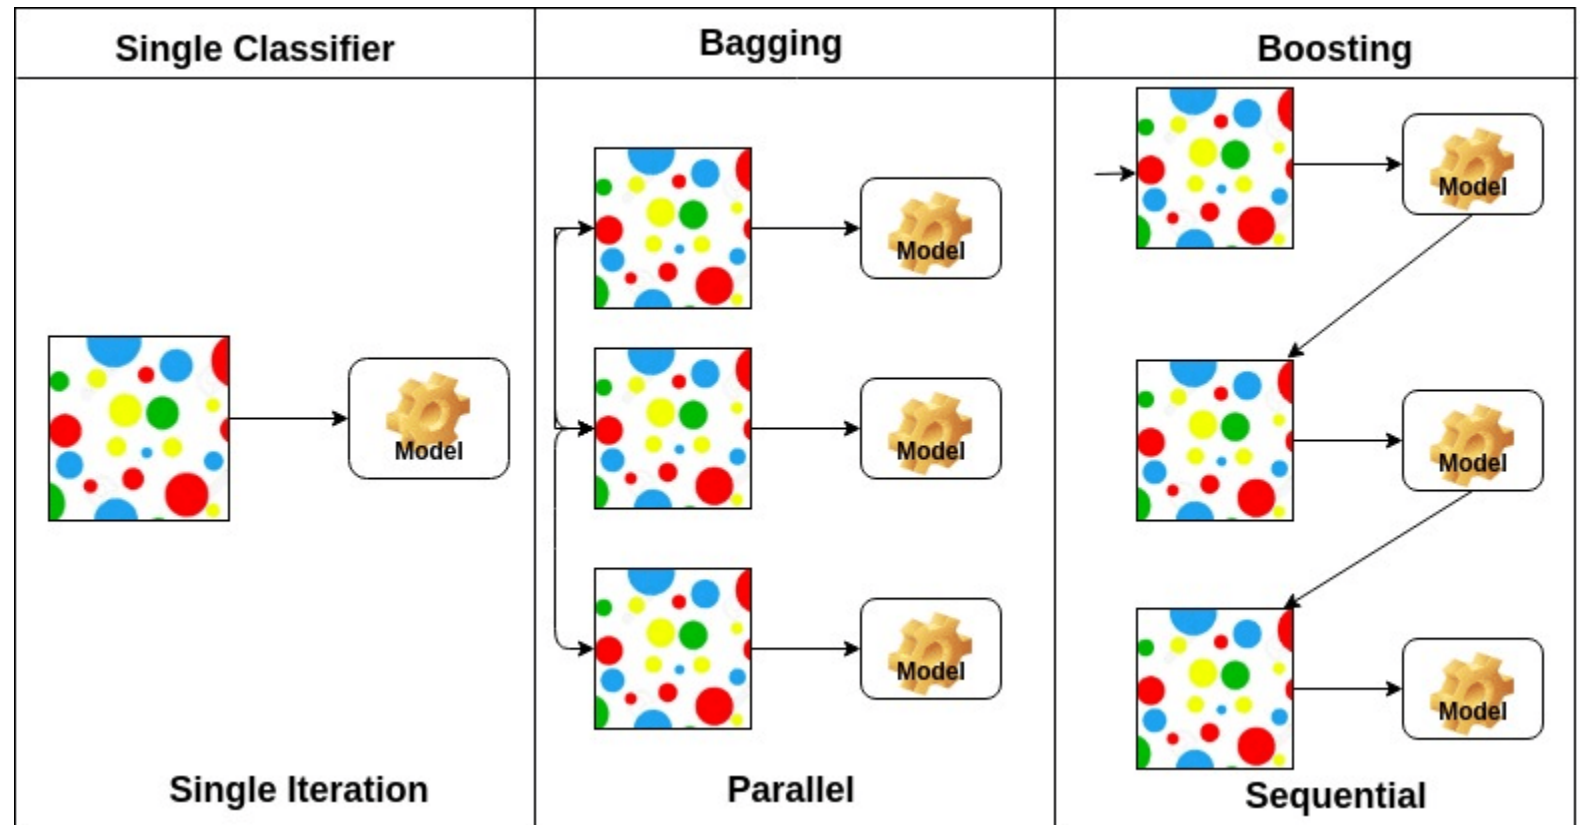
**Model N**

# CLASSIFICATION OVERVIEW

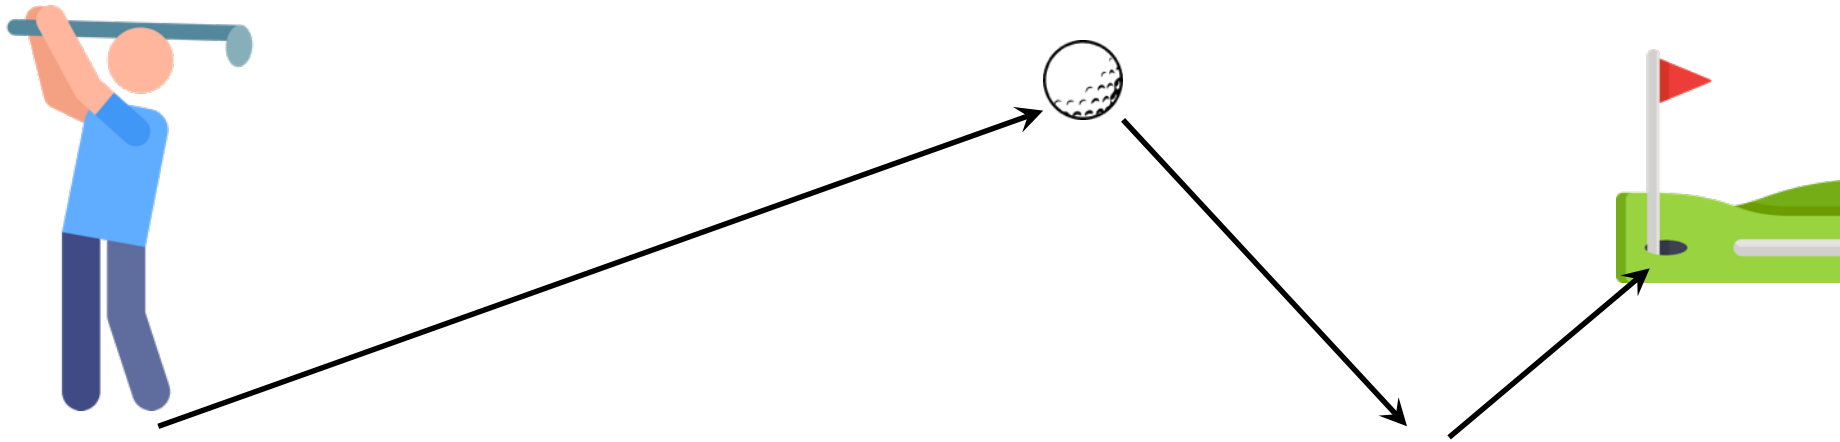
1. Problem Setup
2. Evaluating Classifiers
3. Nearest Neighbor Methods
4. **Trees and Ensembles**
  - a) Decision Trees
  - b) Bagging and Random Forests
  - c) Boosting and Gradient Boosting Machines
5. Logistic Regression
6. Deep Learning

# BAGGING VS BOOSTING

Bagging operates on bootstrap samples **independently**, but boosting is **sequential**: each new classifier “corrects for the mistakes” of the previous classifiers.



# BOOSTING: ITERATIVE PROCESS



Each new classifier is trained to “correct for the mistakes” of the previous set of classifiers, gradually getting closer to the ideal classifier



# GRADIENT BOOSTED DECISION TREES

Current Fitted Model

ID	Home Owner	Marital Status	Annual Income	y	$r_0$
1	Yes	Single	125K	0	0
2	No	Married	100K	0.5	0.5
3	No	Single	170K	0.3	0.3
4	Yes	Married	120K	1	1
5	No	Single	75K	0.9	0.9
6	No	Married	160K	0.4	0.4
7	No	Single	50K	0.3	0.3

**Original Data (X)**

**Residual  
(Initial)**

1. Initial residuals are the original response data.

# GRADIENT BOOSTED DECISION TREES

Current Fitted Model

ID	Home Owner	Marital Status	Annual Income	y	$r_0$	$f_1$
1	Yes	Single	125K	0	0	0.2
2	No	Married	100K	0.5	0.5	0.3
3	No	Single	170K	0.3	0.3	0.3
4	Yes	Married	120K	1	1	1
5	No	Single	75K	0.9	0.9	1
6	No	Married	160K	0.4	0.4	0.4
7	No	Single	50K	0.3	0.3	0.3

Fit tree with  
response  $r_0$

Original Data (X)

Residual  
(Initial)

2. Fit a decision tree to predict residuals  $r_0$ .

# GRADIENT BOOSTED DECISION TREES

Learning  
rate

$\lambda$   
 $= 0.5$



Current Fitted Model

ID	Home Owner	Marital Status	Annual Income	y
1	Yes	Single	125K	0
2	No	Married	100K	0.5
3	No	Single	170K	0.3
4	Yes	Married	120K	1
5	No	Single	75K	0.9
6	No	Married	160K	0.4
7	No	Single	50K	0.3

Original Data (X)

$r_0$
0
0.5
0.3
1
0.9
0.4
0.3

Residual  
(Initial)

Fit tree with  
response  $r_0$

$f_1$
0.2
0.3
0.3
1
1
0.4
0.3

$\lambda f_1$
0.1
0.15
0.15
0.5
0.5
0.2
0.15

3. Compute 'damped predictions', i.e.,  $\lambda$  times decision tree output.

# GRADIENT BOOSTED DECISION TREES

Learning  
rate

$\lambda$   
= 0.5



Current Fitted Model

ID	Home Owner	Marital Status	Annual Income	y
1	Yes	Single	125K	0
2	No	Married	100K	0.5
3	No	Single	170K	0.3
4	Yes	Married	120K	1
5	No	Single	75K	0.9
6	No	Married	160K	0.4
7	No	Single	50K	0.3

Original Data (X)

$r_0$
0
0.5
0.3
1
0.9
0.4
0.3

Residual  
(Initial)

Fit tree with  
response  $r_0$

$f_1$
0.2
0.3
0.3
1
1
0.4
0.3

$\lambda f_1$
0.1
0.15
0.15
0.5
0.5
0.2
0.15

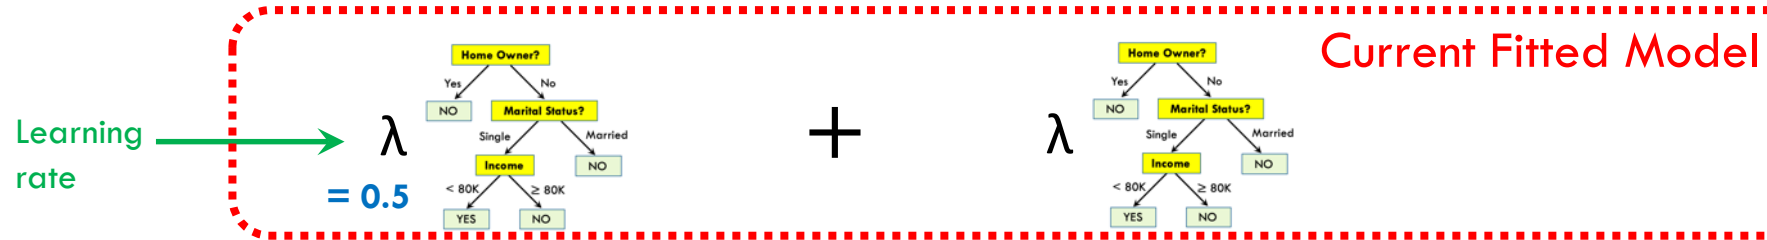
$r_1$
-0.1
0.35
0.15
0.5
0.4
0.2
0.15

Residual

subtract

4. Compute new residual:  $r_1 = r_0 - \lambda f_1$ .

# GRADIENT BOOSTED DECISION TREES



ID	Home Owner	Marital Status	Annual Income	y
1	Yes	Single	125K	0
2	No	Married	100K	0.5
3	No	Single	170K	0.3
4	Yes	Married	120K	1
5	No	Single	75K	0.9
6	No	Married	160K	0.4
7	No	Single	50K	0.3

Original Data (X)

$r_0$
0
0.5
0.3
1
0.9
0.4
0.3

Residual  
(Initial)

Fit tree with  
response  $r_0$

$f_1$
0.2
0.3
0.3
1
1
0.4
0.3

$\lambda f_1$
0.1
0.15
0.15
0.5
0.5
0.2
0.15

$r_1$
-0.1
0.35
0.15
0.5
0.4
0.2
0.15

Fit tree with  
response  $r_1$

$f_2$
-0.2
0.6
0.2
0.6
0.4
0.2
0.2

$\lambda f_2$
-0.1
0.3
0.1
0.3
0.2
0.1
0.1

$r_2$
0
0.05
0.05
0.2
0.2
0.1
0.05

...

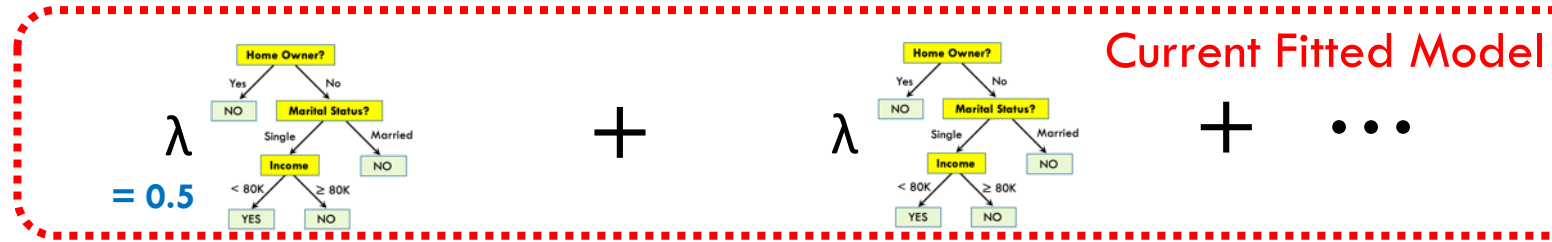
Residual

subtract

Residual

5. Repeat the same process (fitting decision tree, etc.) on the new residuals  $r_1$ .

# GRADIENT BOOSTED DECISION TREES



ID	Home Owner	Marital Status	Annual Income	y
1	Yes	Single	125K	0
2	No	Married	100K	0.5
3	No	Single	170K	0.3
4	Yes	Married	120K	1
5	No	Single	75K	0.9
6	No	Married	160K	0.4
7	No	Single	50K	0.3

Original Data (X)

$r_0$
0
0.5
0.3
1
0.9
0.4
0.3

Residual  
(Initial)

Fit tree with  
response  $r_0$

$f_1$
0.2
0.3
0.3
1
1
0.4
0.3

$\lambda f_1$
0.1
0.15
0.15
0.5
0.5
0.2
0.15

$r_1$
-0.1
0.35
0.15
0.5
0.4
0.2
0.15

Residual

Fit tree with  
response  $r_1$

$f_2$
-0.2
0.6
0.2
0.6
0.4
0.2
0.2

$\lambda f_2$
-0.1
0.3
0.1
0.3
0.2
0.1
0.1

$r_2$
0
0.05
0.05
0.2
0.2
0.1
0.05

Residual

...

6. After fitting B trees, the final prediction is the sum of damped trees.

# SUMMARY: GRADIENT TREES BOOSTING (REGRESSION CASE)

**Initialization:**  $\hat{f}(x) = 0$ , and  $r_i = y_i$  for all  $i$

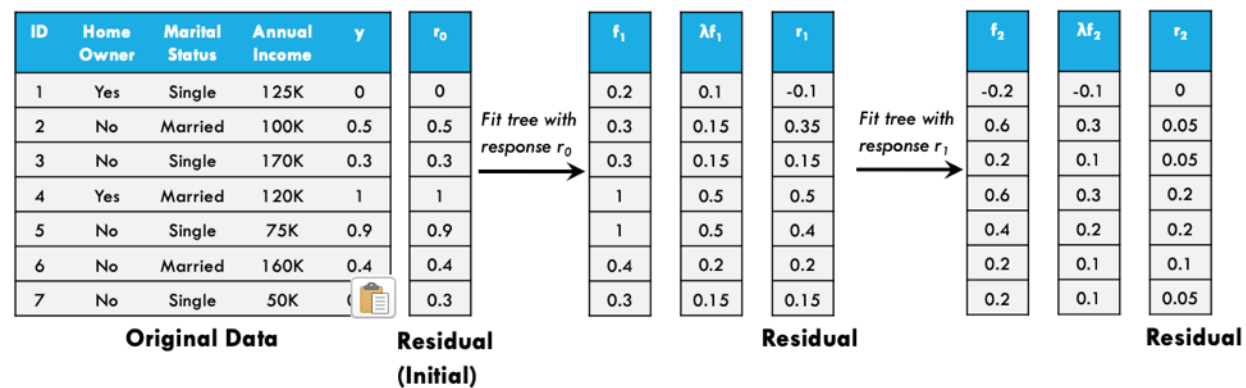
For  $b = 1, \dots, B$ :

- Fit Decision Tree  $\hat{f}^b$  to residuals  $r_1, \dots, r_n$ .
- Update Residual

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i).$$

**Output:**

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x).$$

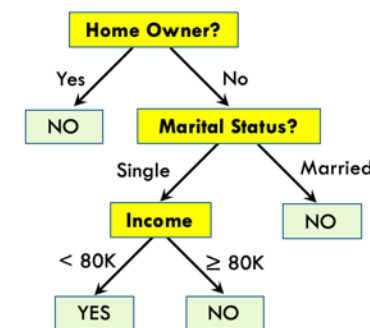


# QUIZ: GRADIENT BOOSTING HYPERPARAMETERS

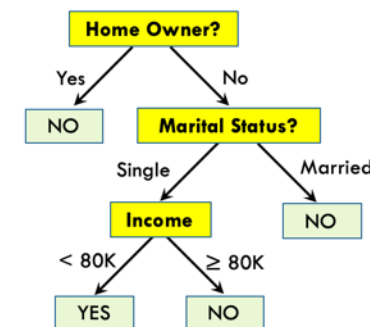


**Q:** While tuning a gradient boosting tree model, you decide to decrease the value of  $\lambda$ . How does this affect the number of trees you should use?

1. Increase
2. Decrease



...





# QUIZ: GRADIENT BOOSTING HYPERPARAMETERS

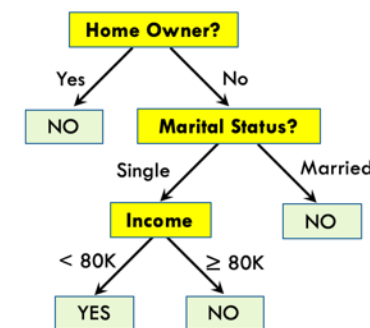


**Q:** While tuning a gradient boosting tree model, you decide to decrease the value of  $\lambda$ . How does this affect the number of trees you should use?

1. Increase

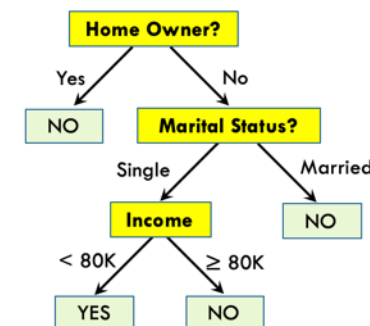


2. Decrease



**Model 1**

...



**Model N**

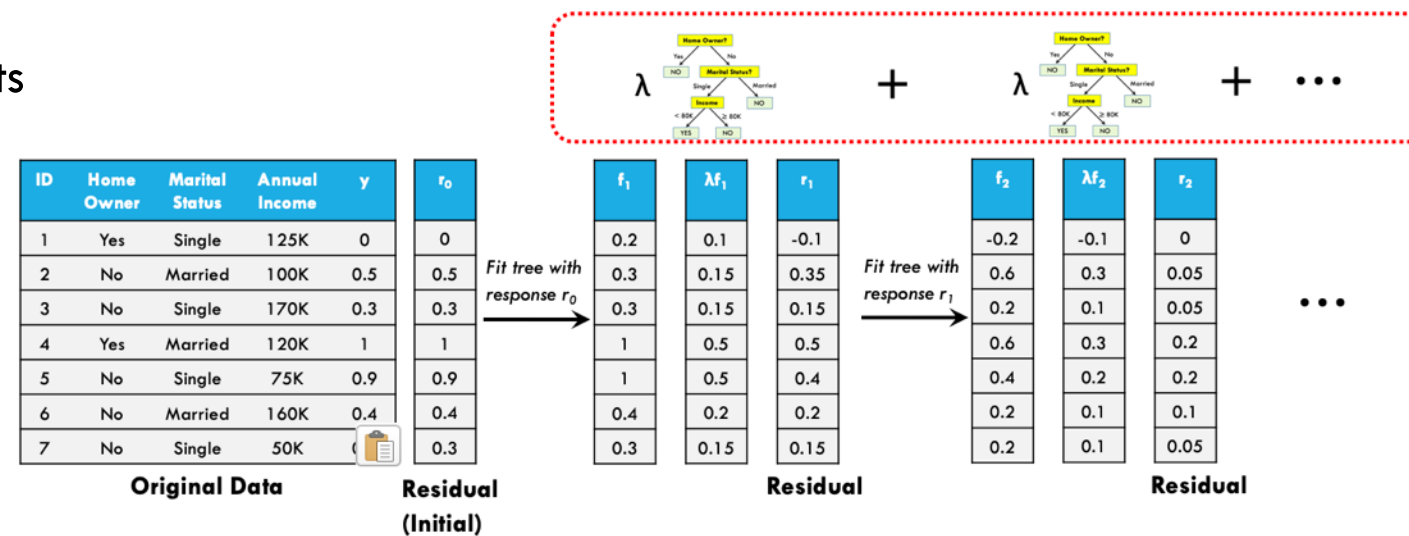
# PROS AND CONS

## Pros

- *Variance Reduction*: Ensembling many decision trees leads to more accurate predictions
- Generally, more accurate than random forests

## Cons

- Less interpretable than decision trees
- Slower than decision trees
- Significant tuning required<sup>1</sup> (for tree depth, learning rate, number of trees)



<sup>1</sup>**General tuning guidelines:** tree depth is usually between 1 to 10, and it is important for controlling overfitting. Tuning it by selecting from {2, 4, 6, 8, 10} based on cross validation is generally fine. Learning rate and no. of trees have to be tuned together: learning rate is usually around 0.01 to 0.1. Smaller learning rate requires more trees (and thus longer training time). If speed is not an issue, a common approach is to set a low learning rate (e.g. 0.01), and train it while selecting the number of trees using early stopping (most gradient boosting tree packages have early stopping already implemented). If this is too slow, it may be advisable to start with a higher learning rate (e.g. 0.1) and later assess whether lowering the learning rate provides improvement in accuracy. See <https://machinelearningmastery.com/configure-gradient-boosting-algorithm/> (and other guides that exist online) for more details (including other tuning parameters)

# CLASSIFICATION OVERVIEW

1. Problem Setup
2. Evaluating Classifiers
3. Nearest Neighbor Methods
4. Trees and Ensembles
5. Logistic Regression
6. Deep Learning



# CLASSIFICATION OVERVIEW

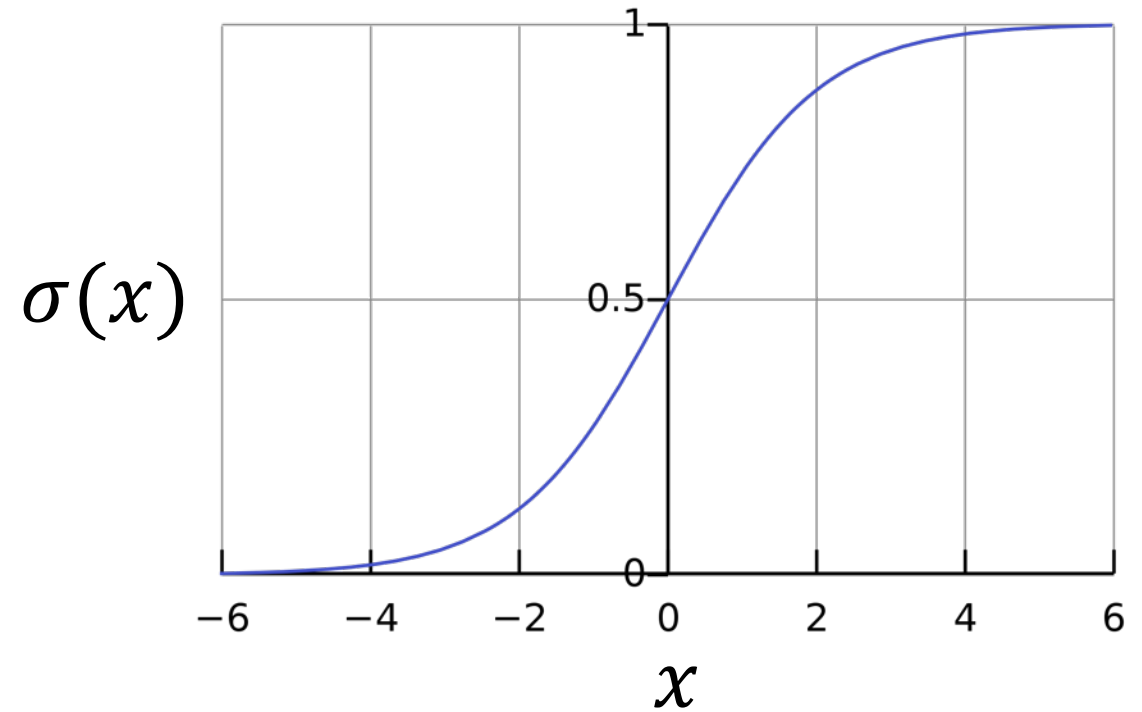
1. Problem Setup
2. Evaluating Classifiers
3. Nearest Neighbor Methods
4. Trees and Ensembles
5. Logistic Regression
6. Deep Learning

# SIGMOID FUNCTION

The sigmoid function  $\sigma(x)$  maps the real numbers to the range  $(0, 1)$

It is defined as:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



# SIGMOID FUNCTION

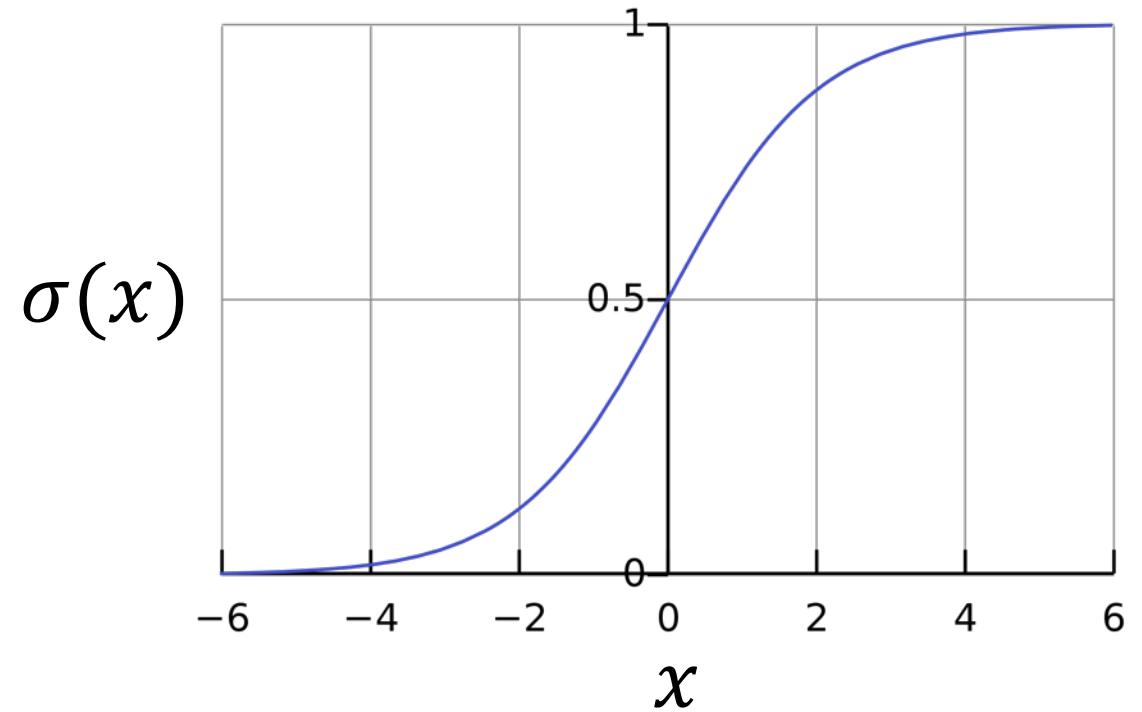
The sigmoid function  $\sigma(x)$  maps the real numbers to the range  $(0, 1)$

It is defined as:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Check:

$$\text{As } x \rightarrow -\infty, \sigma(x) \rightarrow \frac{1}{1+e^{\infty}} = 0$$



# SIGMOID FUNCTION

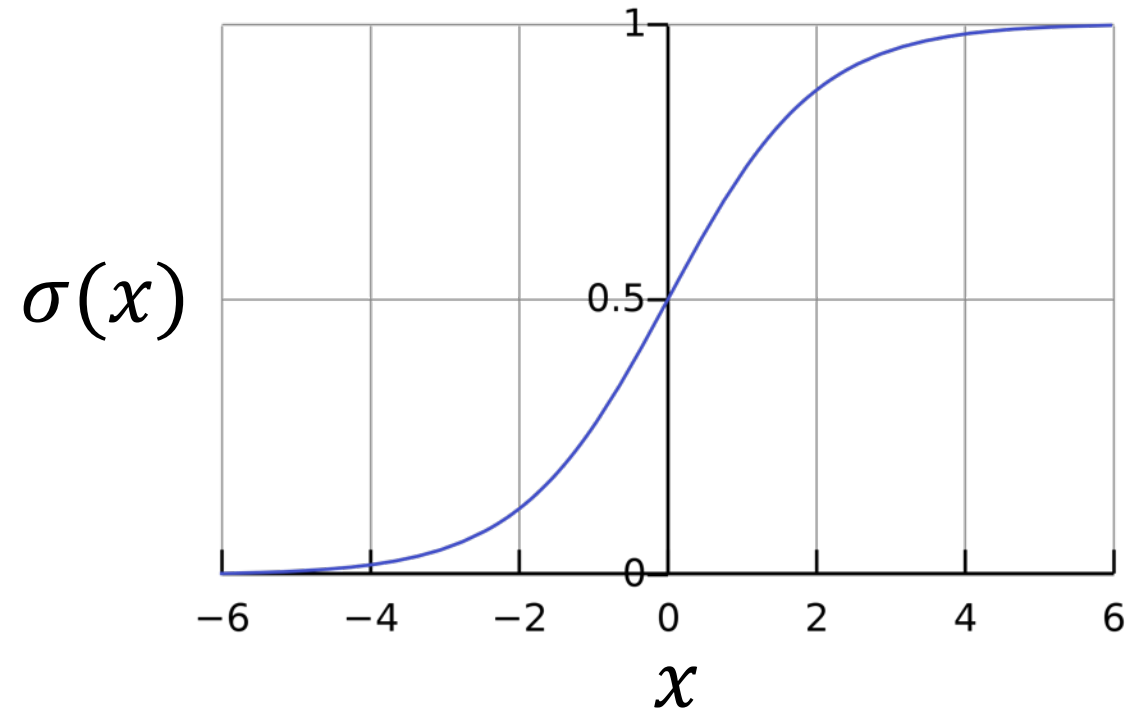
The sigmoid function  $\sigma(x)$  maps the real numbers to the range  $(0, 1)$

It is defined as:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Check:

$$\sigma(0) = \frac{1}{1 + e^0} = 0.5$$



# SIGMOID FUNCTION

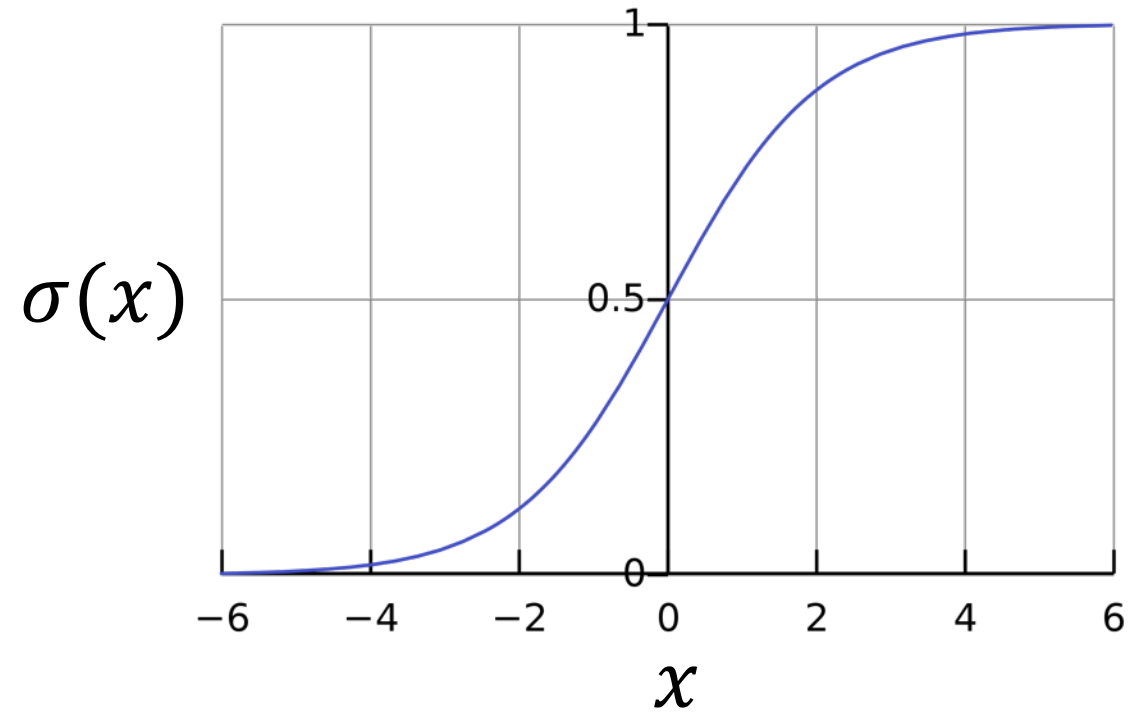
The sigmoid function  $\sigma(x)$  maps the real numbers to the range  $(0, 1)$

It is defined as:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Check:

$$\text{As } x \rightarrow \infty, \sigma(x) \rightarrow \frac{1}{1 + e^{-\infty}} = 1$$





# SIGMOID FUNCTION

Main benefit: can be interpreted as a probability

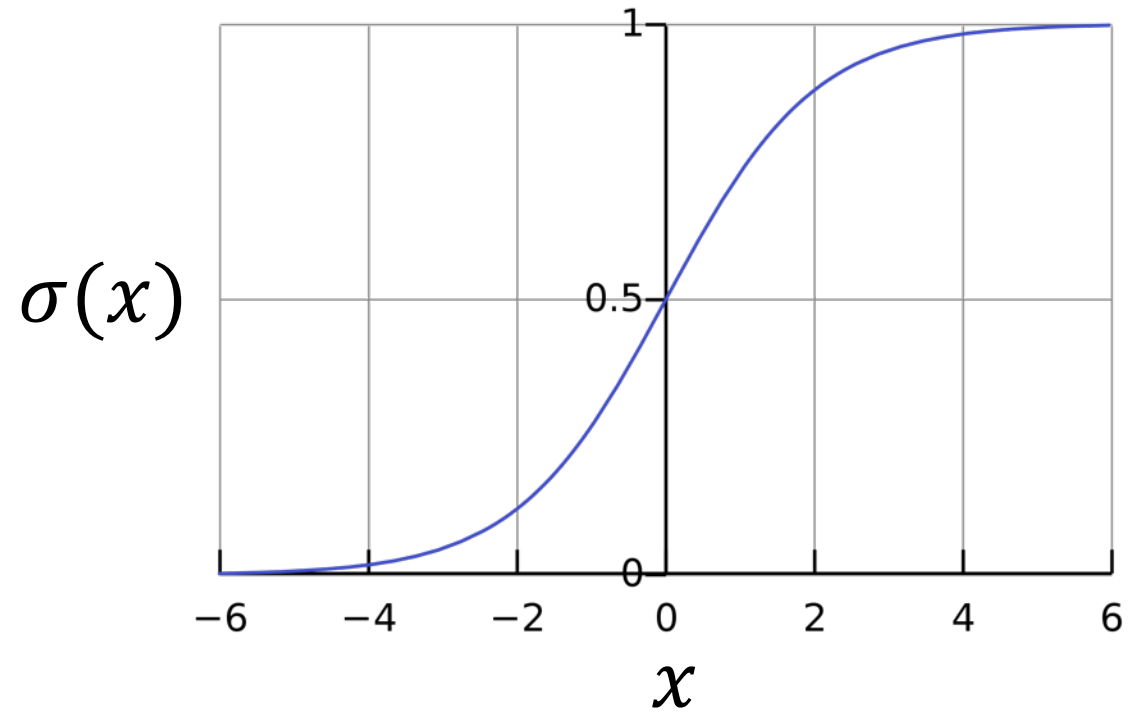
The sigmoid function  $\sigma(x)$  maps the real numbers to the range  $(0, 1)$

It is defined as:

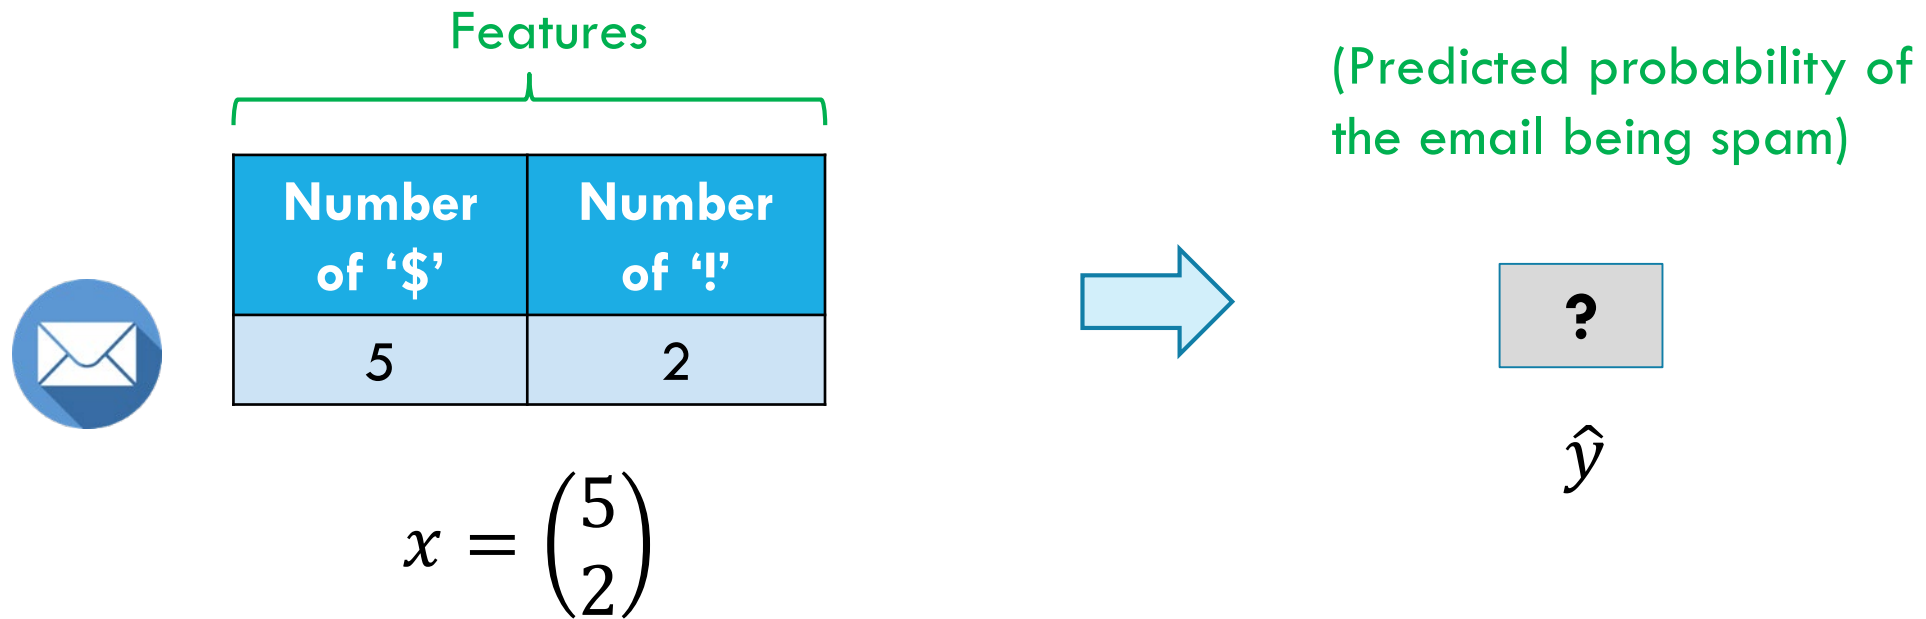
$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Check:

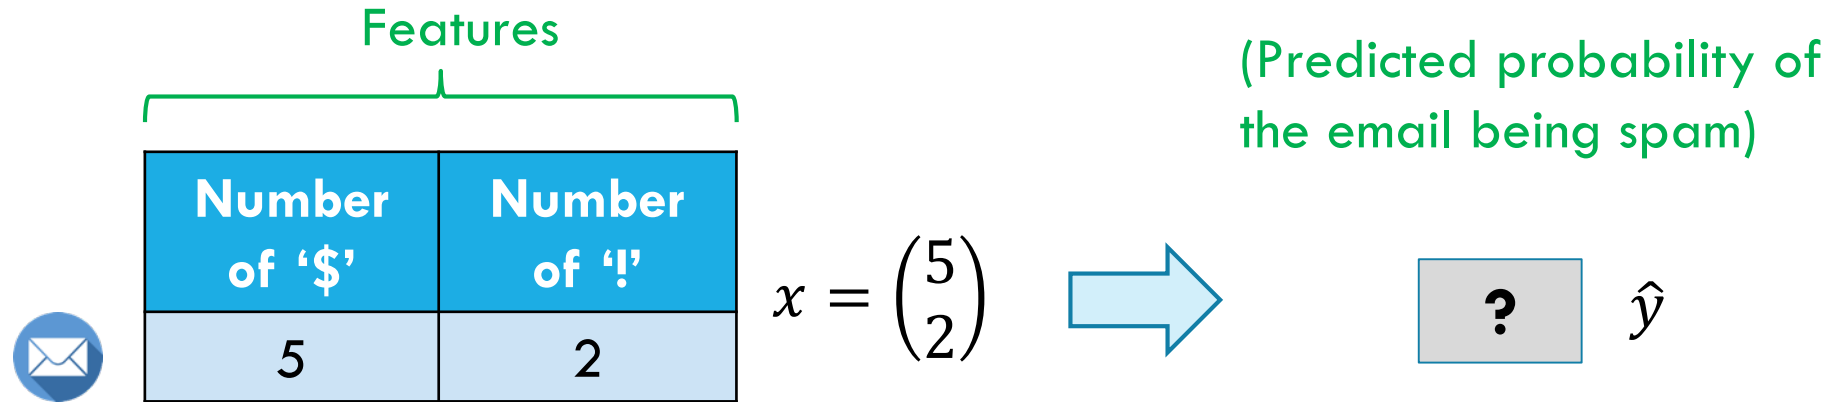
$$\text{As } x \rightarrow \infty, \sigma(x) \rightarrow \frac{1}{1 + e^{-\infty}} = 1$$



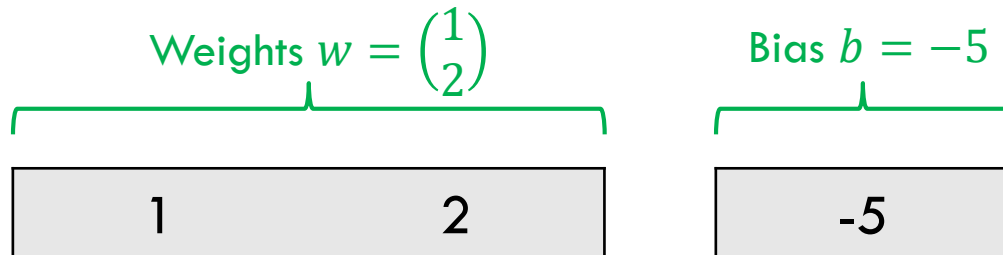
# RUNNING EXAMPLE: SPAM CLASSIFICATION



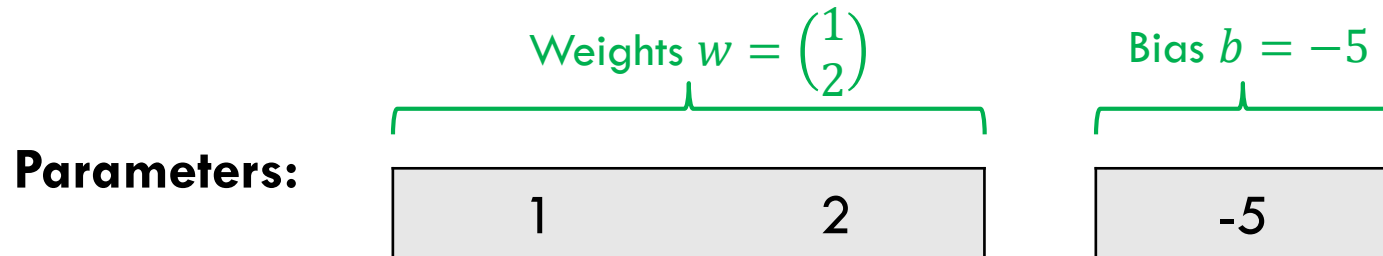
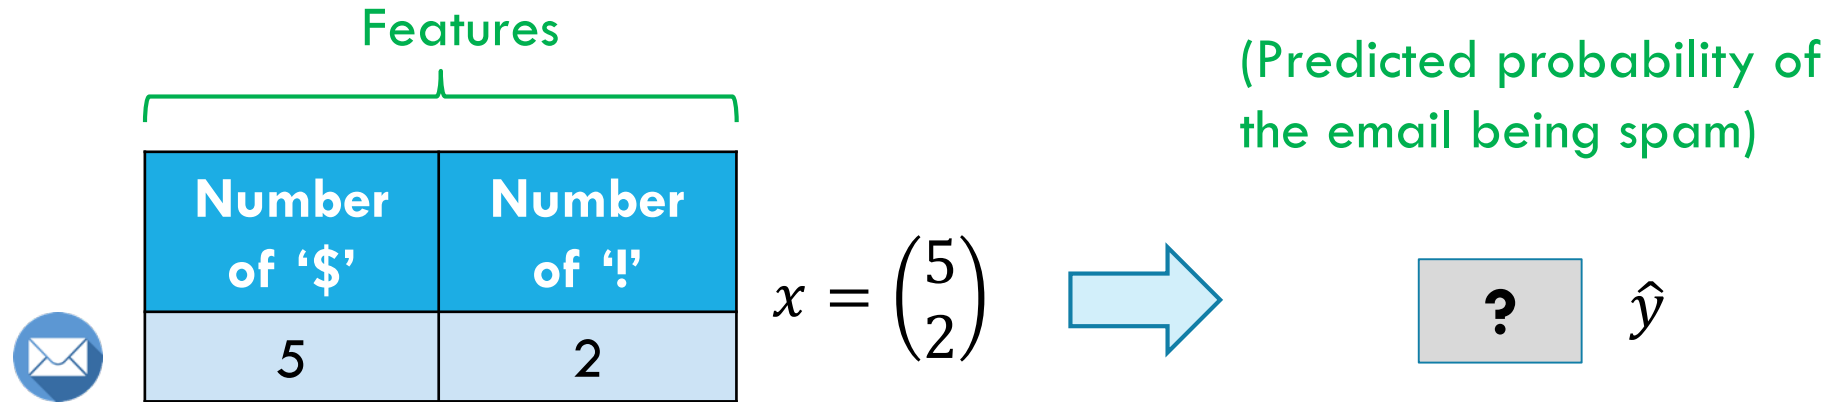
# PARAMETERS OF LOGISTIC REGRESSION



Parameters:



# PREDICTION FUNCTION OF LOGISTIC REGRESSION



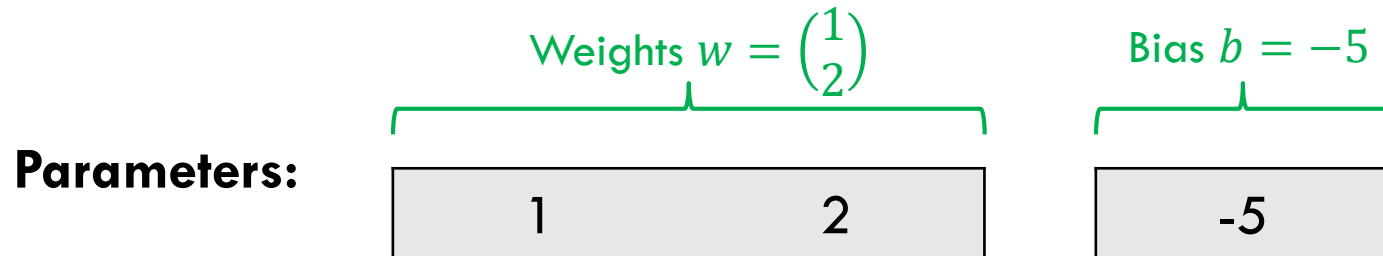
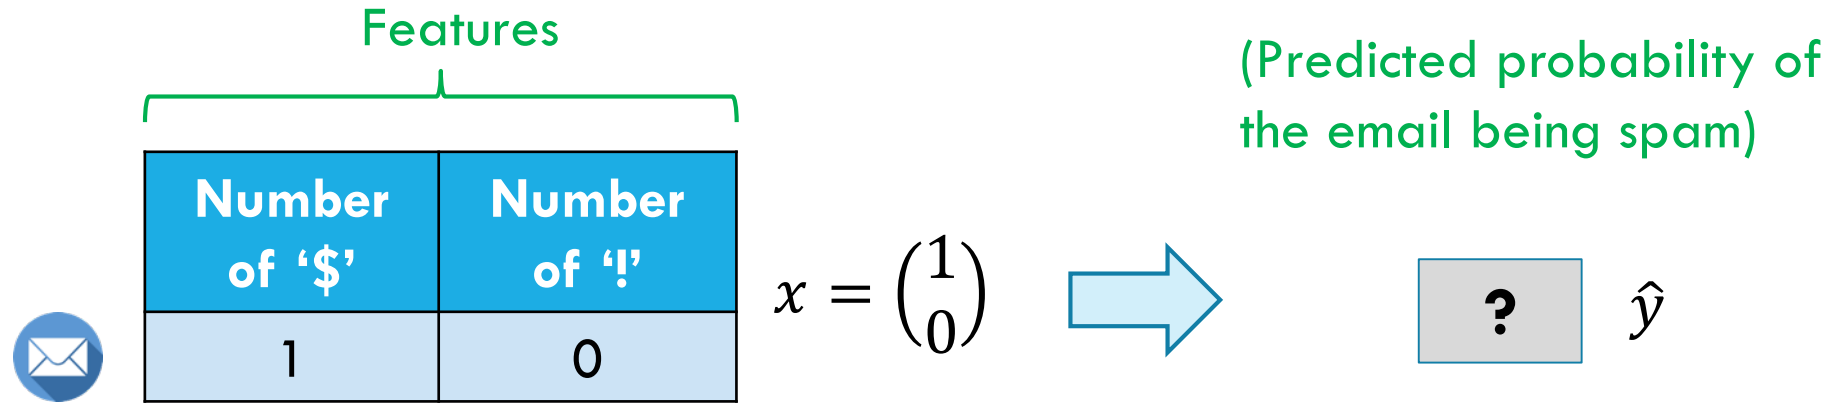
Prediction:

$$\hat{y} = \sigma(x \cdot w + b) = \sigma\left(\begin{pmatrix} 5 \\ 2 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix} - 5\right) = \sigma(9 - 5) = \frac{1}{1 + e^{-4}} = 0.982$$

Sigmoid function      Dot product



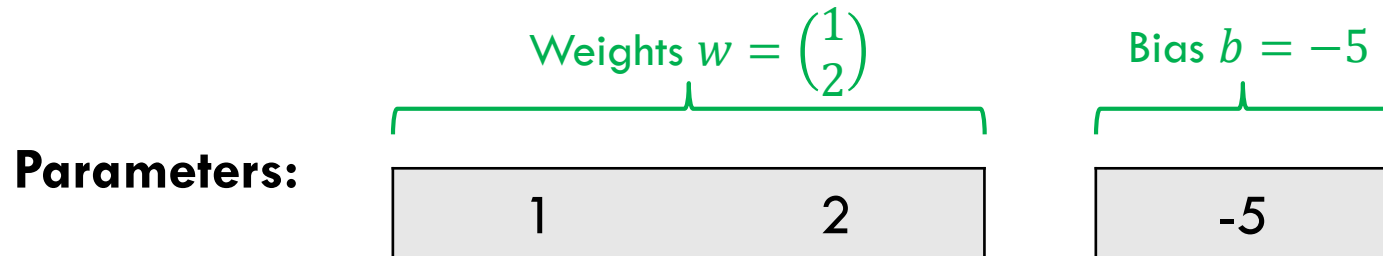
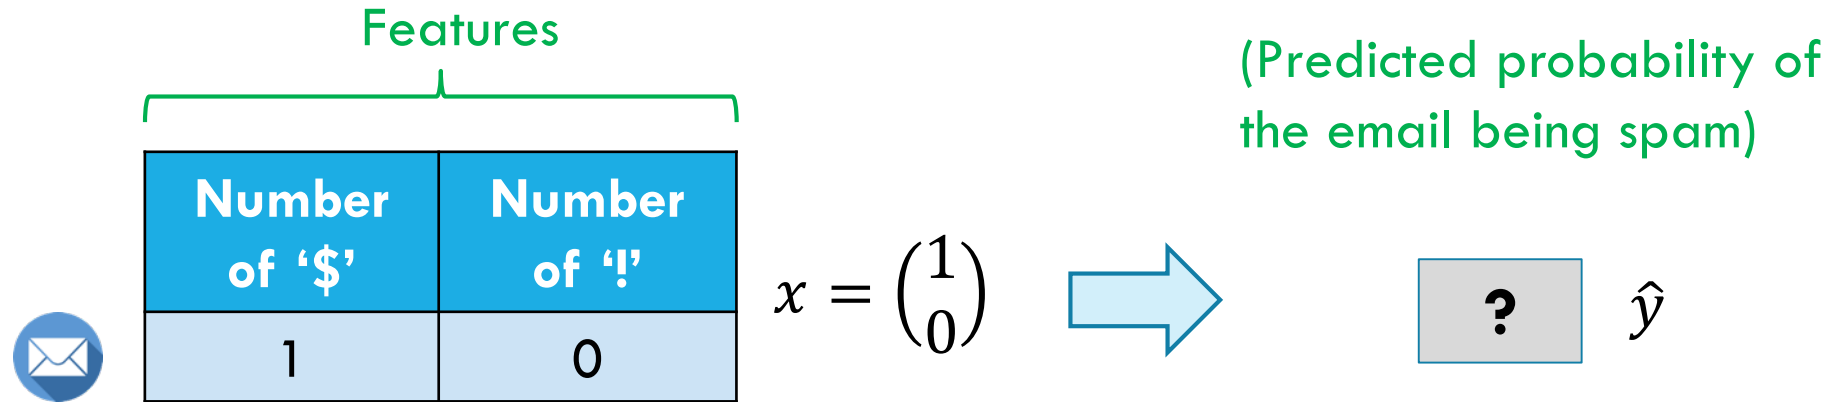
# QUIZ: SPAM CLASSIFICATION



**Prediction:** What is the predicted probability  $\hat{y}$  for this new email  $x = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$ ?



# QUIZ: SPAM CLASSIFICATION



Prediction:

$$\hat{y} = \sigma(x \cdot w + b) = \sigma\left(\begin{pmatrix} 1 \\ 0 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix} - 5\right) = \sigma(1 - 5) = \frac{1}{1 + e^4} = 0.018$$

Sigmoid function      Dot product

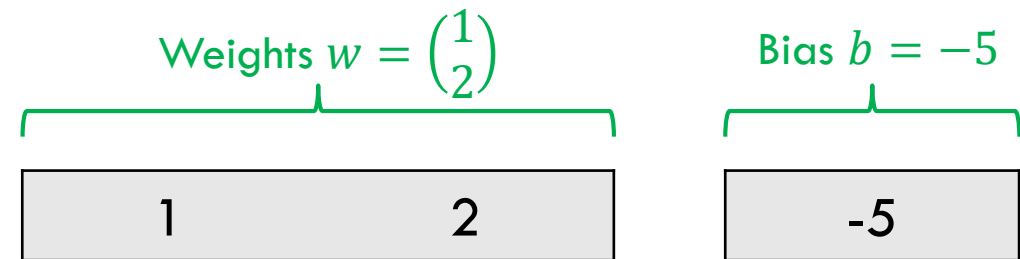
# TRAINING LOGISTIC REGRESSION

## Training Data

samples / observations	features		label
	\$	!	spam
	5	2	1
	$X_{\text{train}}$		$Y_{\text{train}}$

Minimize Cost  
Function  $J(w, b)$

## Logistic Regression Parameters



# COST FUNCTION J FOR PARAMETERS

**Cost function**  $J(w, b)$  for parameters  $w, b$  is defined as the **Cross Entropy Loss** of the predictions  $\hat{y}_i$  obtained from  $w, b$ :

$$J(w, b) = L(\hat{\mathbf{y}}, \mathbf{y}) = \sum_{i=1}^n -y_i \log \hat{y}_i - (1 - y_i) \log(1 - \hat{y}_i)$$

## Interpretation:

Cross Entropy Loss  $L(\hat{\mathbf{y}}, \mathbf{y})$  represents the “**disagreement**” between our predictions  $\hat{\mathbf{y}}$  and the labels  $\mathbf{y}$ .

Now, we are trying to find the “ideal” values of  $w$  and  $b$  that **minimize** this disagreement (note that  $\hat{\mathbf{y}}$  is a function of  $w$  and  $b$ ).



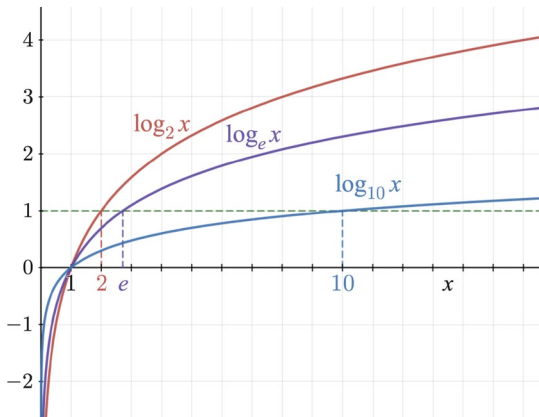
# CROSS-ENTROPY LOSS (FOR A SINGLE DATA POINT)

**Cross Entropy Loss** (for a single  $y$  and  $\hat{y}$ ):

$$L(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

Predictions      Labels

$\hat{y}$	$y$
0.2	1



**How to interpret this?**

The log probability of observing  $y$  given  $\hat{y}$  is

$$\log P(y|\hat{y}) = \begin{cases} \log \hat{y} & \text{if } y = 1 \\ \log(1 - \hat{y}) & \text{if } y = 0 \end{cases}$$

This is exactly the negative of cross entropy loss.

Thus,  $L(\hat{y}, y)$  can be interpreted as **negative log-likelihood** of observing  $y$  given  $\hat{y}$

# CROSS-ENTROPY LOSS (FOR FULL DATA)

Now the labels and predictions are **vectors**,  $\mathbf{y}$  and  $\hat{\mathbf{y}}$ :

Predictions ↓	Labels ↓
$\hat{y}$	$y$
0.2	1
0.5	0
0.9	1

**Cross Entropy Loss** (for full data  $\mathbf{y}$  and  $\hat{\mathbf{y}}$ ):

$$L(\hat{\mathbf{y}}, \mathbf{y}) = \sum_{i=1}^n -y_i \log \hat{y}_i - (1 - y_i) \log(1 - \hat{y}_i)$$

# COST FUNCTION J FOR PARAMETERS

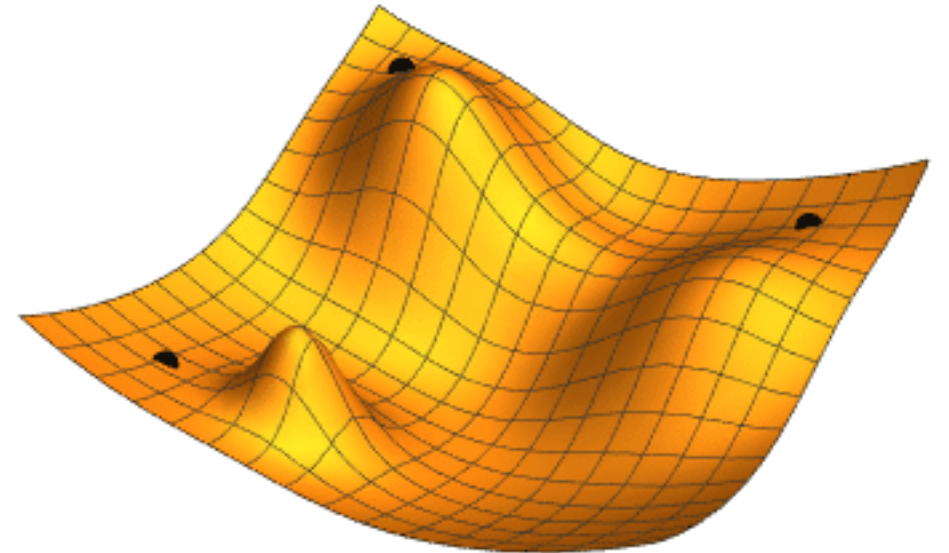
**Cost function**  $J(w, b)$  for parameters  $w, b$  is defined as the Cross Entropy Loss of the predictions  $\hat{y}_i$  obtained from  $w, b$ :

$$J(w, b) = L(\hat{\mathbf{y}}, \mathbf{y}) = \sum_{i=1}^n -y_i \log \hat{y}_i - (1 - y_i) \log(1 - \hat{y}_i)$$

**Goal:** find  $w, b$  to minimize  $J(w, b)$

**Approach: gradient descent**, an incremental approach that repeatedly makes small changes to  $w, b$  to gradually decrease  $J(w, b)$

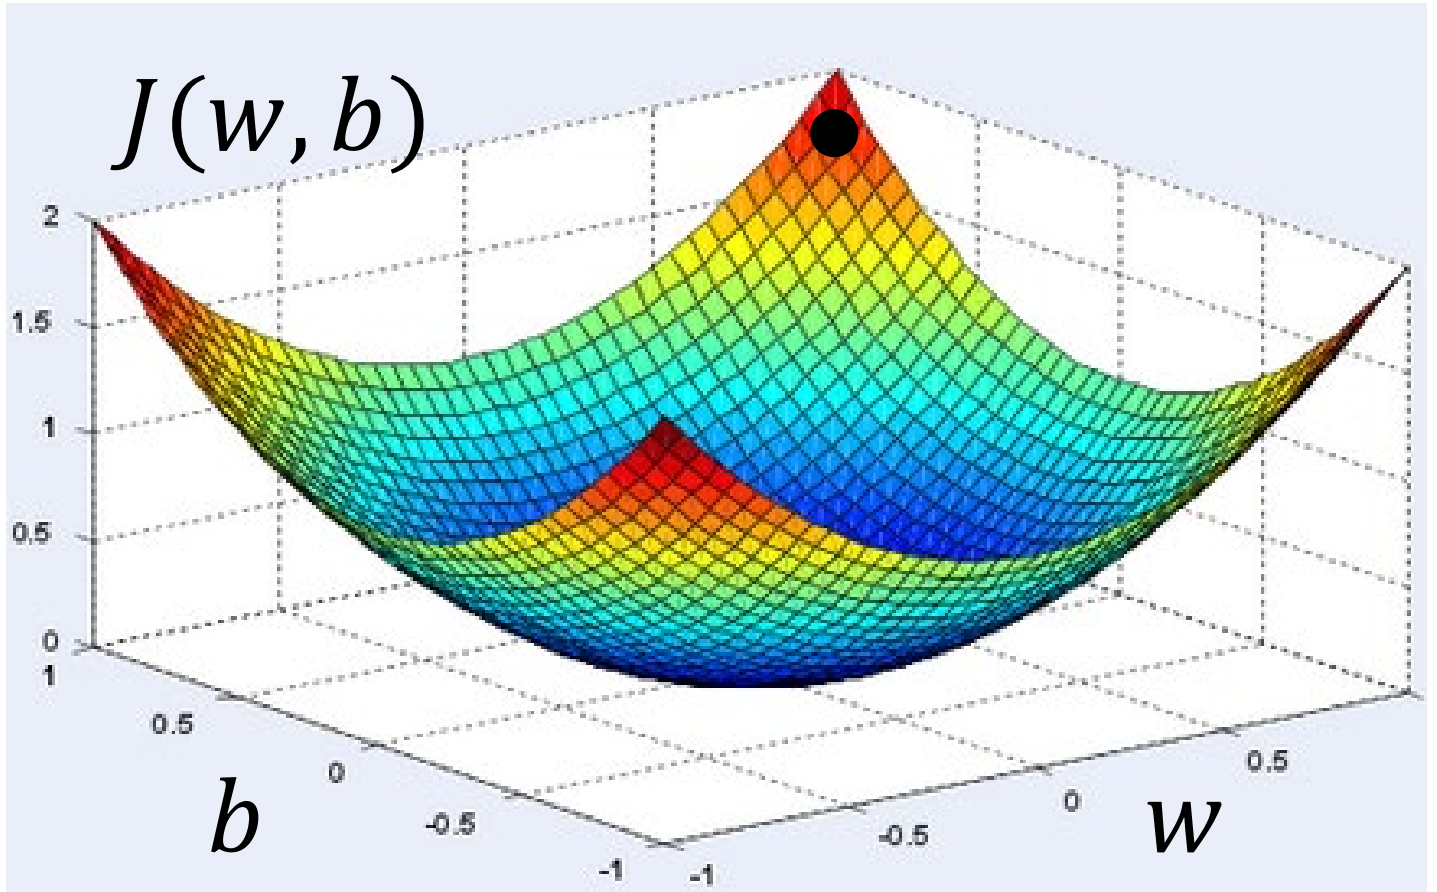
# MINIMIZING COST FUNCTIONS: GRADIENT DESCENT



# MINIMIZING COST FUNCTIONS: GRADIENT DESCENT

We want to find  $w, b$  to minimize  $J(w, b)$

1. Start at an arbitrary point, then keep moving in the negative gradient direction until convergence



# MINIMIZING COST FUNCTIONS: GRADIENT DESCENT

We want to find  $w, b$  to minimize  $J(w, b)$

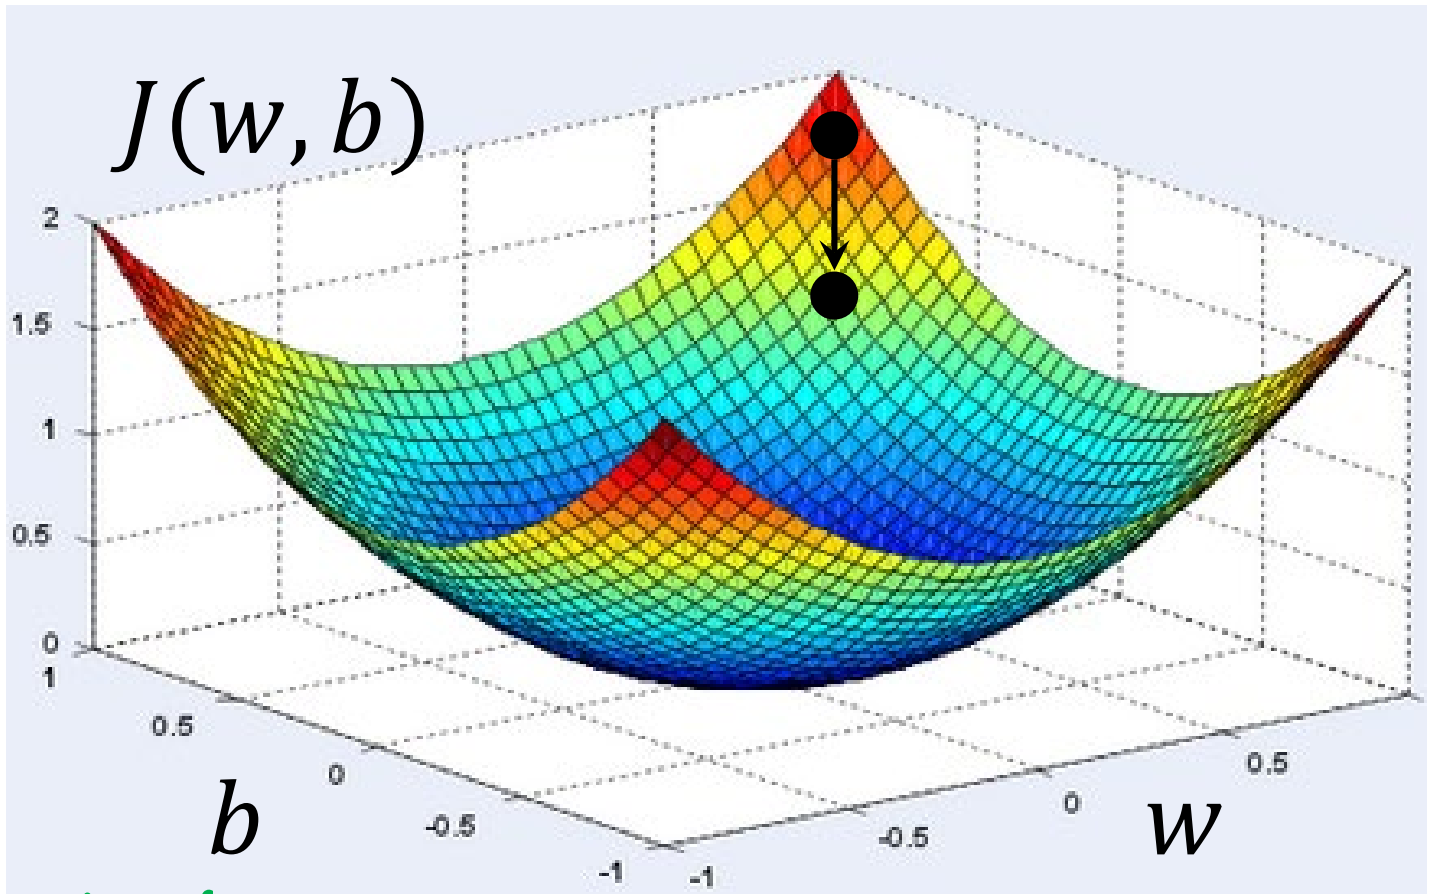
1. Start at an arbitrary point, then move in the negative gradient direction

$$w \leftarrow w - \eta \frac{\partial J(w, b)}{\partial w}$$

$$b \leftarrow b - \eta \frac{\partial J(w, b)}{\partial b}$$

“Learning Rate”

“Slope”: direction of steepest decrease of  $J$



# MINIMIZING COST FUNCTIONS: GRADIENT DESCENT

We want to find  $w, b$  to minimize  $J(w, b)$

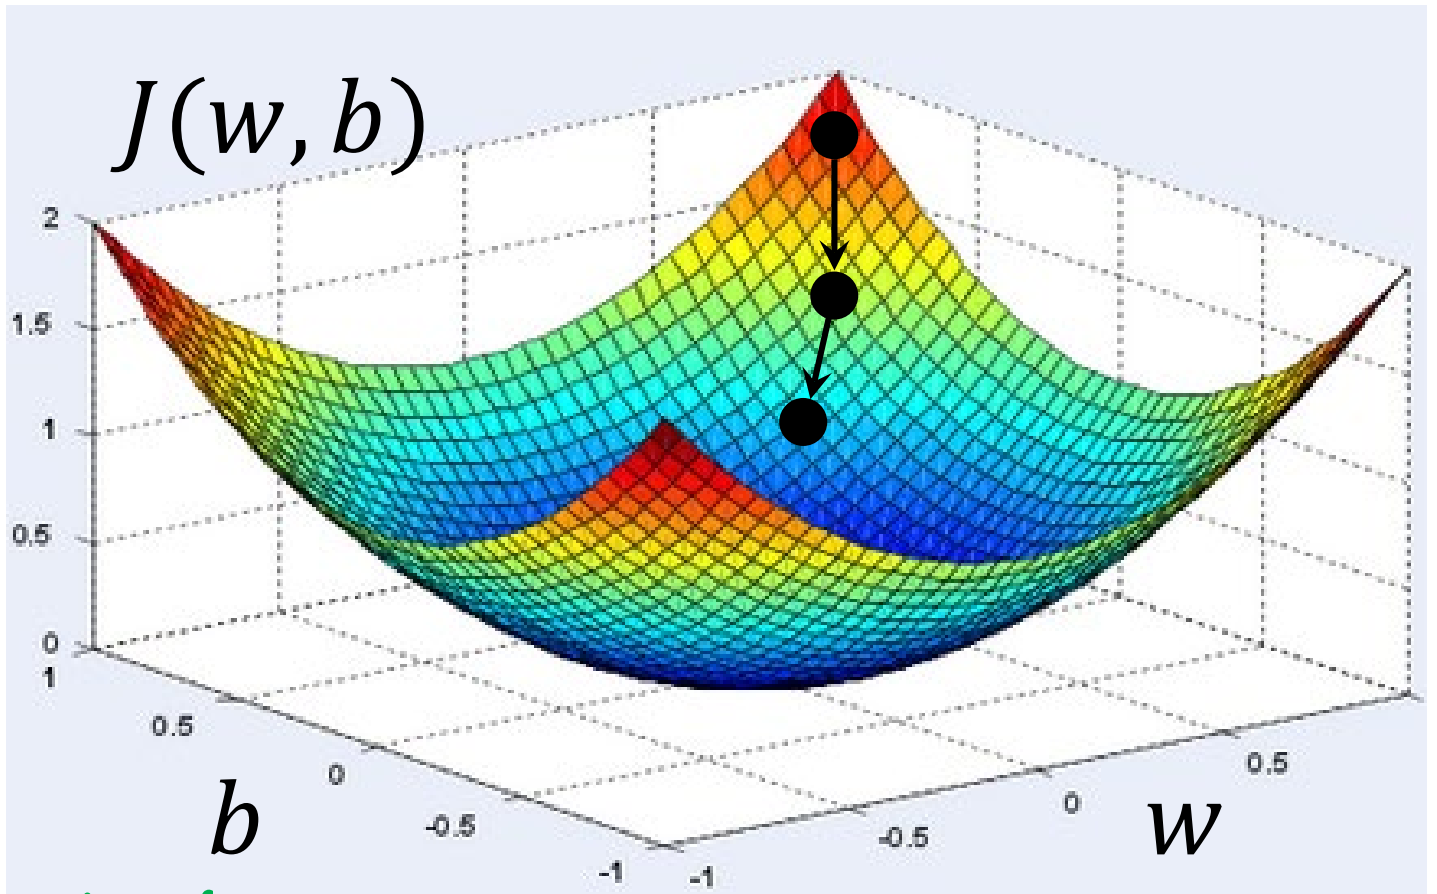
1. Start at an arbitrary point, then move in the negative gradient direction

$$w \leftarrow w - \eta \frac{\partial J(w, b)}{\partial w}$$

$$b \leftarrow b - \eta \frac{\partial J(w, b)}{\partial b}$$

“Learning Rate”

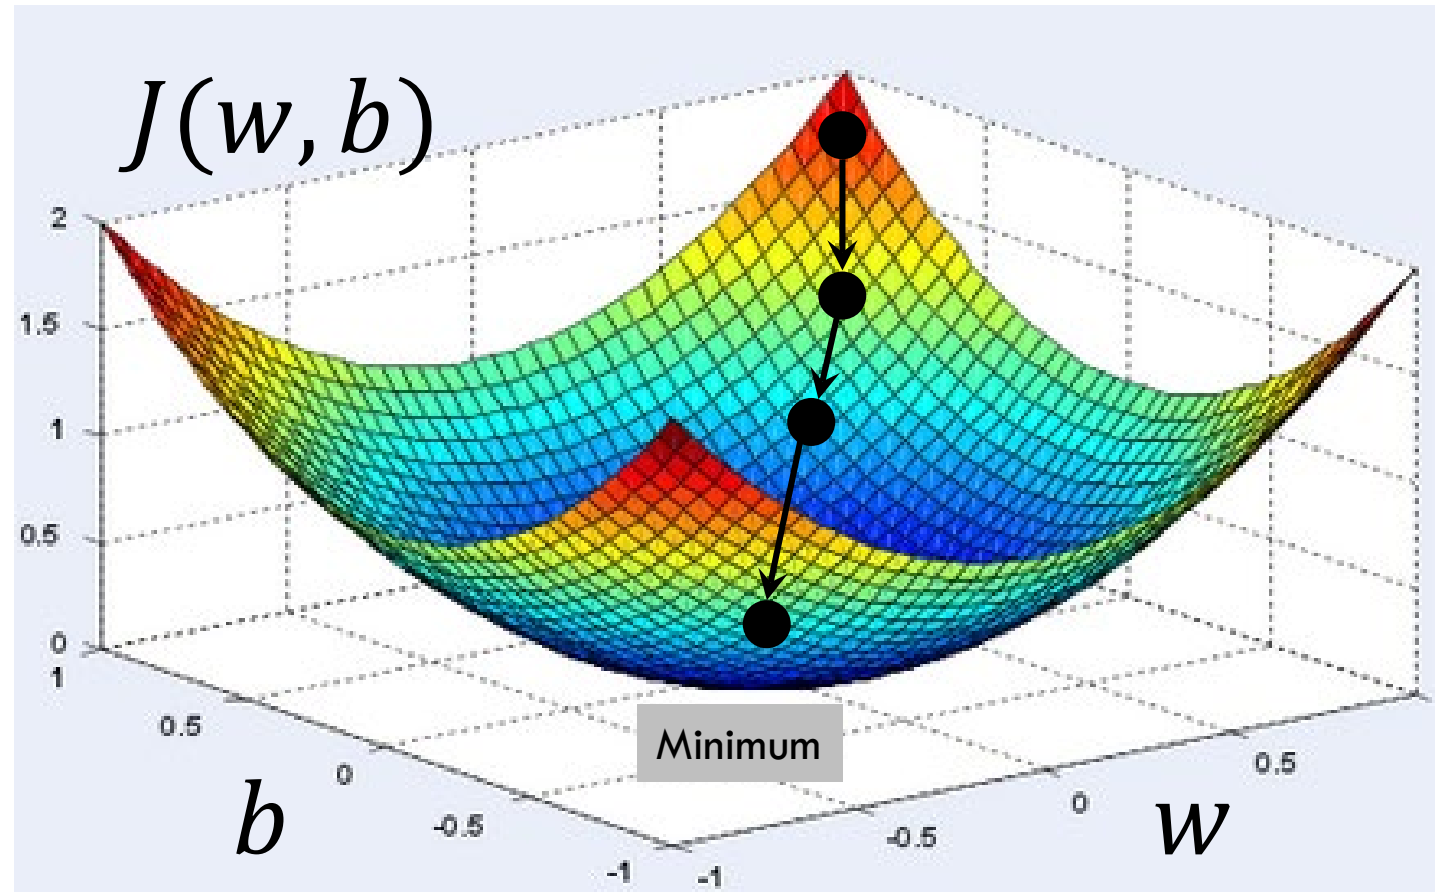
“Slope”: direction of steepest decrease of  $J$



# MINIMIZING COST FUNCTIONS: GRADIENT DESCENT

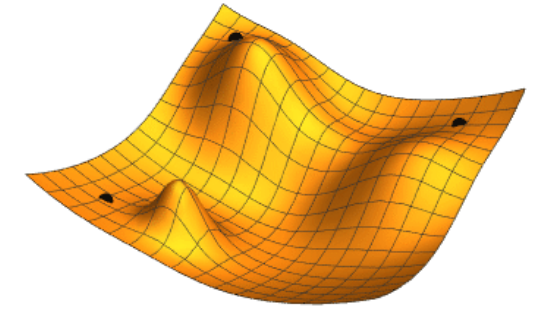
We want to find  $w, b$  to minimize  $J(w, b)$

1. Start at an arbitrary point, then move in the negative gradient direction
2. Continue until convergence
  - Stop when improvement in  $J$  is below a fixed threshold





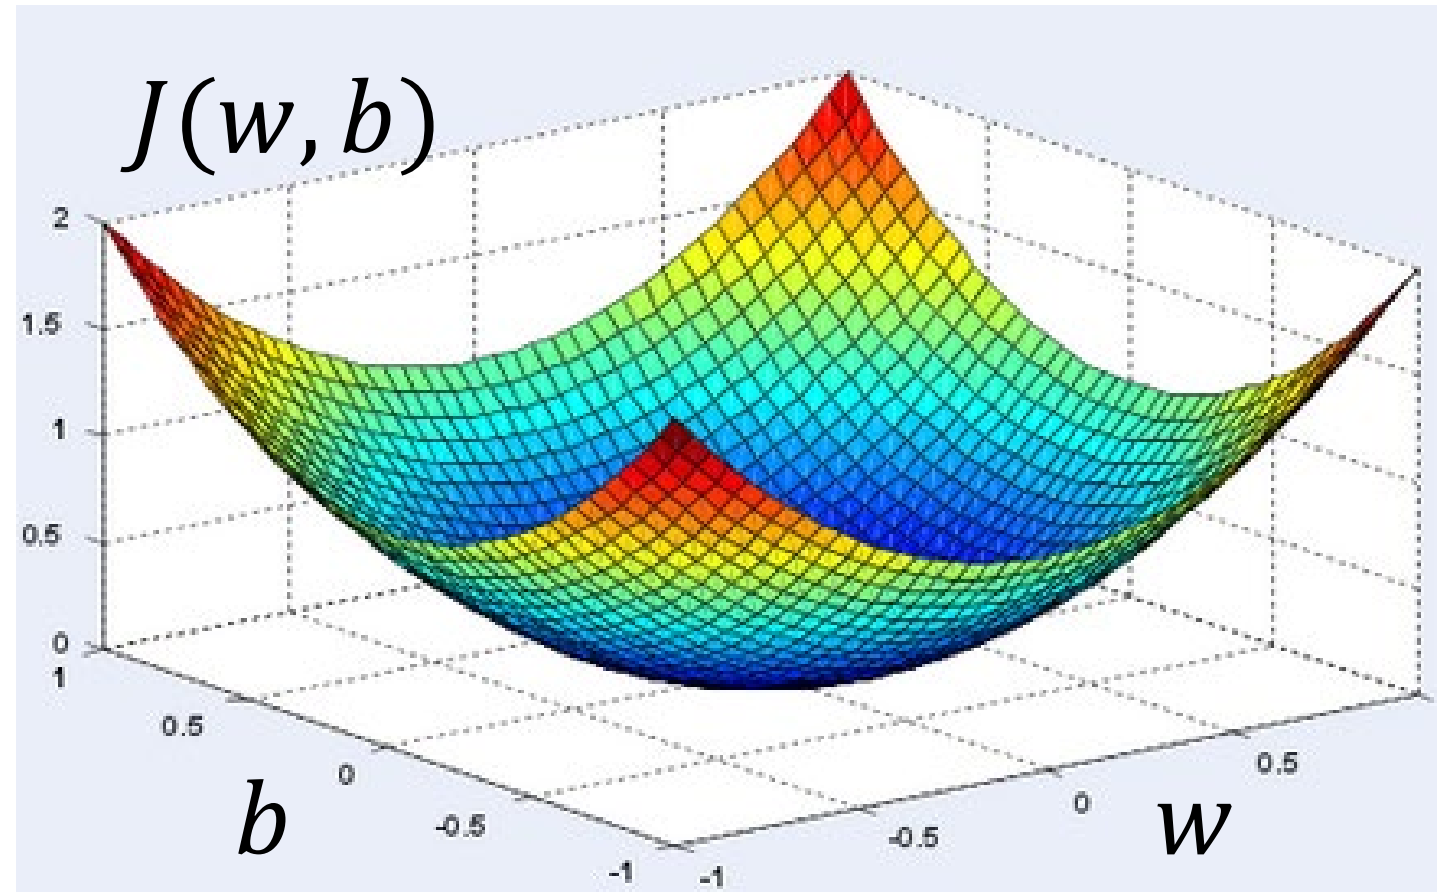
# GLOBAL MINIMUM VS LOCAL MINIMA



This does not always reach the “global minimum”;

It may get stuck in a “local minimum”

But for logistic regression, it always reaches the global minimum (due to “**convexity**” of  $J$ )



# GRADIENT DESCENT: GENERAL ALGORITHM

Important concept: learning rate  $\eta$

- Scaling factor for gradient (typical range: 0.01 - 0.0001)

$$\nabla_{\theta} L = \frac{\partial L}{\partial \theta} = \begin{bmatrix} \frac{\partial L}{\partial \theta_0} \\ \frac{\partial L}{\partial \theta_1} \\ \frac{\partial L}{\partial \theta_2} \\ \vdots \\ \frac{\partial L}{\partial \theta_d} \end{bmatrix}$$

**Input** : data  $(X, y)$ , loss function  $L$ , learning rate  $\eta$

**Initialization** : Set  $\theta$  to random values

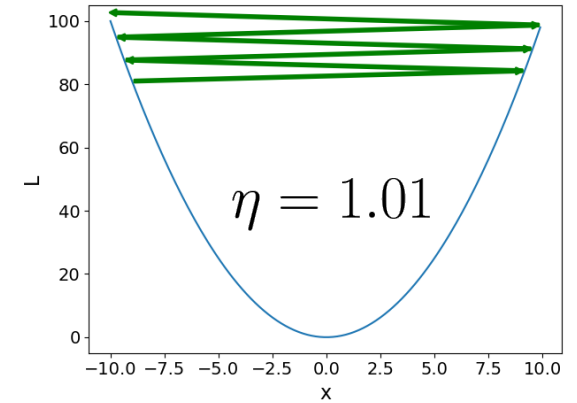
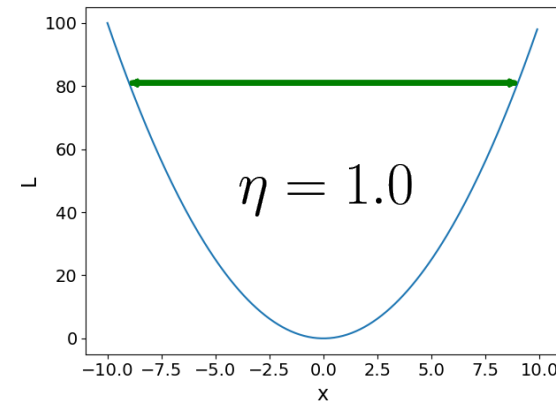
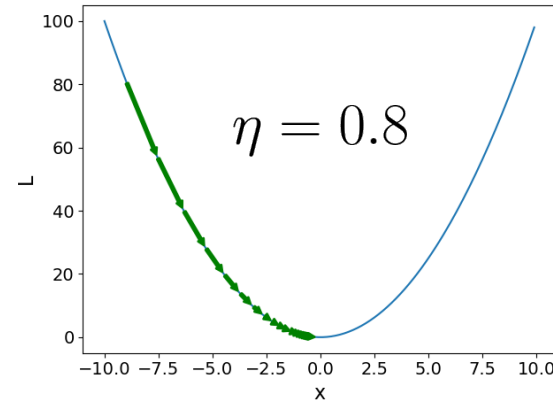
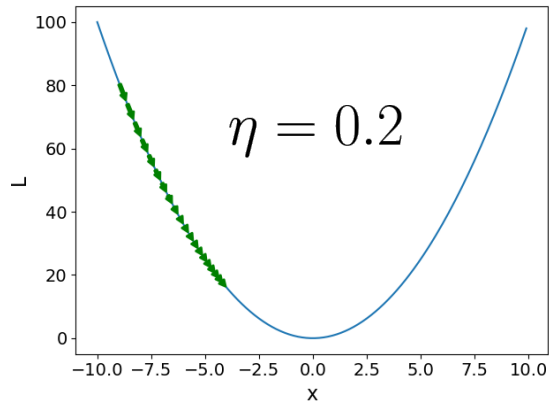
**while true** :

    Calculate gradient  $\nabla_{\theta} L$

$\theta \leftarrow \theta - (\eta \cdot \nabla_{\theta} L)$

In practice: stop loop  
when loss converges

# EFFECTS OF LEARNING RATE



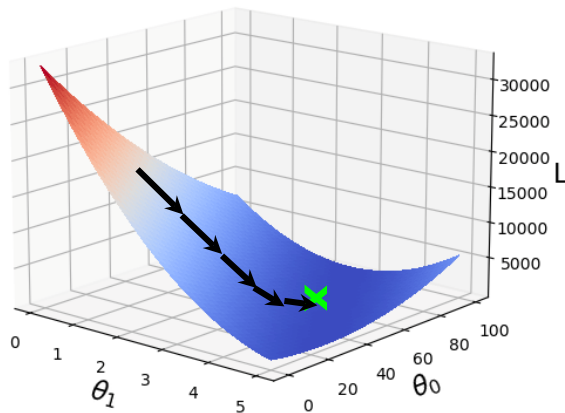
Too low learning rate: slow progress  
Too high learning rate: unstable progress

$$L = x^2, \frac{\partial L}{\partial x} = 2x, 20 \text{ steps}$$

# GRADIENT DESCENT: VARIATIONS

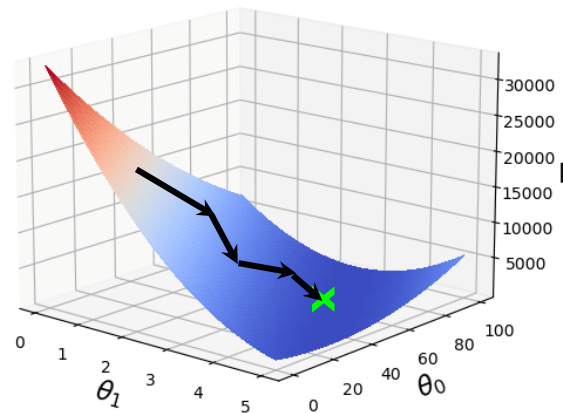
## Full-Batch Gradient Descent

- calculate gradient and update  $\theta$  over **whole** training dataset



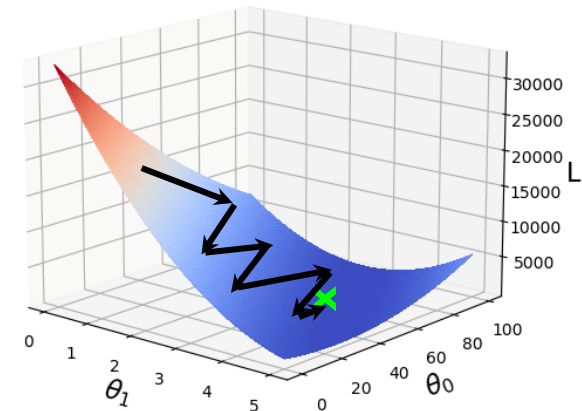
## Mini-batch Gradient Descent

- calculate gradient and update  $\theta$  over **randomly sampled batch** of samples



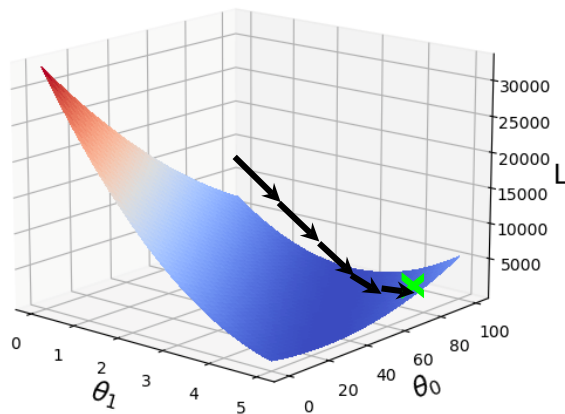
## Stochastic Gradient Descent (SGD)

- calculate gradient and update  $\theta$  over **single** training sample



# GRADIENT DESCENT: VARIATIONS

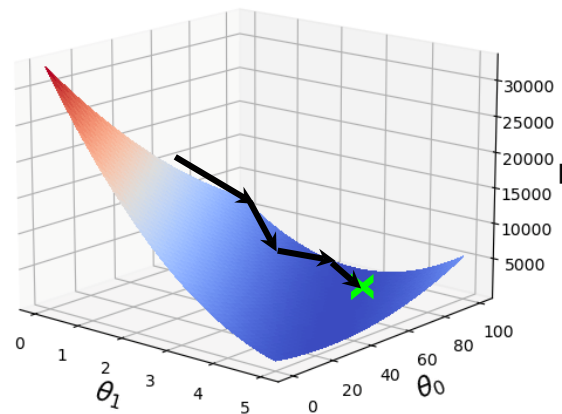
**Full-Batch Gradient Descent**



Gradient averaged over all data items

- Smooth descent
- Small(er) gradients
- Small(er) update steps

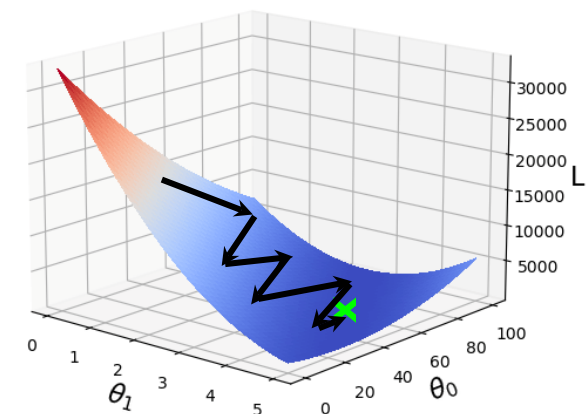
**Mini-Batch Gradient Descent**



Gradient averaged over some data items

Trade-off for gradients  
and update steps

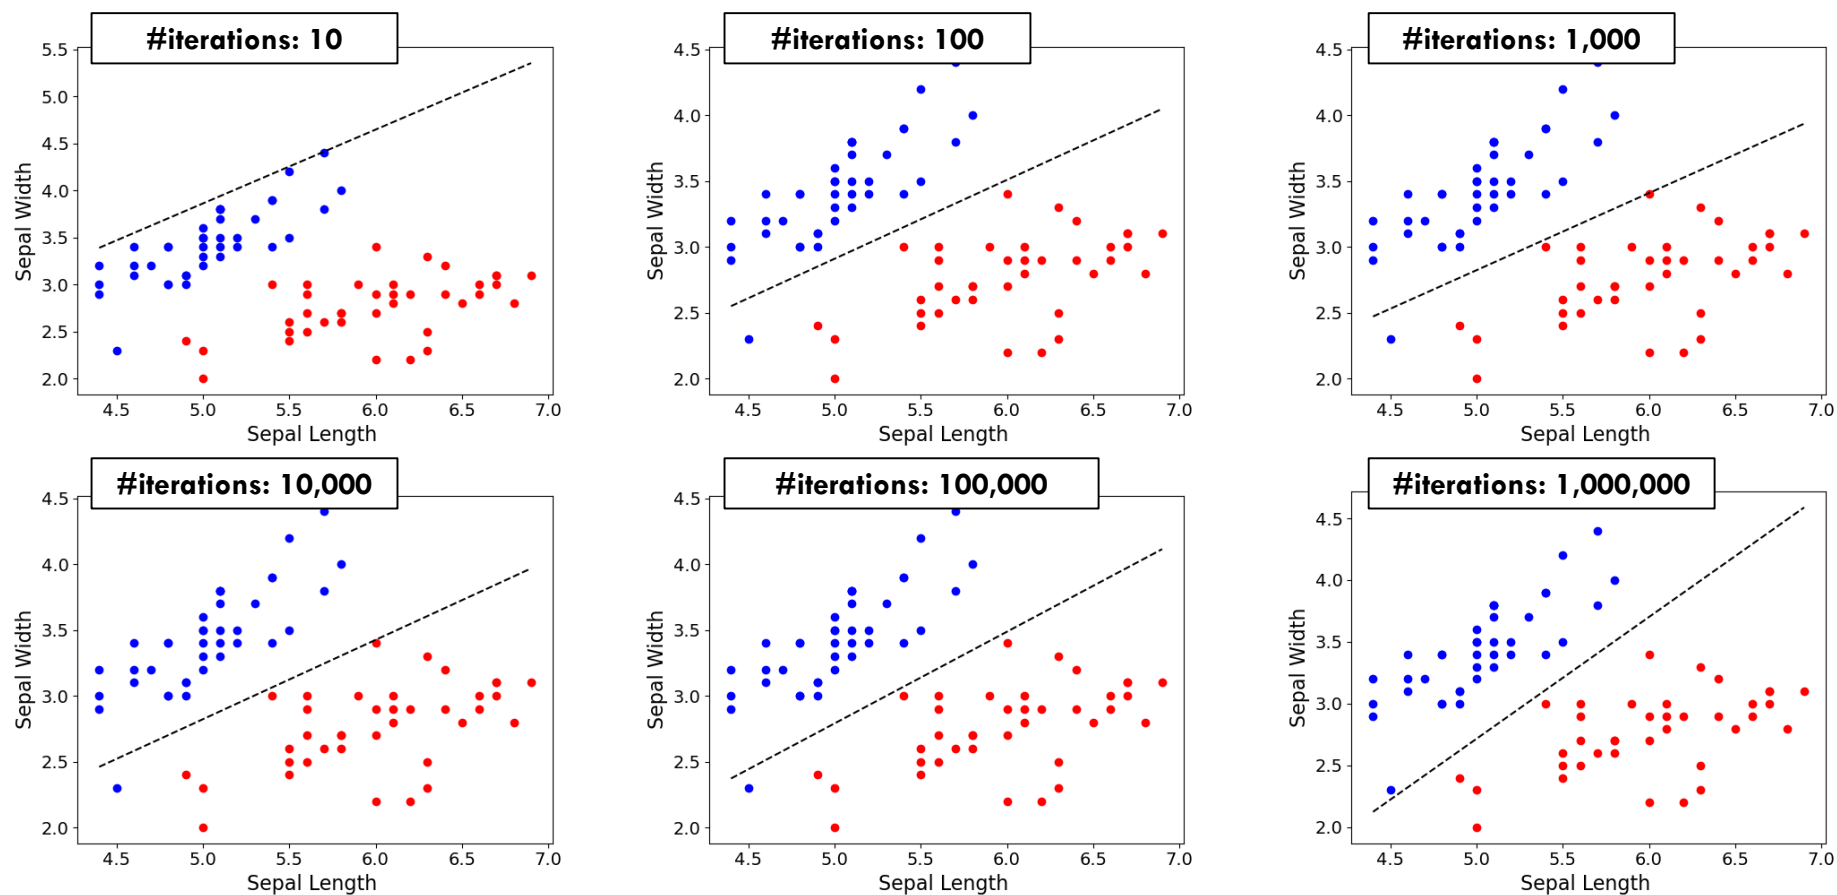
**Stochastic Gradient Descent**



Gradient for each sample

- Choppy descent
- Large(r) gradients
- Large(r) steps

# LOGISTIC REGRESSION: 2D EXAMPLE



(Full-Batch Gradient Descent)

# L1 AND L2 REGULARIZATION

L1 and L2 regularization are commonly used to **control overfitting** in logistic regression.

Given a **regularization parameter**  $\lambda$ , we modify the cost function  $J(w, b)$  to:

- $J(w, b) + \lambda \|w\|_1 = J(w, b) + \lambda \sum_{i=1}^p |w_i|$  (L1 regularization)

- $J(w, b) + \frac{\lambda}{2} \|w\|_2^2 = J(w, b) + \frac{\lambda}{2} \sum_{i=1}^p w_i^2$  (L2 regularization)

We can fit these using gradient descent as before.

- The larger  $\lambda$  is, the stronger the effect of the regularization. However, too large regularization can cause underfitting.
- L1 regularization induces **sparsity**: i.e., generally, many entries of the fitted  $w$  will be exactly 0. However, L2 regularization does **not** induce sparsity.



# PROS AND CONS OF LOGISTIC REGRESSION

## Pros

- Fast and simple
- Loss function is convex
- Interpretable

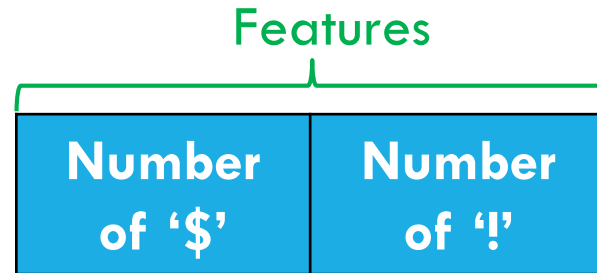
## Cons

- Linear model (up to before sigmoid layer)
- Cannot directly handle categorical features (need one-hot encoding)

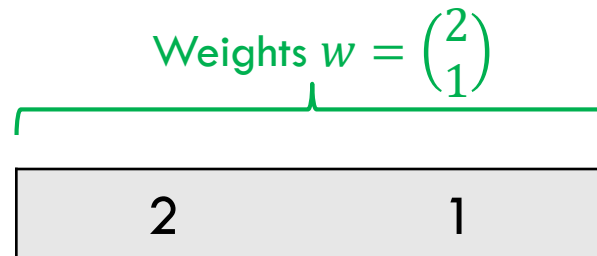




# QUIZ: INTERPRETABILITY OF COEFFICIENTS



**Q:** In the spam classification example, assume we fit the following weight vector:

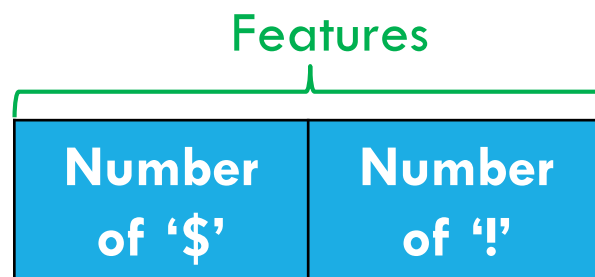


These weights can be interpreted as the “strength” of each feature. Which of the following emails will be given a higher probability of being spam?

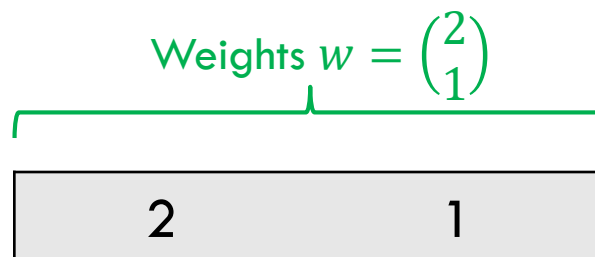
**Email A:** \$\$

**Email B:** !!!

# QUIZ: INTERPRETABILITY OF COEFFICIENTS



**Q:** In the spam classification example, assume we fit the following weight vector:



These weights can be interpreted as the “strength” of each feature. Which of the following emails will be given a higher probability of being spam?

**Email A:** \$\$

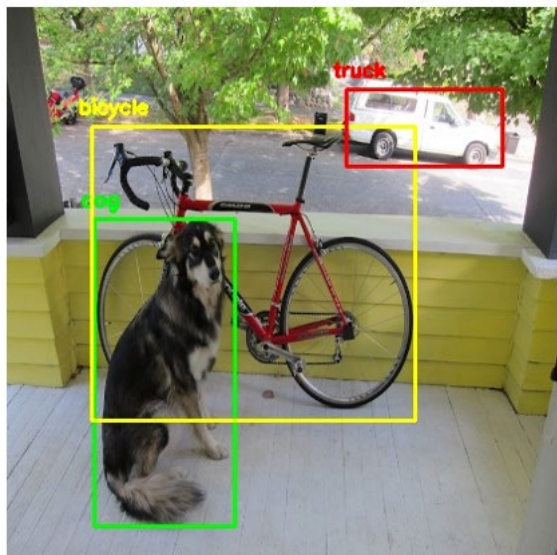
**Email B:** !!!

**A:** Email A

# CLASSIFICATION OVERVIEW

1. Problem Setup
2. Evaluating Classifiers
3. Nearest Neighbor Methods
4. Trees and Ensembles
5. Logistic Regression
6. Deep Learning

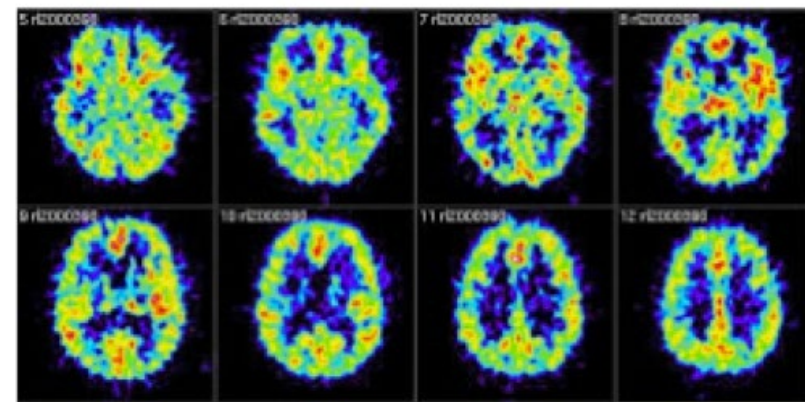
# APPLICATIONS OF DEEP LEARNING



Object Detection



Voice Recognition

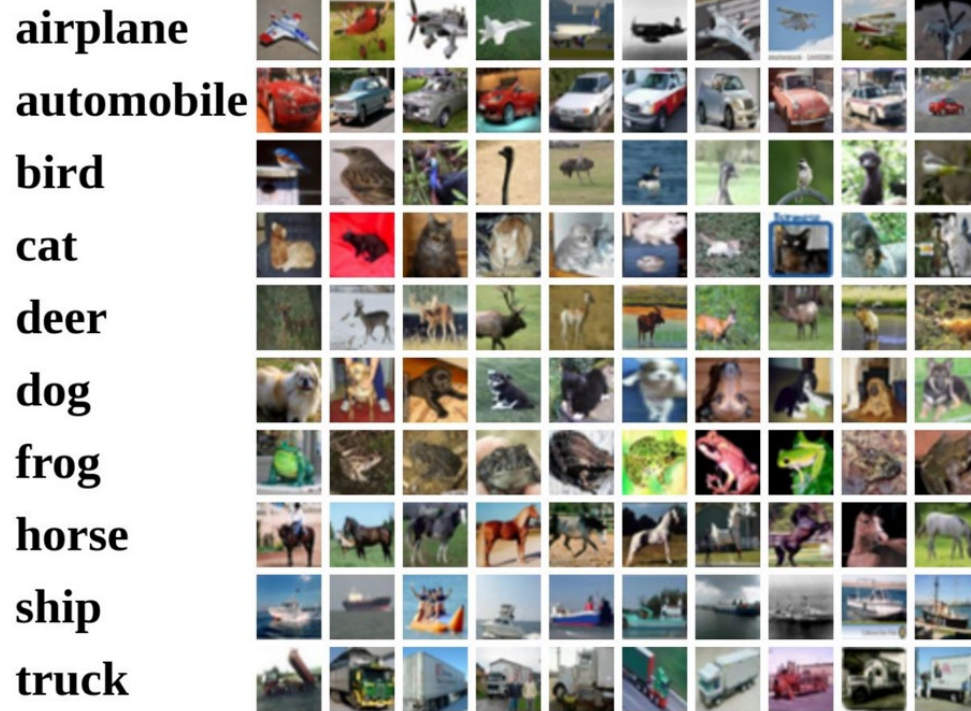


Medical Imaging

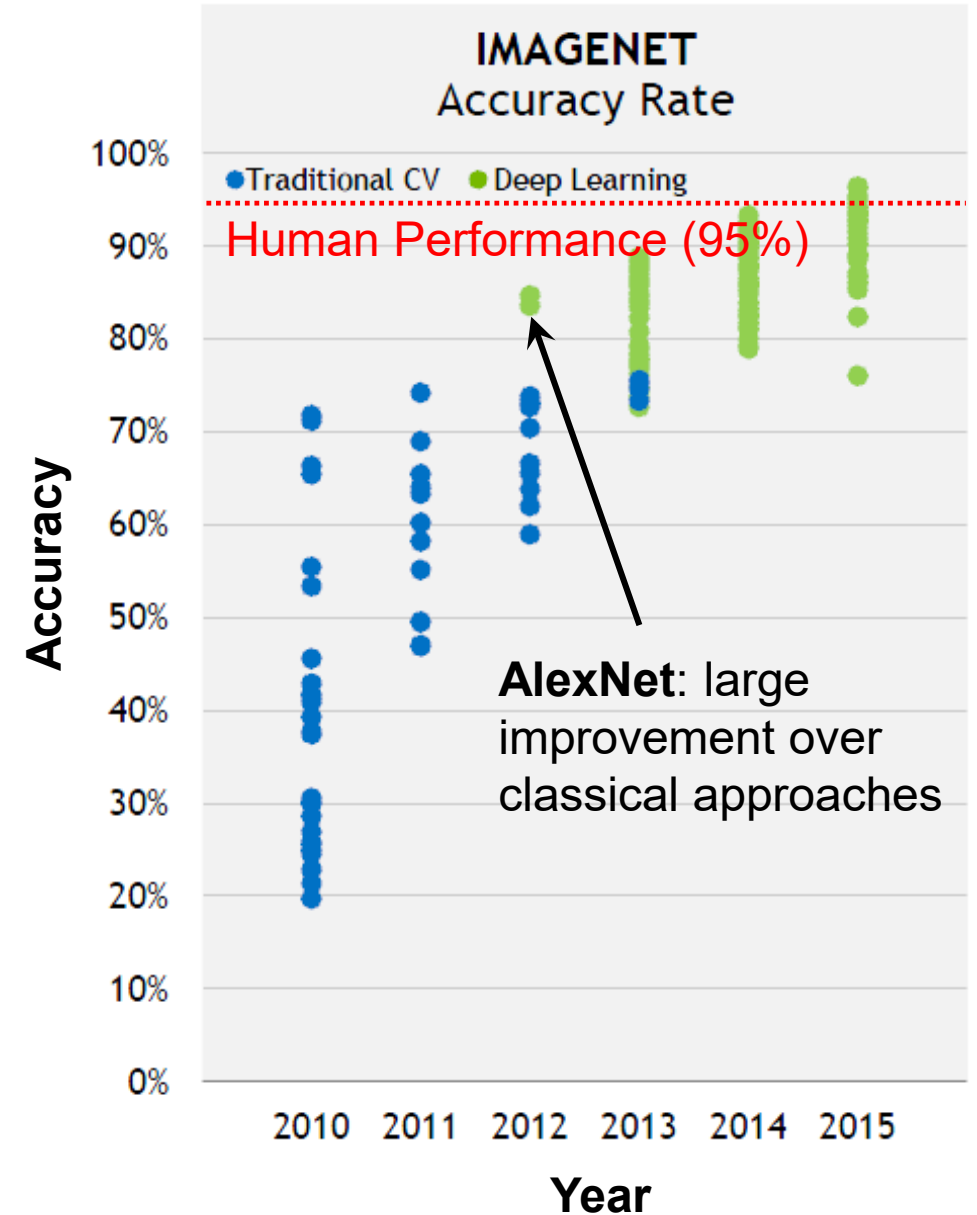


Translation

# DEEP LEARNING AND IMAGE CLASSIFICATION



Alex Krizhevsky, "Learning Multiple Layers of Features from Tiny Images", 2009.

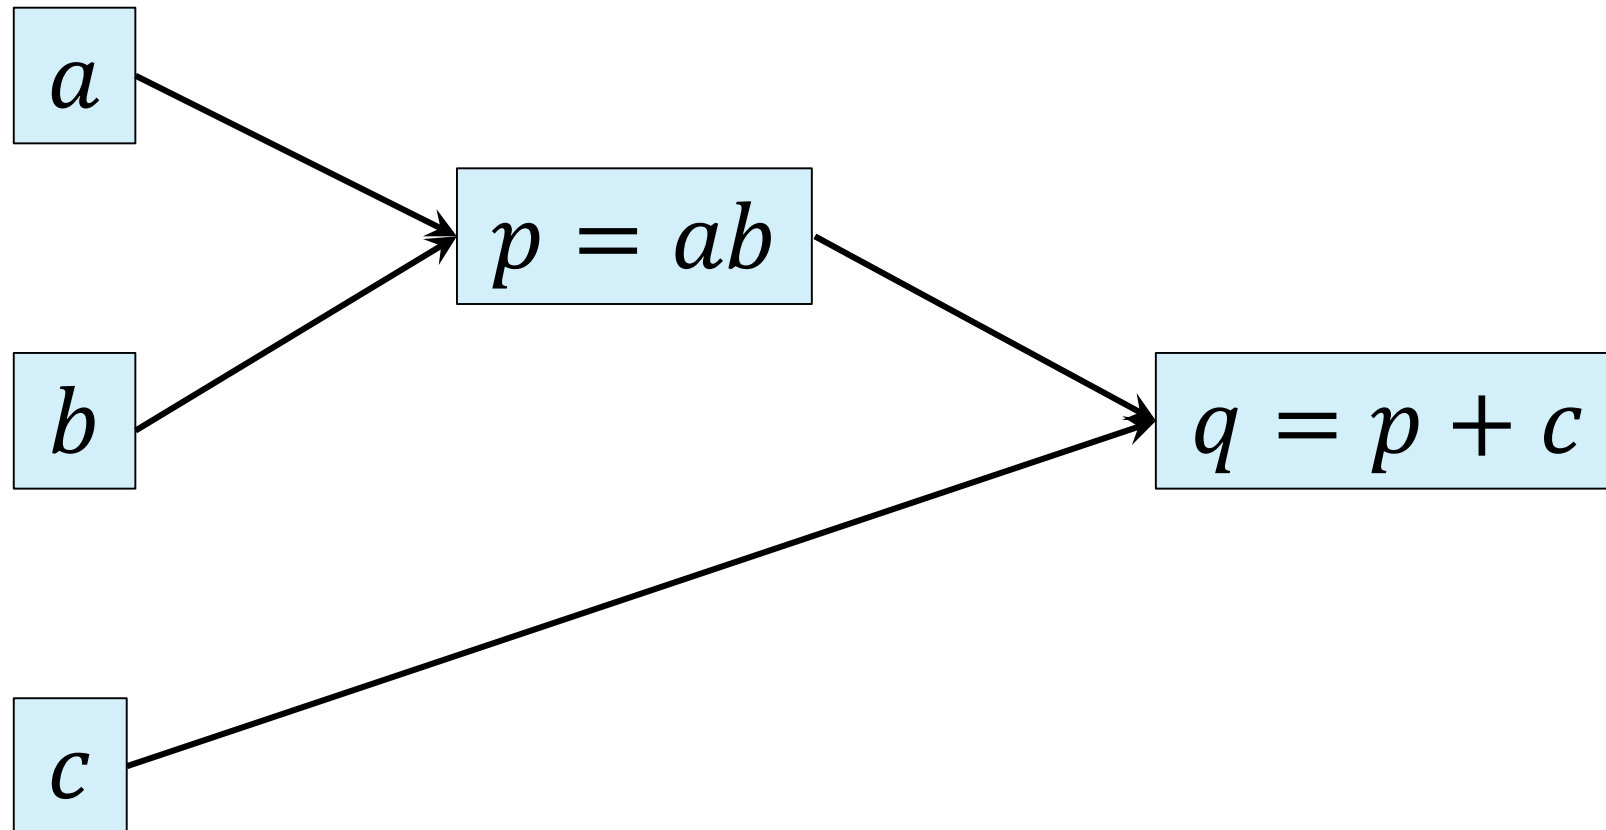




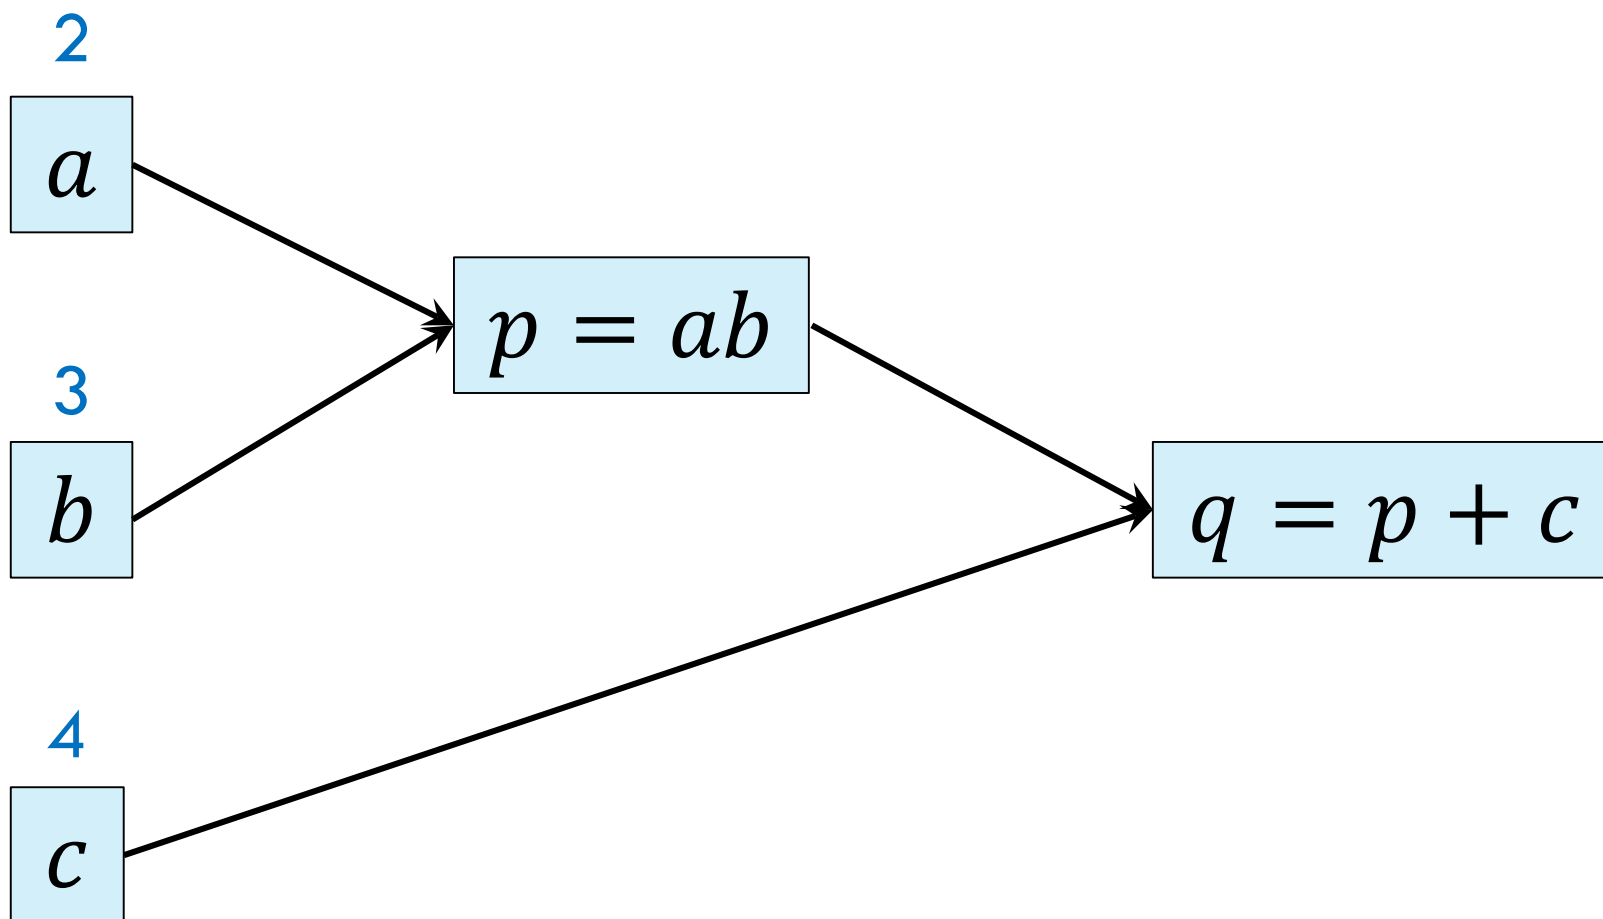
# BUILDING COMPLEX CLASSIFIERS FROM SIMPLE PARTS



# COMPUTATIONAL GRAPH

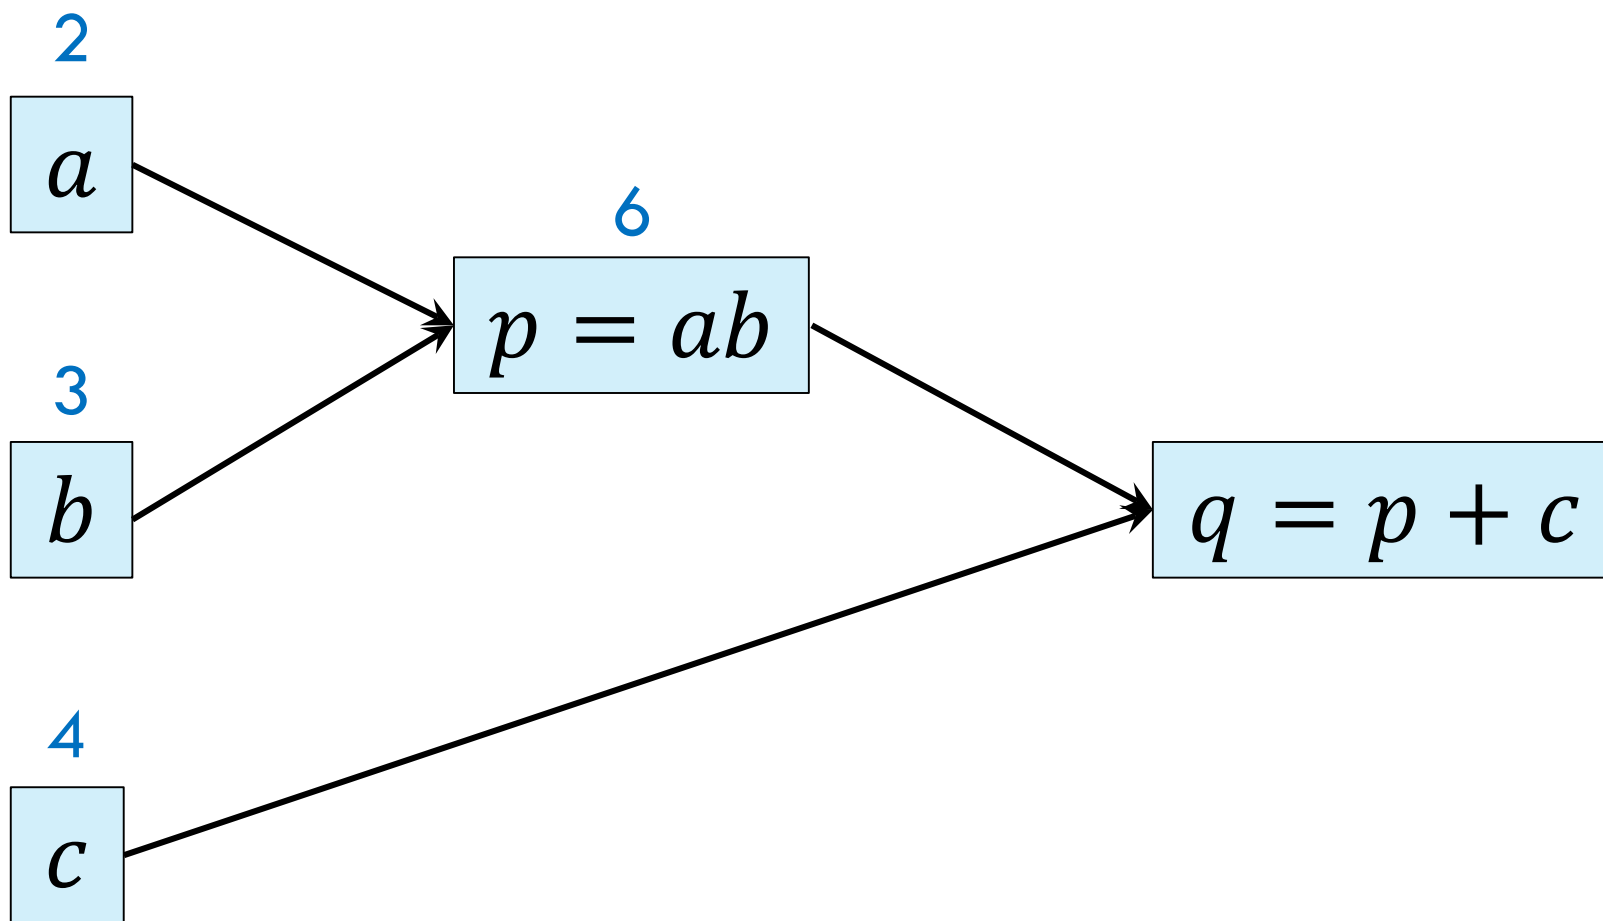


# COMPUTATIONAL GRAPH: FORWARD PASS

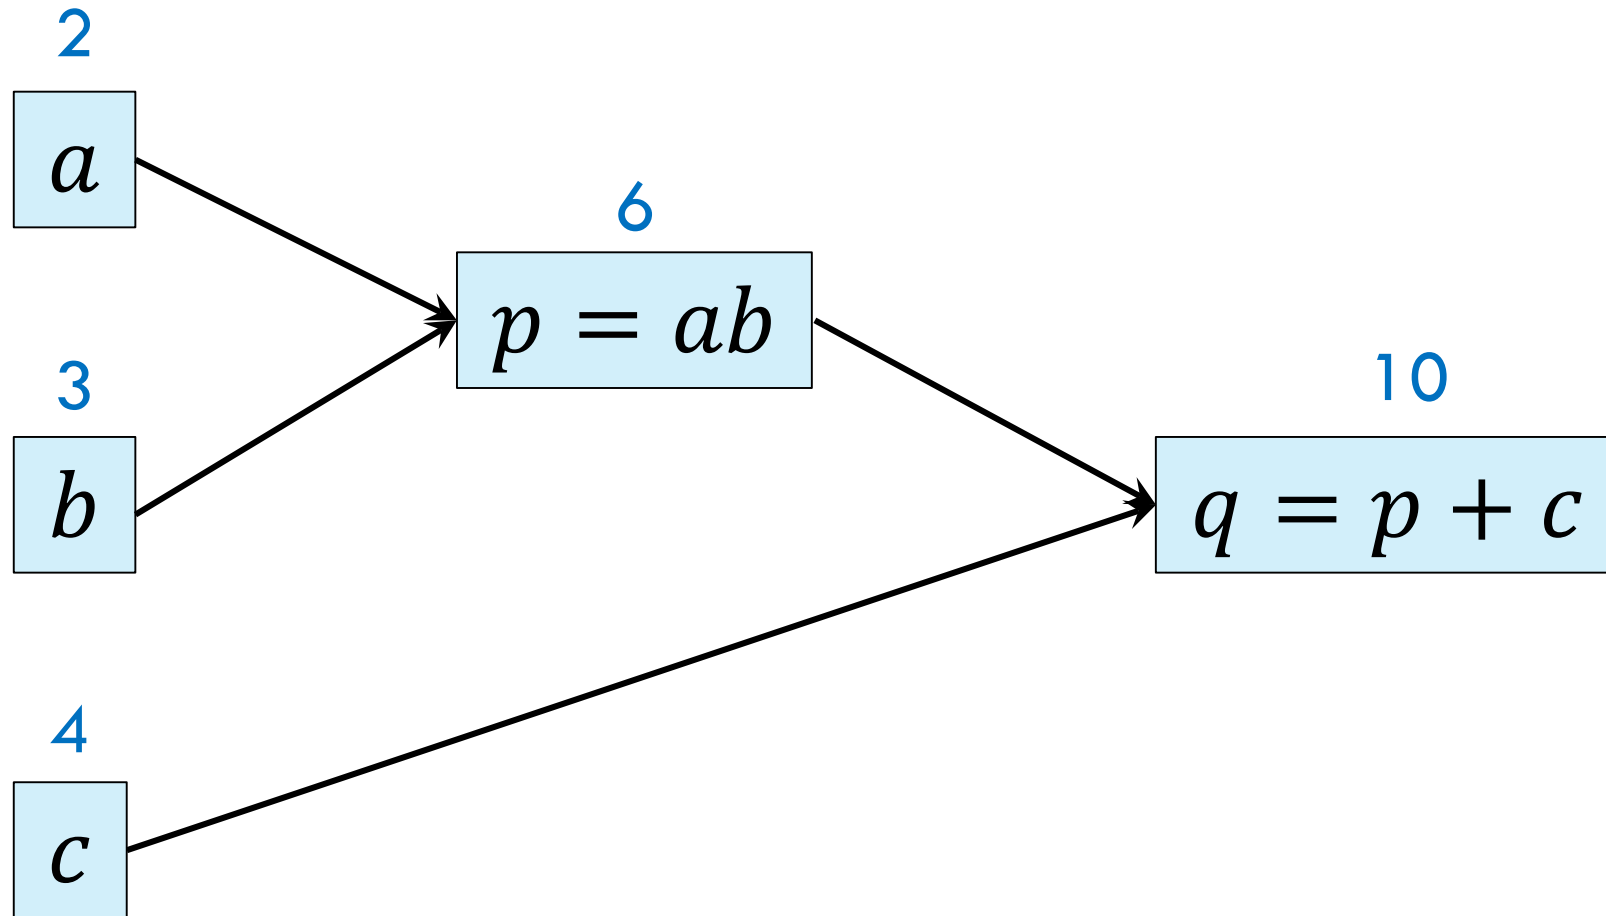




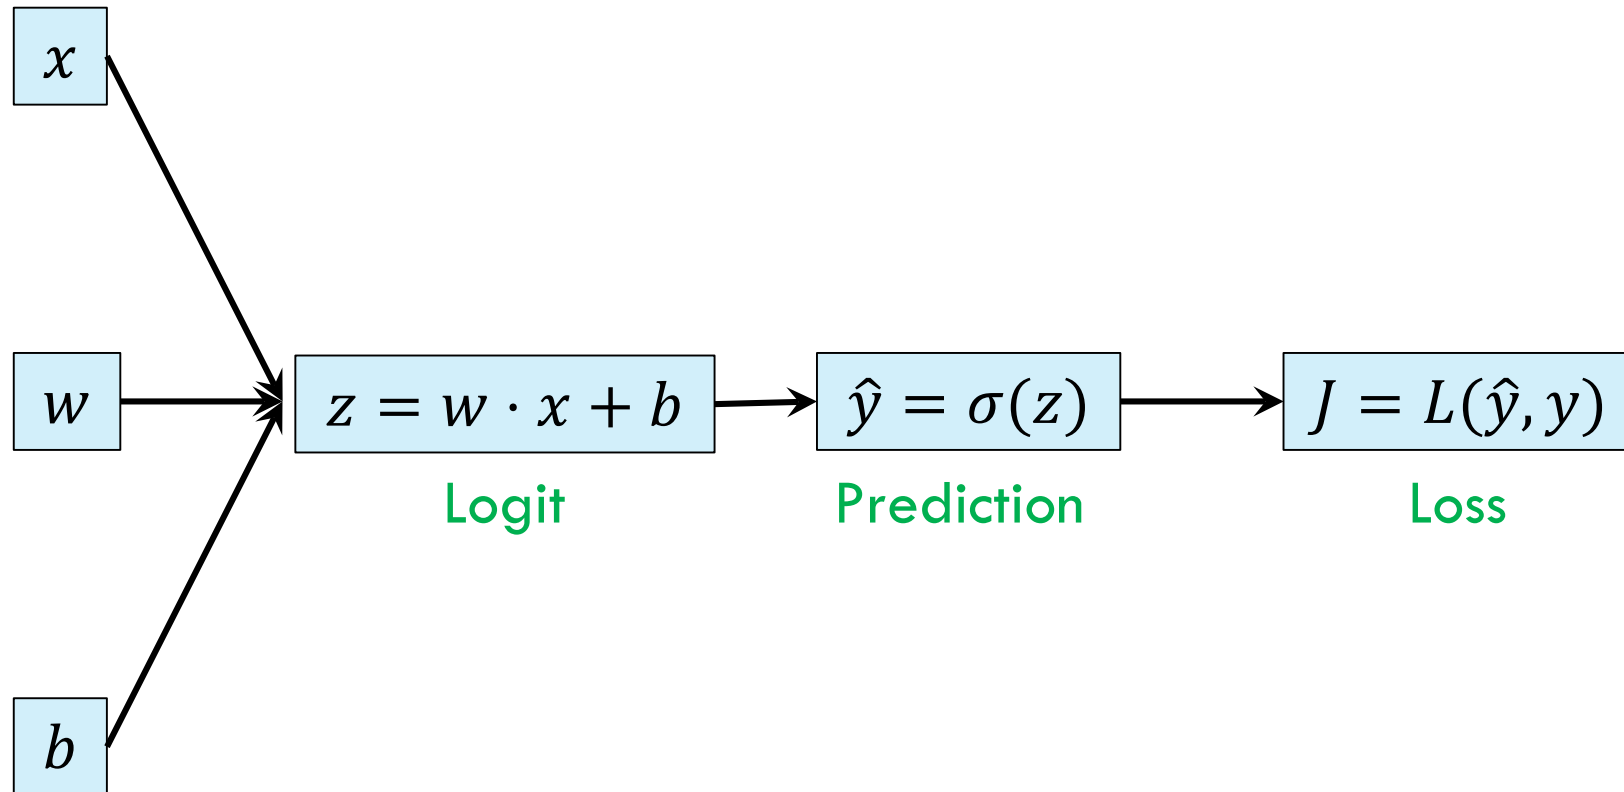
# COMPUTATIONAL GRAPH: FORWARD PASS



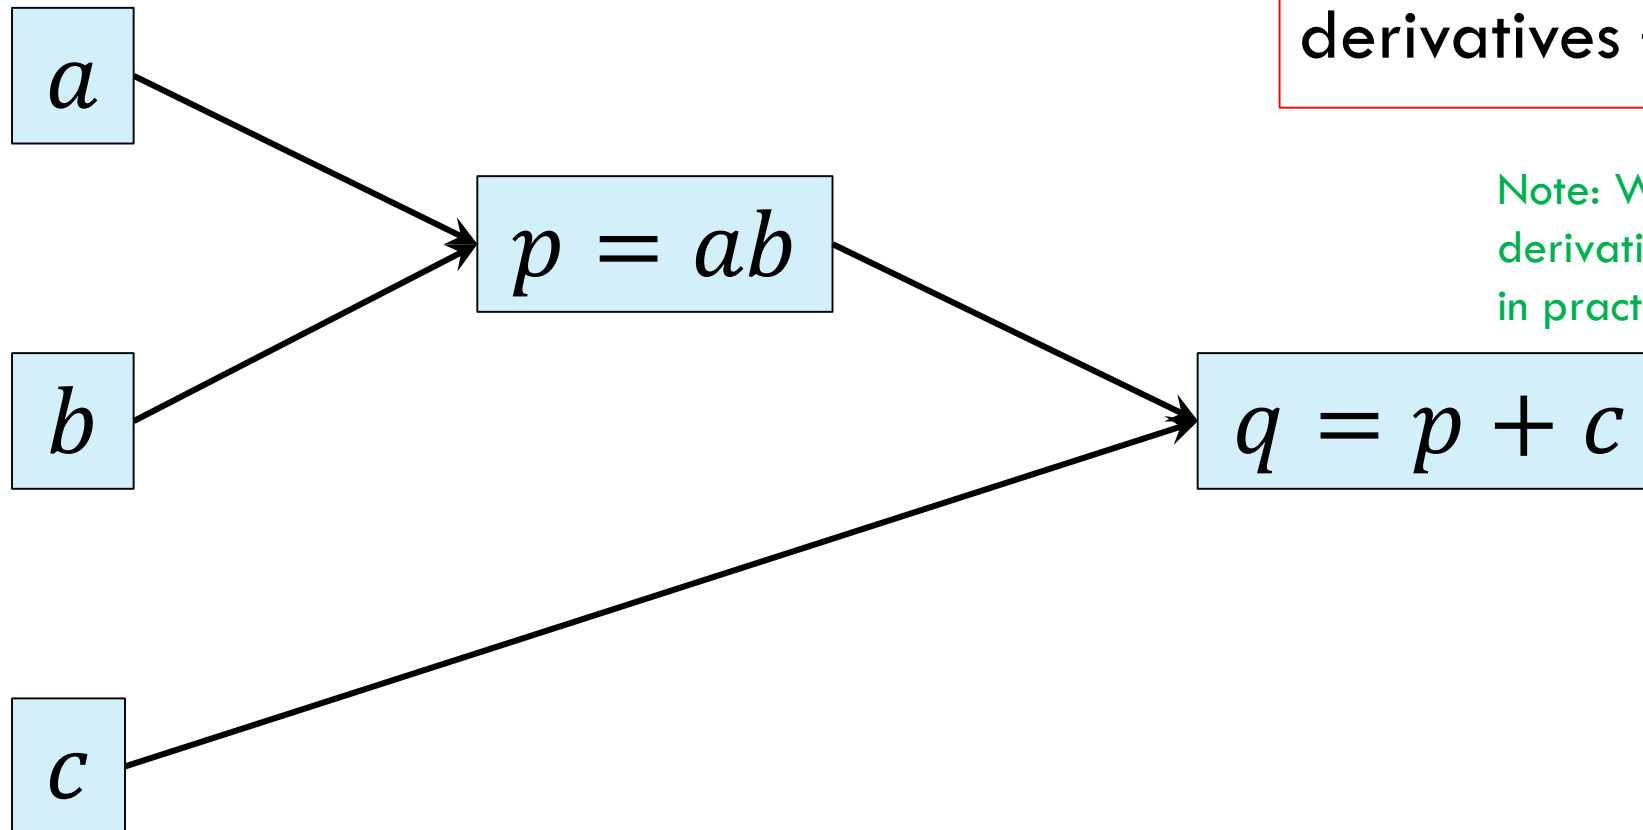
# COMPUTATIONAL GRAPH: FORWARD PASS



# LOGISTIC REGRESSION AS A COMPUTATIONAL GRAPH



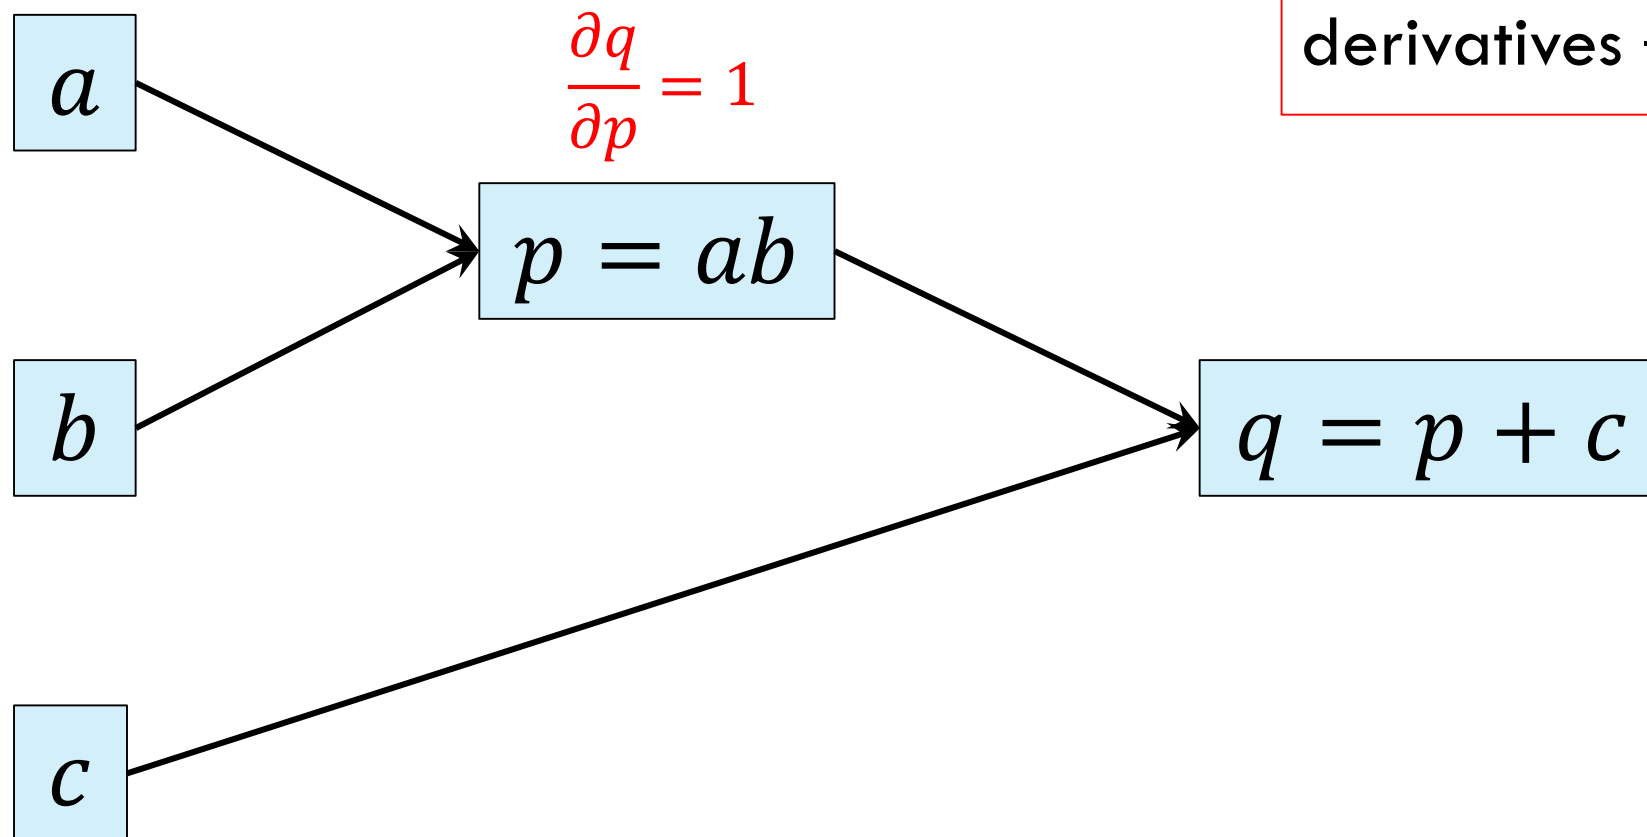
# COMPUTATIONAL GRAPH: BACKWARD PASS



How to compute the partial derivatives  $\frac{\partial q}{\partial a}$ ,  $\frac{\partial q}{\partial b}$ , etc.?

Note: We are mostly interested in the derivatives of the last variable (which in practice is almost always the loss)

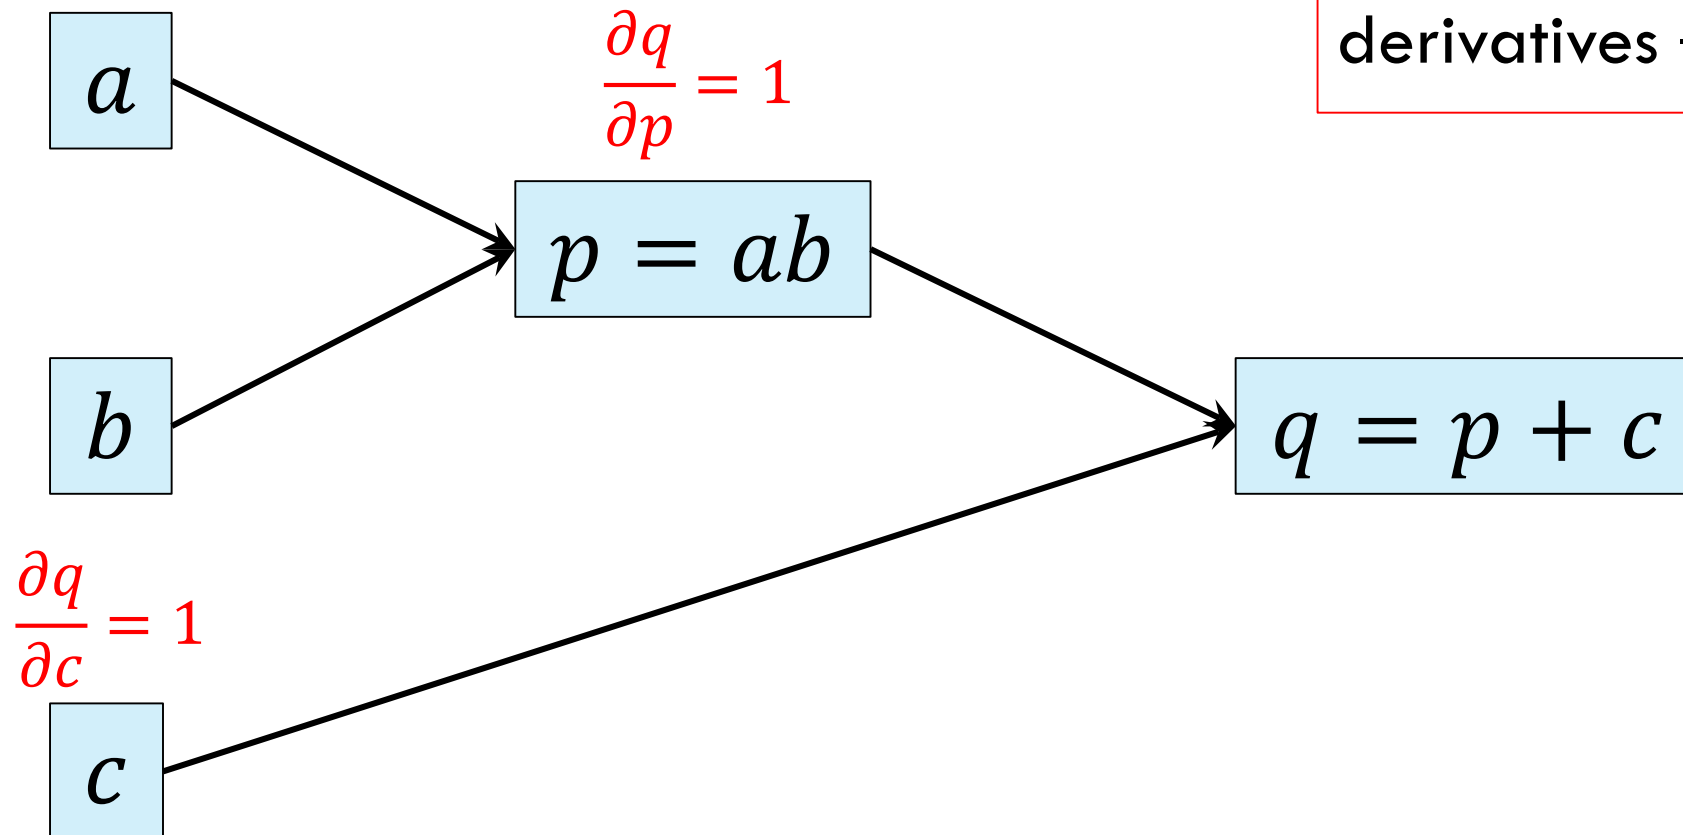
# COMPUTATIONAL GRAPH: BACKWARD PASS



How to compute the partial derivatives  $\frac{\partial q}{\partial a}$ ,  $\frac{\partial q}{\partial b}$ , etc.?

# COMPUTATIONAL GRAPH: BACKWARD PASS

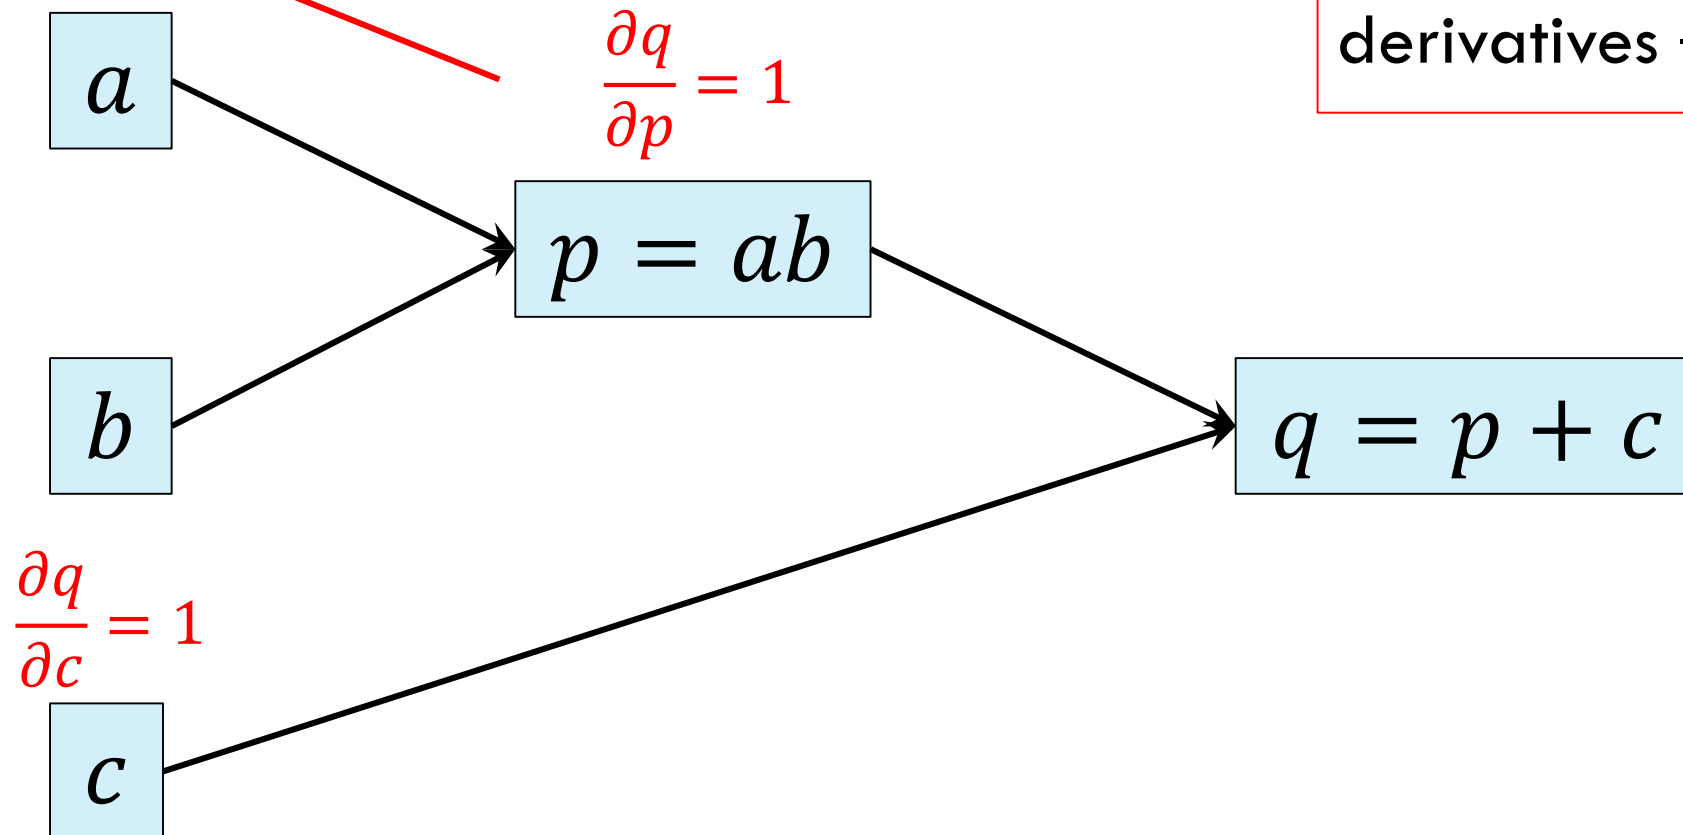
How to compute the partial derivatives  $\frac{\partial q}{\partial a}$ ,  $\frac{\partial q}{\partial b}$ , etc.?



# COMPUTATIONAL GRAPH: BACKWARD PASS

$$\frac{\partial q}{\partial a} = \frac{\partial q}{\partial p} \cdot \frac{\partial p}{\partial a} = 1 \cdot b$$

How to compute the partial derivatives  $\frac{\partial q}{\partial a}$ ,  $\frac{\partial q}{\partial b}$ , etc.?



# COMPUTATIONAL GRAPH: BACKWARD PASS

$$\frac{\partial q}{\partial a} = \frac{\partial q}{\partial p} \cdot \frac{\partial p}{\partial a} = 1 \cdot b$$

$a$

$$\frac{\partial q}{\partial p} = 1$$

$$\frac{\partial q}{\partial b} = \frac{\partial q}{\partial p} \cdot \frac{\partial p}{\partial b} = 1 \cdot a$$

$b$

$p = ab$

$$\frac{\partial q}{\partial c} = 1$$

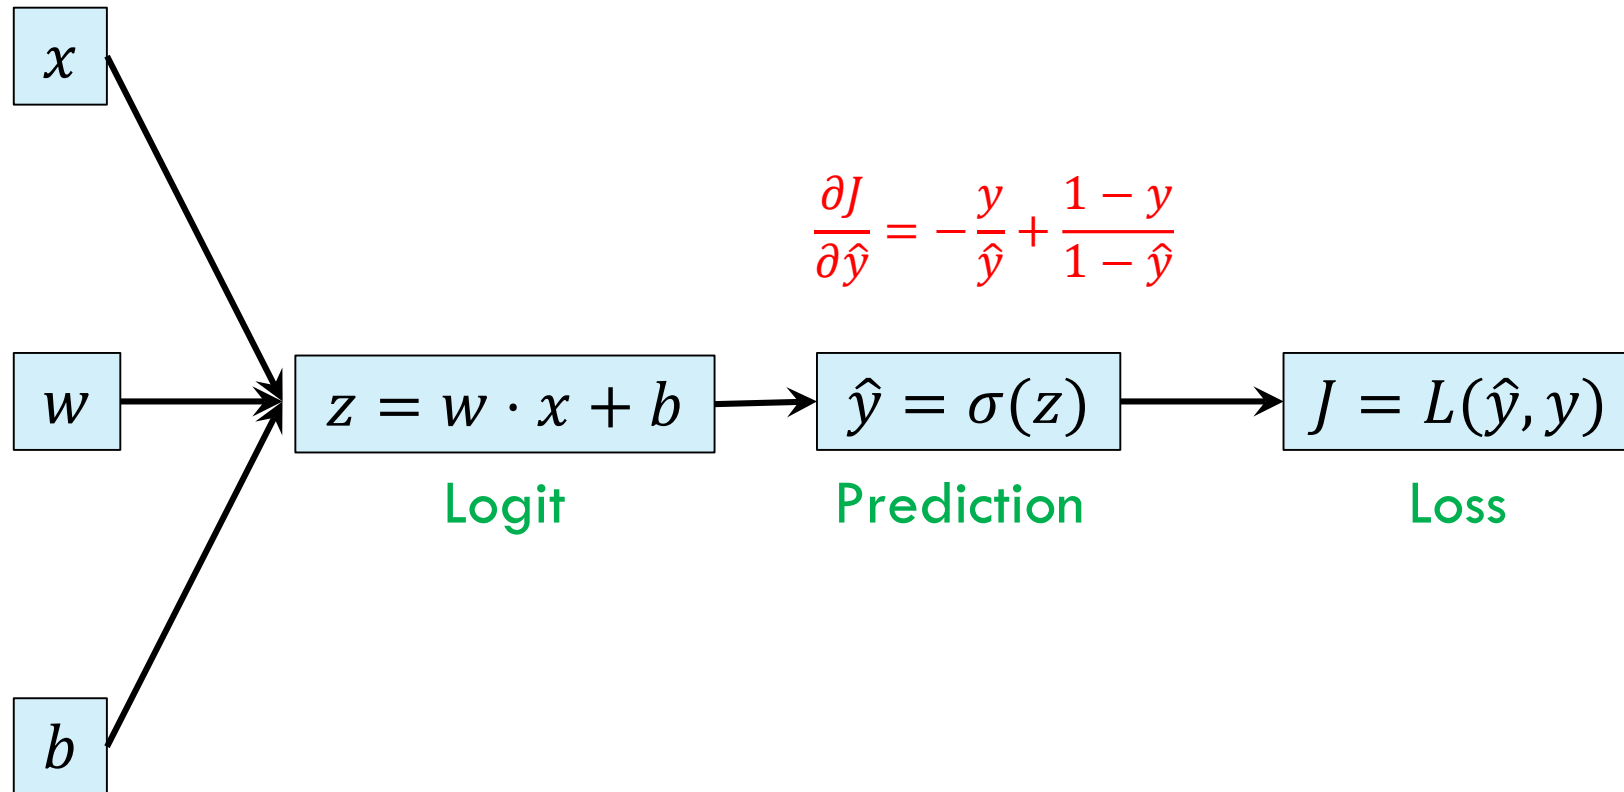
$c$

$q = p + c$

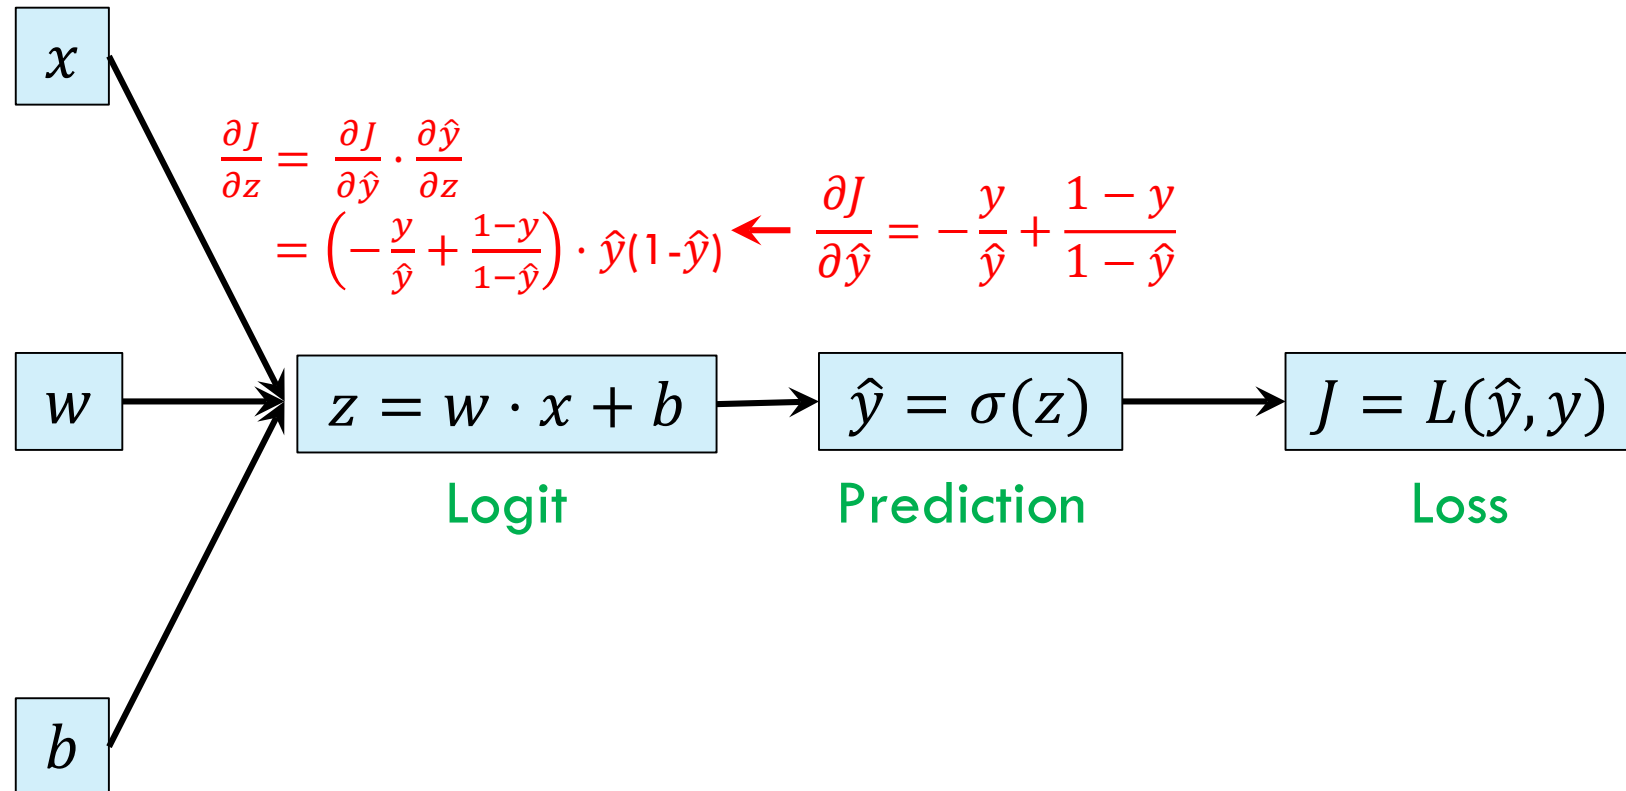
How to compute the partial derivatives  $\frac{\partial q}{\partial a}$ ,  $\frac{\partial q}{\partial b}$ , etc.?



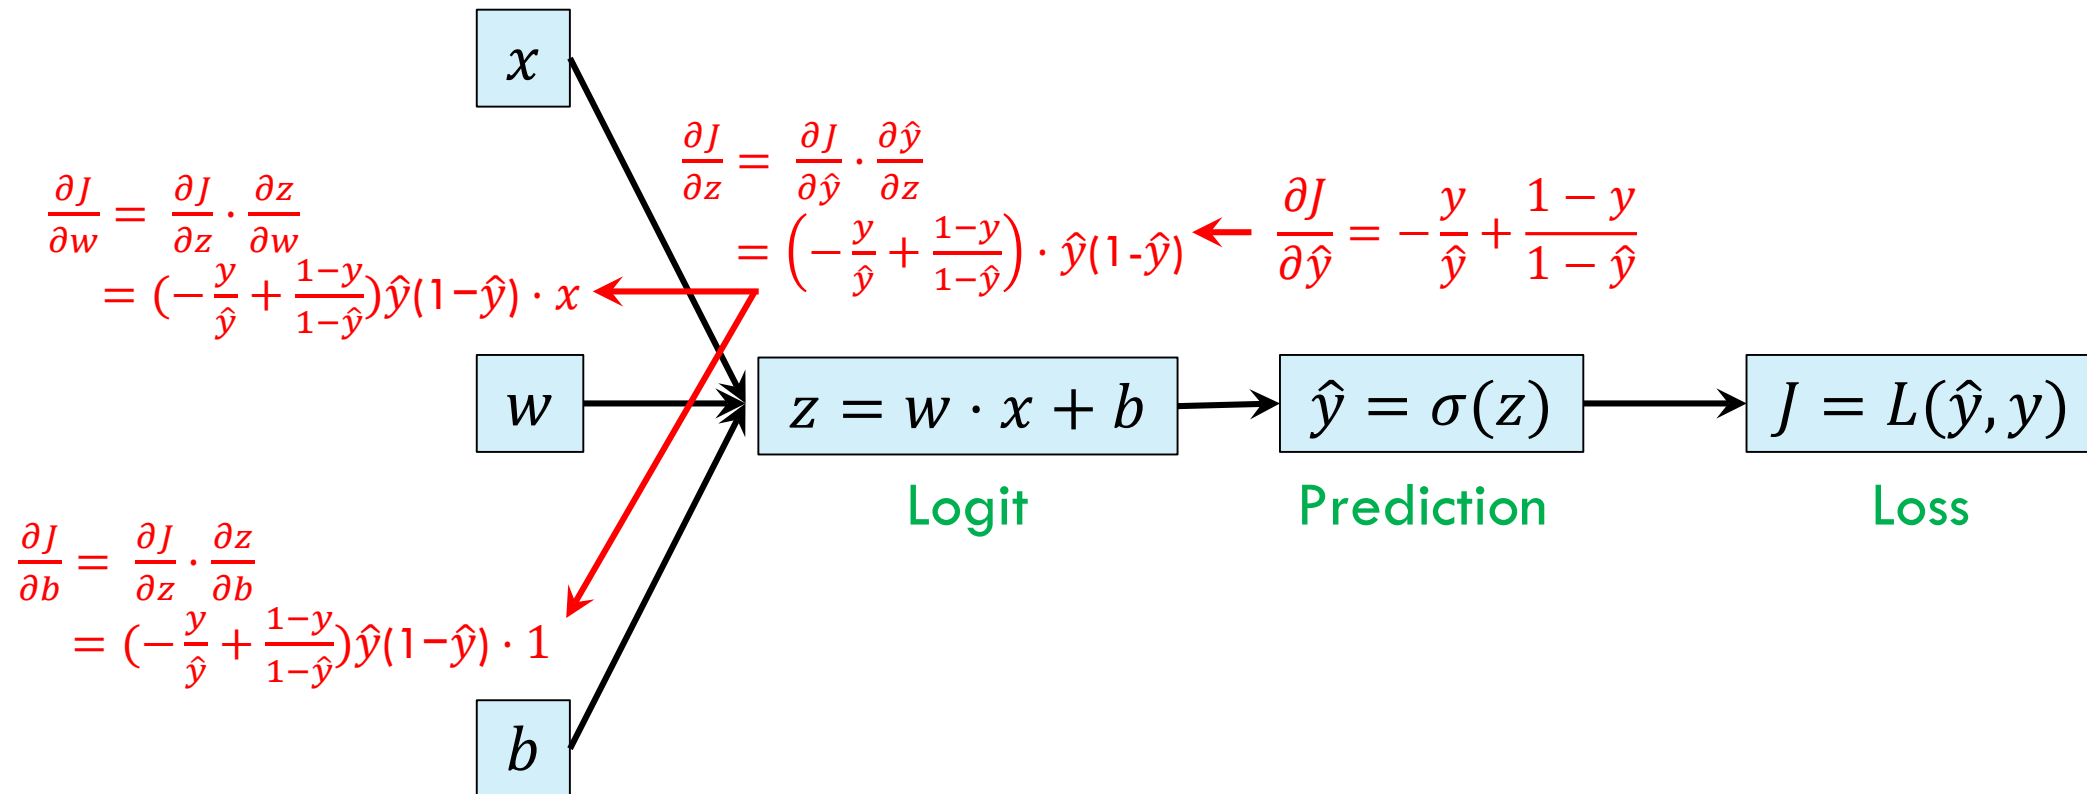
# LOGISTIC REGRESSION AS A COMPUTATIONAL GRAPH



# LOGISTIC REGRESSION AS A COMPUTATIONAL GRAPH



# LOGISTIC REGRESSION AS A COMPUTATIONAL GRAPH

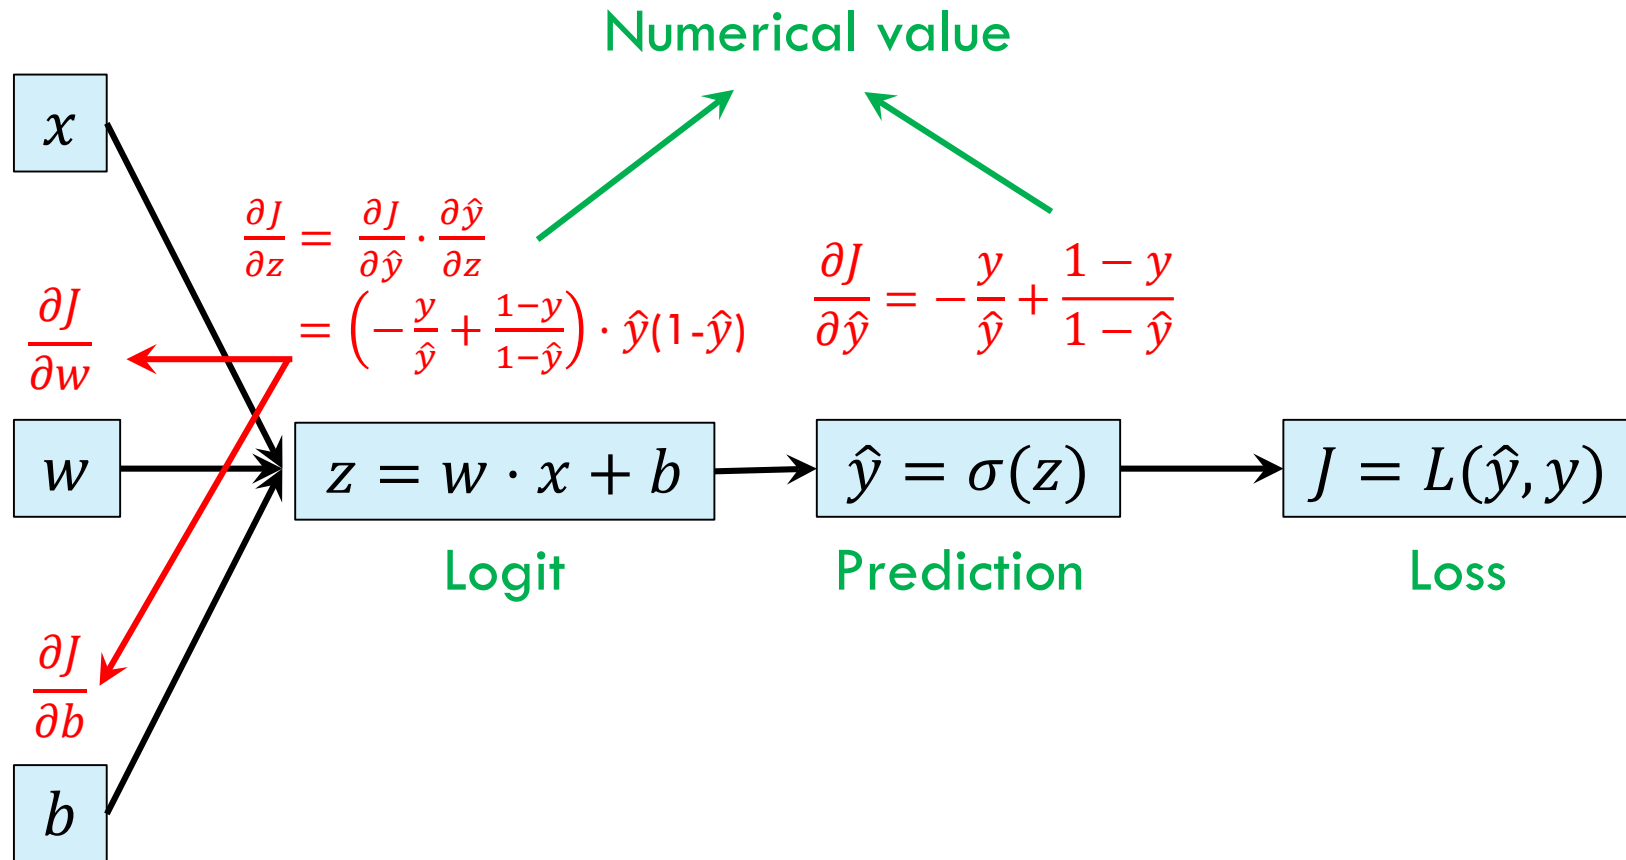


# LOGISTIC REGRESSION AS A COMPUTATIONAL GRAPH

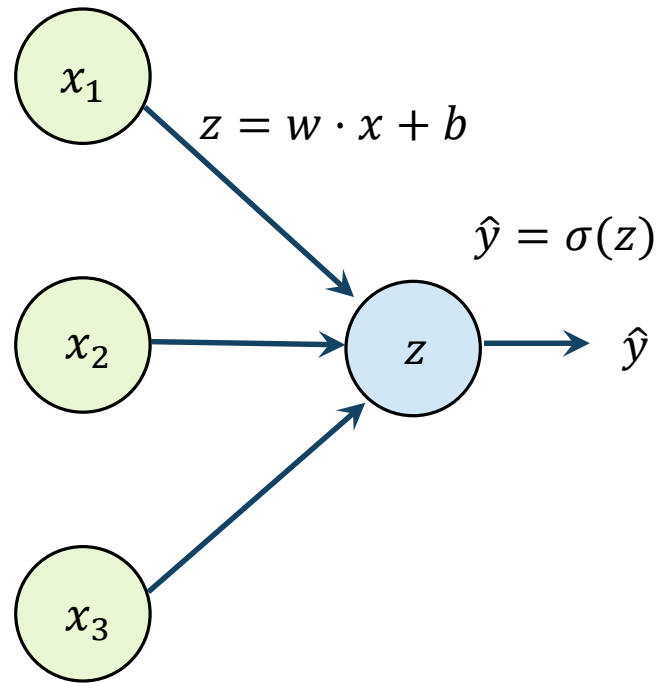
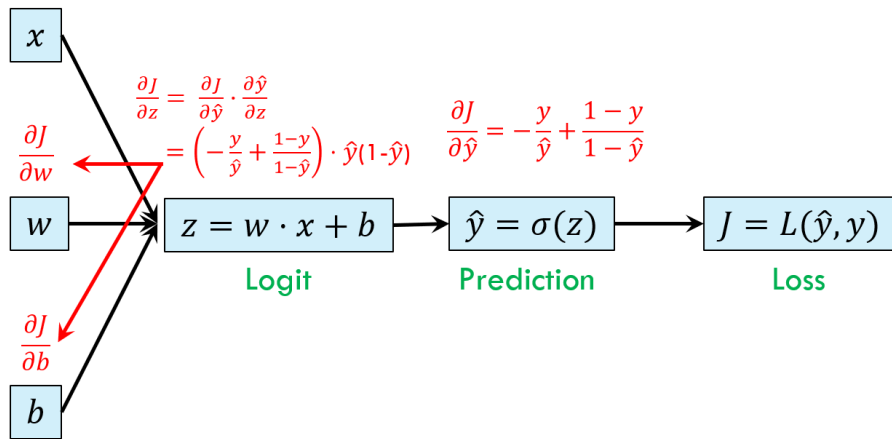
## Gradient Descent Steps

$$w \leftarrow w - \eta \frac{\partial J(w, b)}{\partial w}$$

$$b \leftarrow b - \eta \frac{\partial J(w, b)}{\partial b}$$



# LOGISTIC REGRESSION AS A NEURAL NETWORK



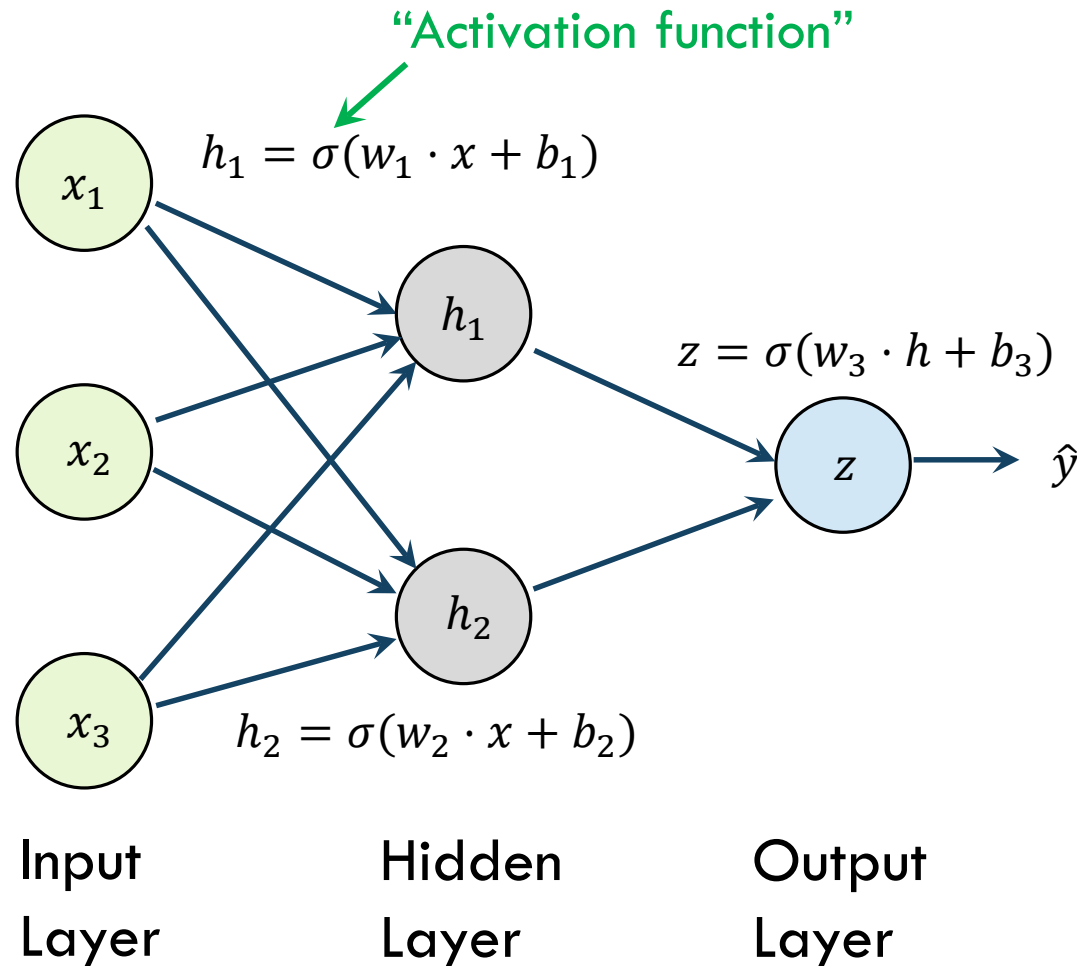
Input  
Layer

Output  
Layer

**Layers of neurons:** each layer computes a function of the previous layer

This is a 1-layer network (input layer is not counted)

# DEEPER NEURAL NETWORKS



Each hidden unit (“neuron”) has its own weights and biases

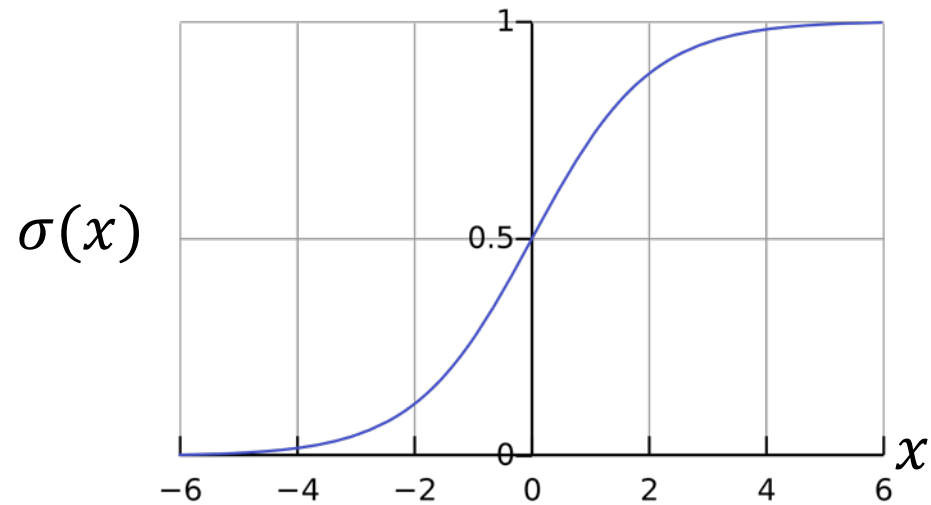
These parameters are used to compute a linear function of the previous layer

But they apply a **nonlinear activation function**, similar to the sigmoid in logistic regression

- The resulting models are much **more expressive** than logistic regression / linear models

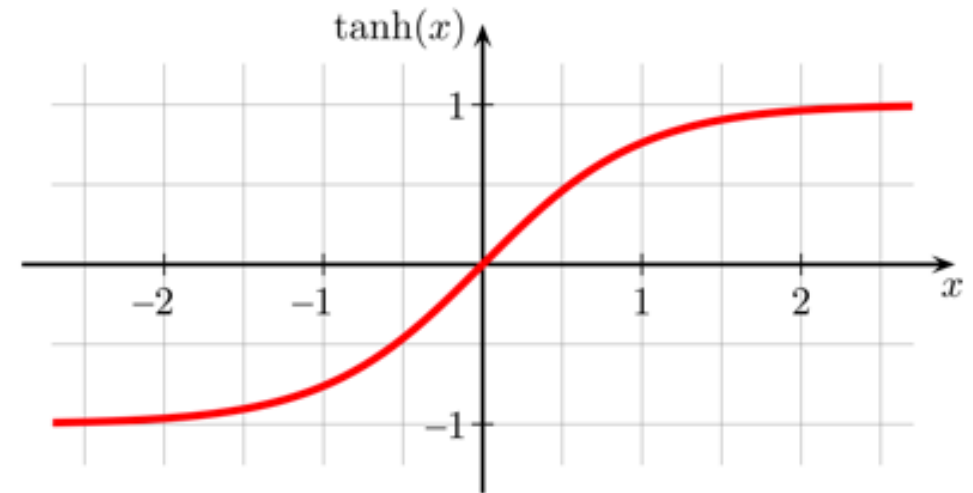
# ACTIVATION FUNCTIONS

## Sigmoid



$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

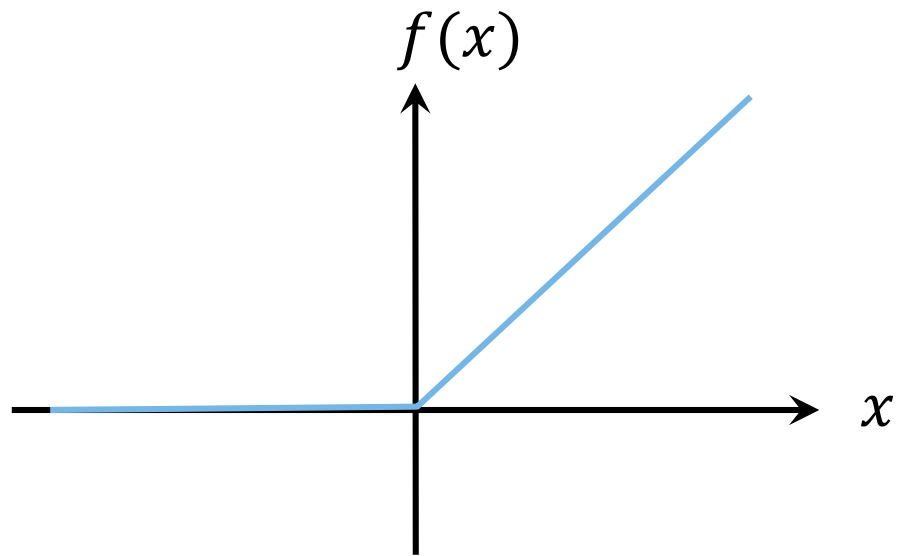
## Tanh



$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

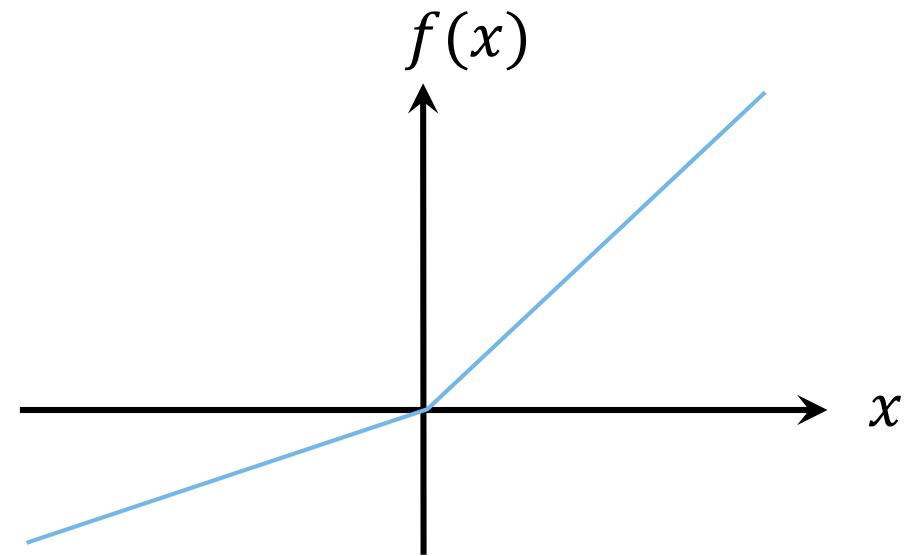
# ACTIVATION FUNCTIONS

## Rectified Linear Unit (ReLU)



$$f(x) = \max(0, x)$$

## Parametric ReLU



$$f(x) = \max(ax, x) \quad (0 < a < 1)$$



# MANY TRICKS

## Initialization

- Random, Weight initialization, Xavier initialization, ...

## Adaptive learning rates

- Decay, Momentum, RMSProp, Adam, ...

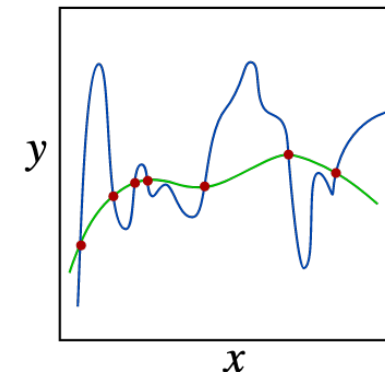
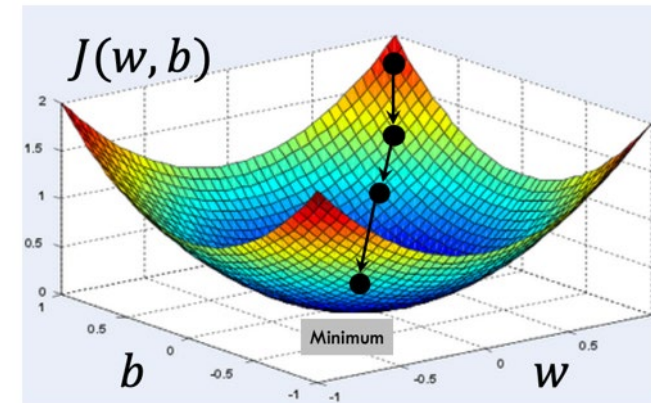
## Batch normalization

## Regularization

- Early stopping, dropout, L1 / L2 regularization, data augmentation, ...

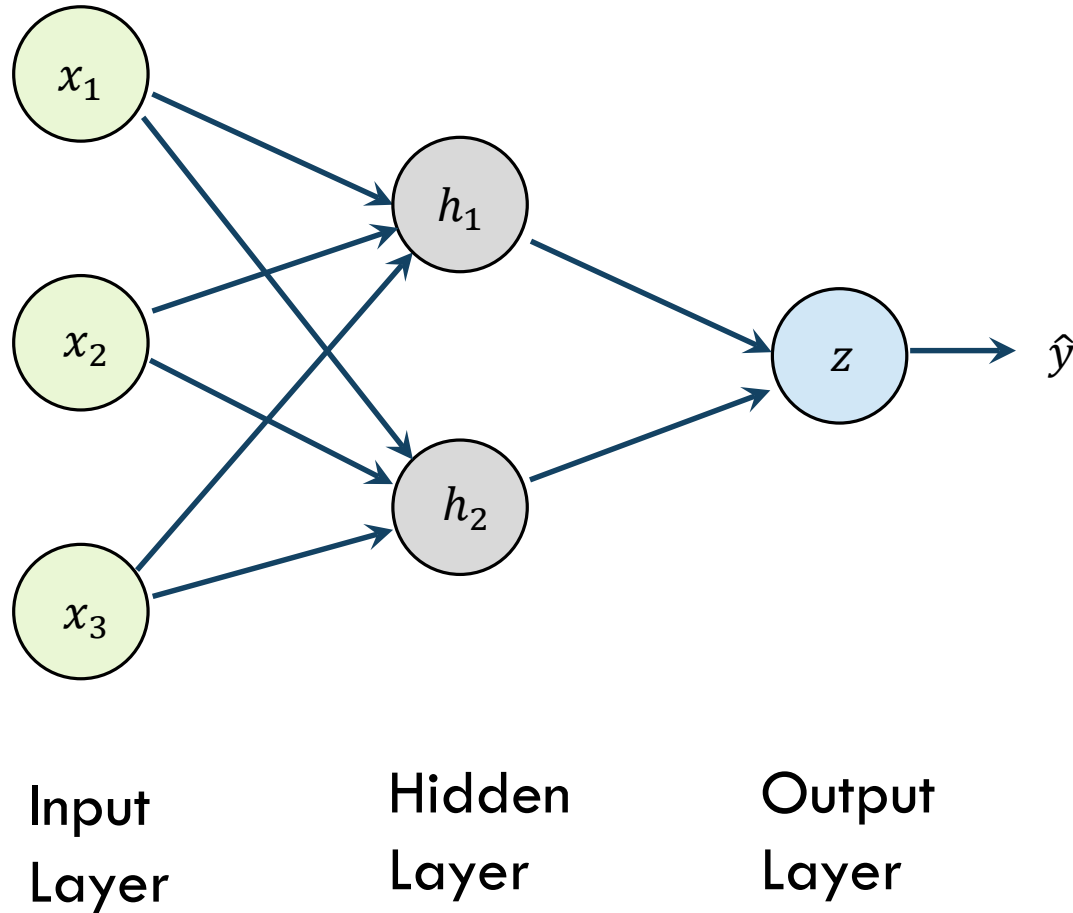
Speed up  
learning

Reduce  
overfitting





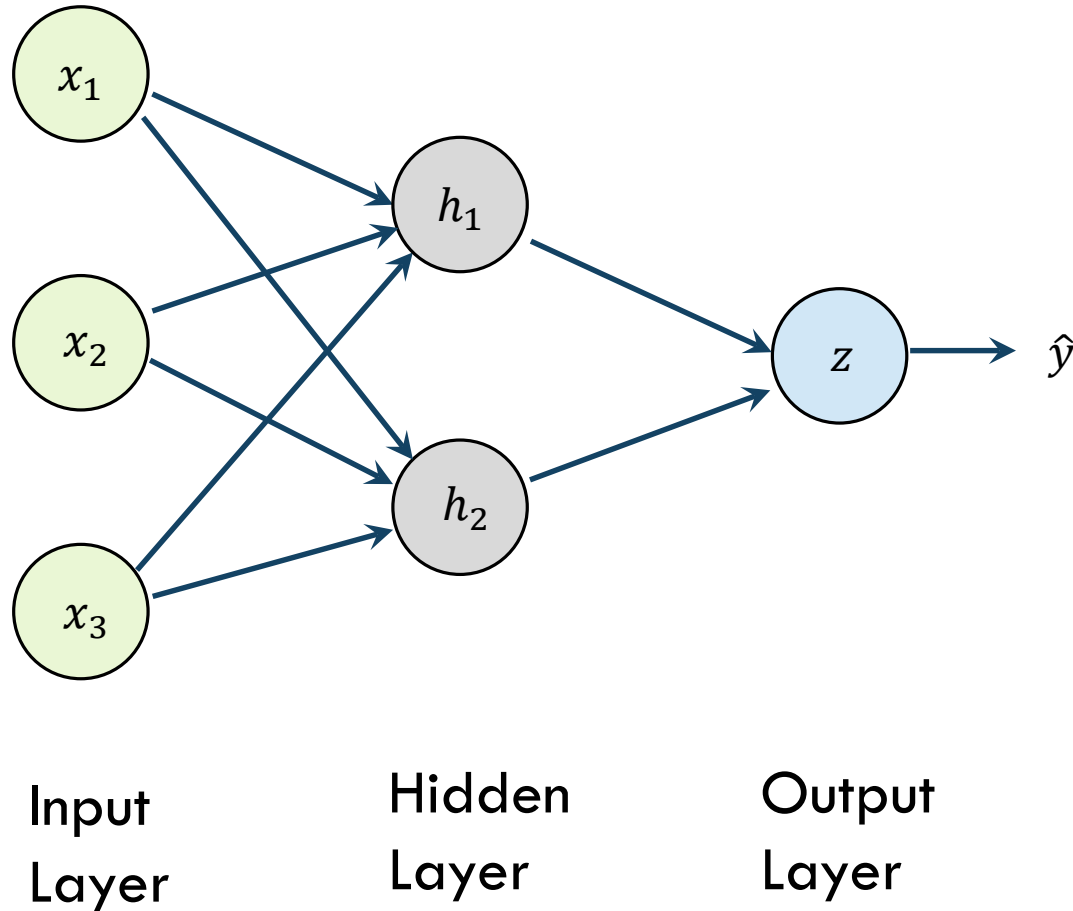
# QUIZ: HOW MANY PARAMETERS?



**Q:** In this neural network, how many weight parameters (i.e. not bias parameters) in total?



# QUIZ: HOW MANY PARAMETERS?

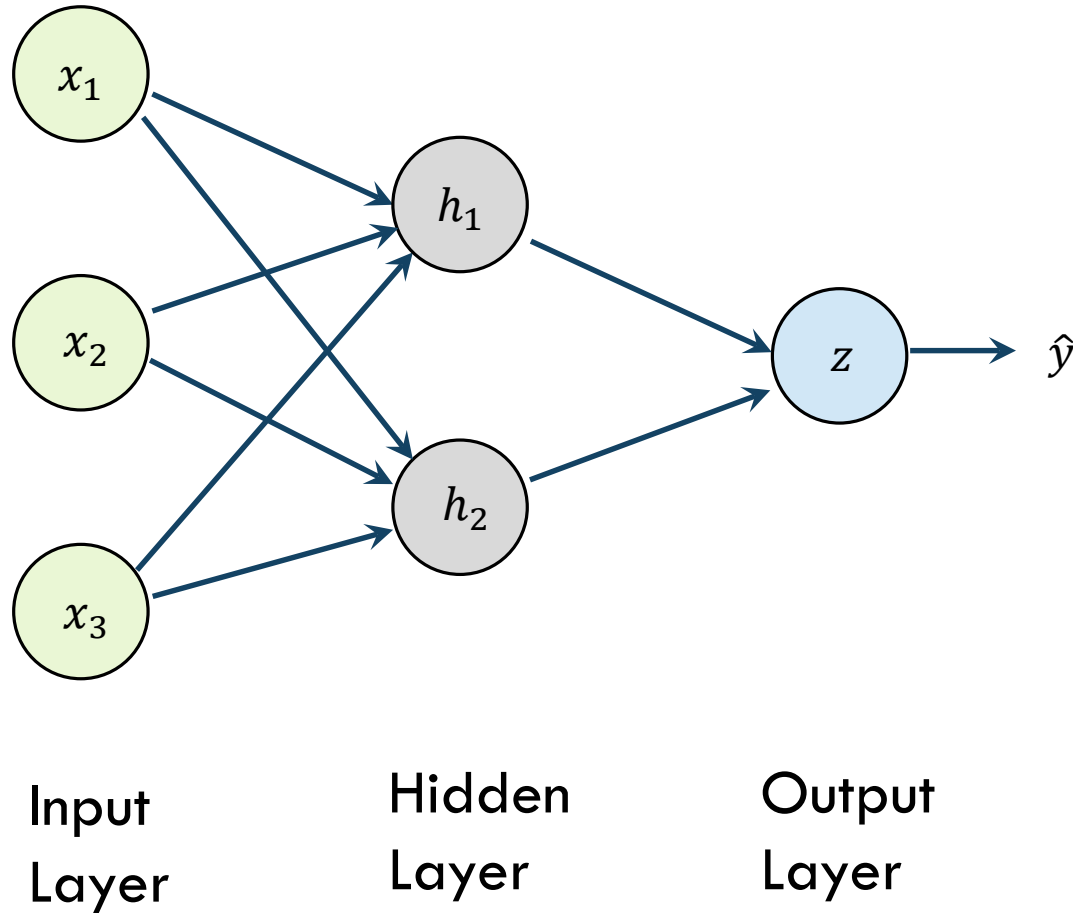


**Q:** In this neural network, how many weight parameters (i.e. not bias parameters) in total?

**A:** 8 ( $3 \times 2 = 6$  between the input and hidden layer, and 2 between the hidden and output layer)



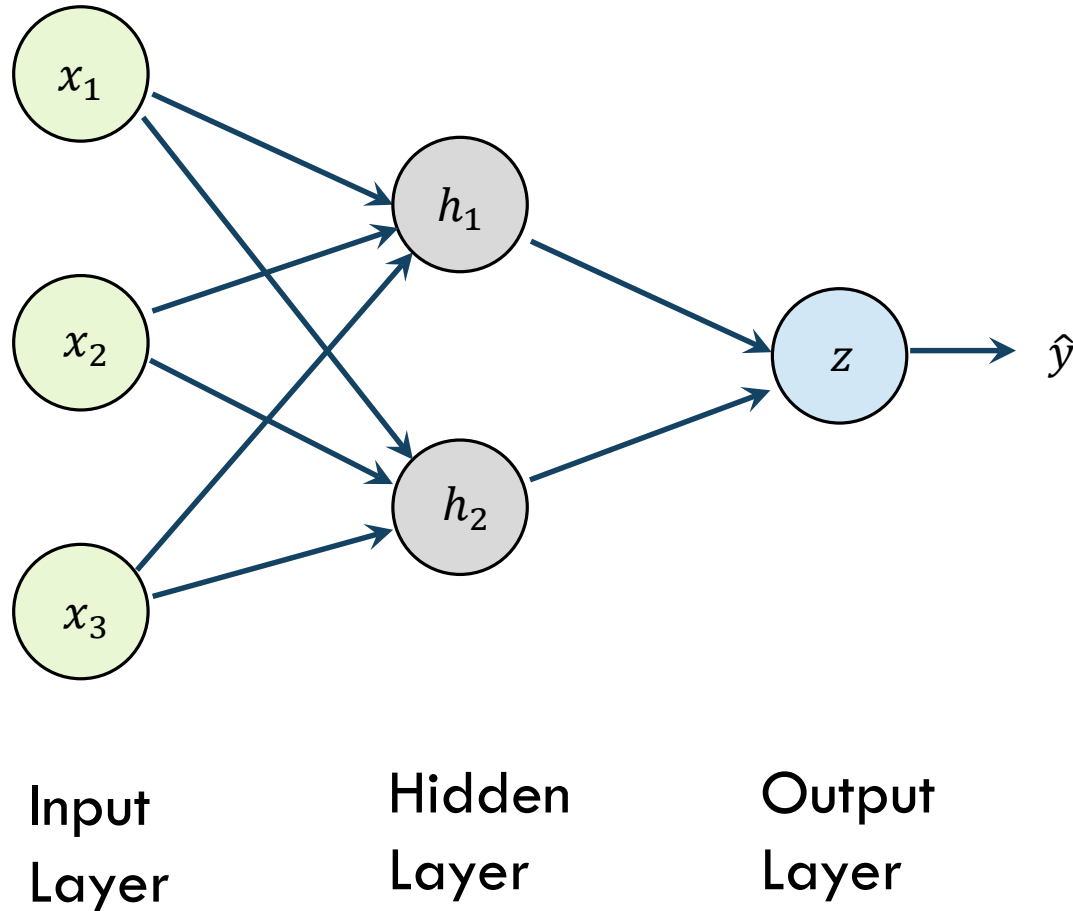
# QUIZ: HOW MANY PARAMETERS?



**Q:** In this neural network, how many bias parameters in total are there?



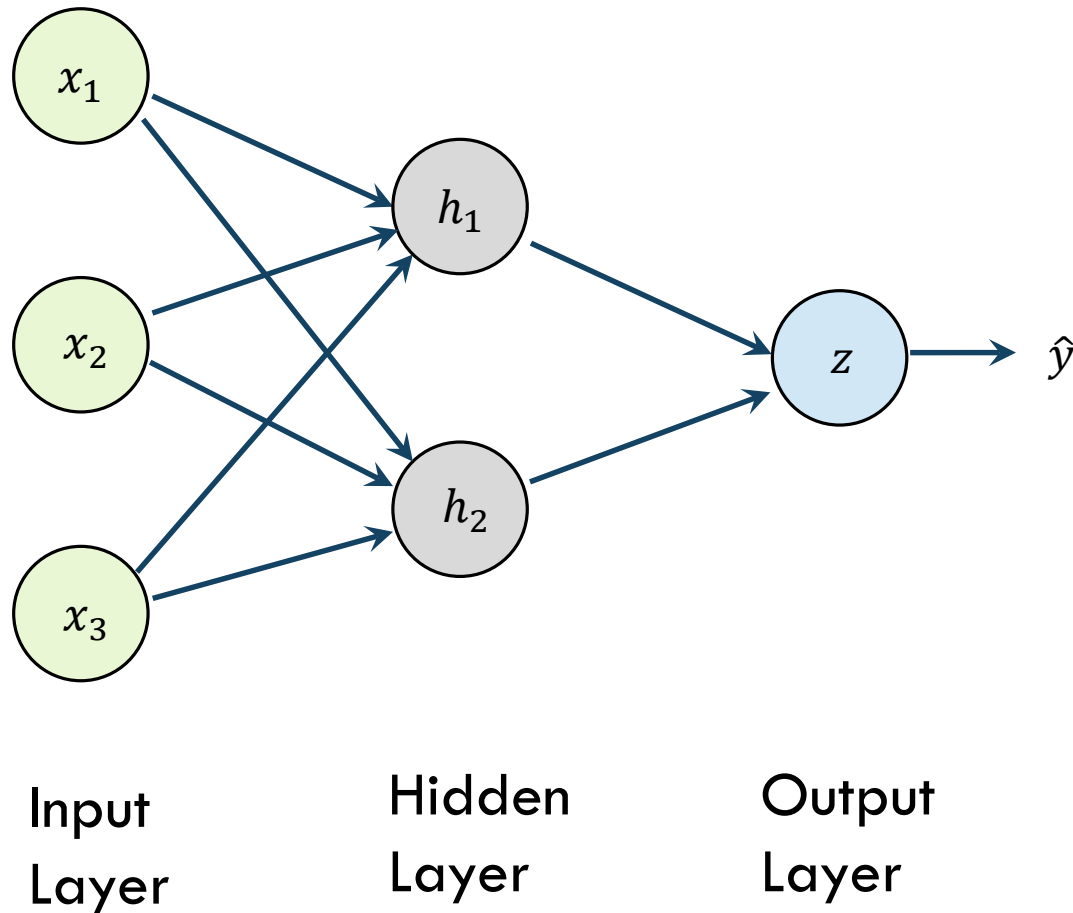
# QUIZ: HOW MANY PARAMETERS?



**Q:** In this neural network, how many bias parameters in total are there?

**A:** 3 (once for each hidden and output layer neuron)

# HOW MANY PARAMETERS?



## Fully connected layers:

- Number of weights:  $\# \text{ inputs} \times \# \text{ outputs}$
- Number of bias terms:  $\# \text{ outputs}$

## Memory usage:

- $4 \text{ bytes} \times \text{number of parameters}$
- Should fit in your GPU memory